BIG DATA SPAIN

Adding intelligence to stream-processing applications

Michael McCune and William Benton
msm@redhat.com • @FOSSJunkie
willb@redhat.com • @willb

Forecast

Introducing intelligent applications

Introducing microservice architectures

Stream processing with Apache Kafka and Apache Spark

Interactivity time

BIG DATA SPAIN

Intelligent applications

Intelligent applications **collect** and **learn from** data in order to provide **improved functionality** with **longevity** and **popularity**.

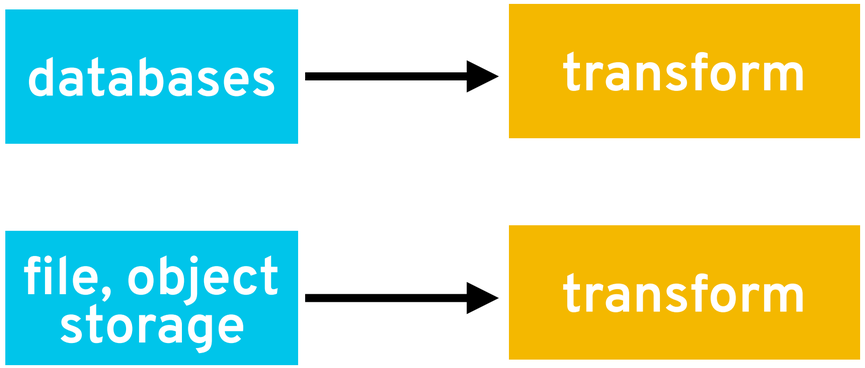
databases
transform
file, object storage
transform

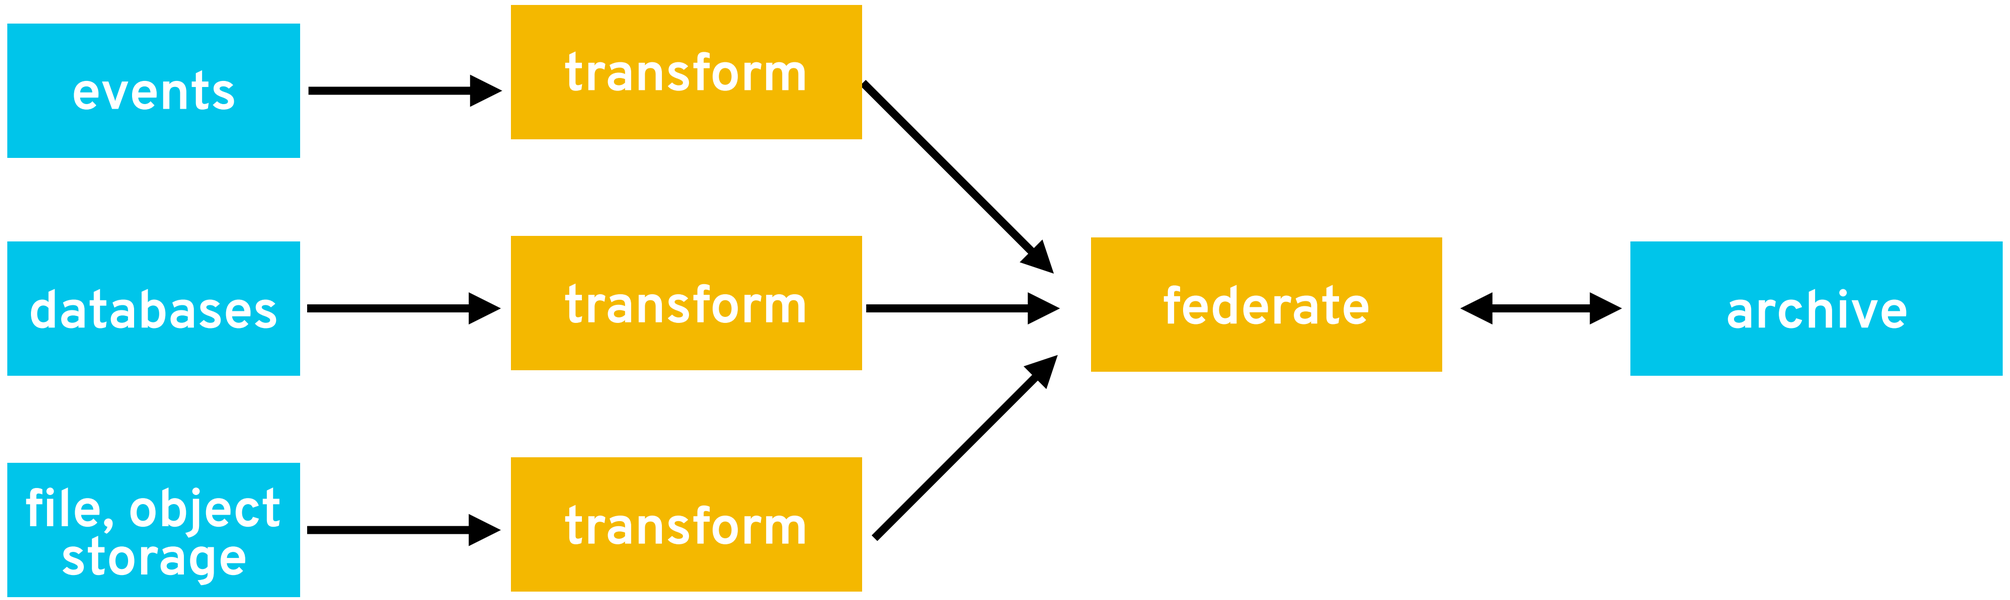
events
databases
file, object storage
transform
transform
transform
federate
archive

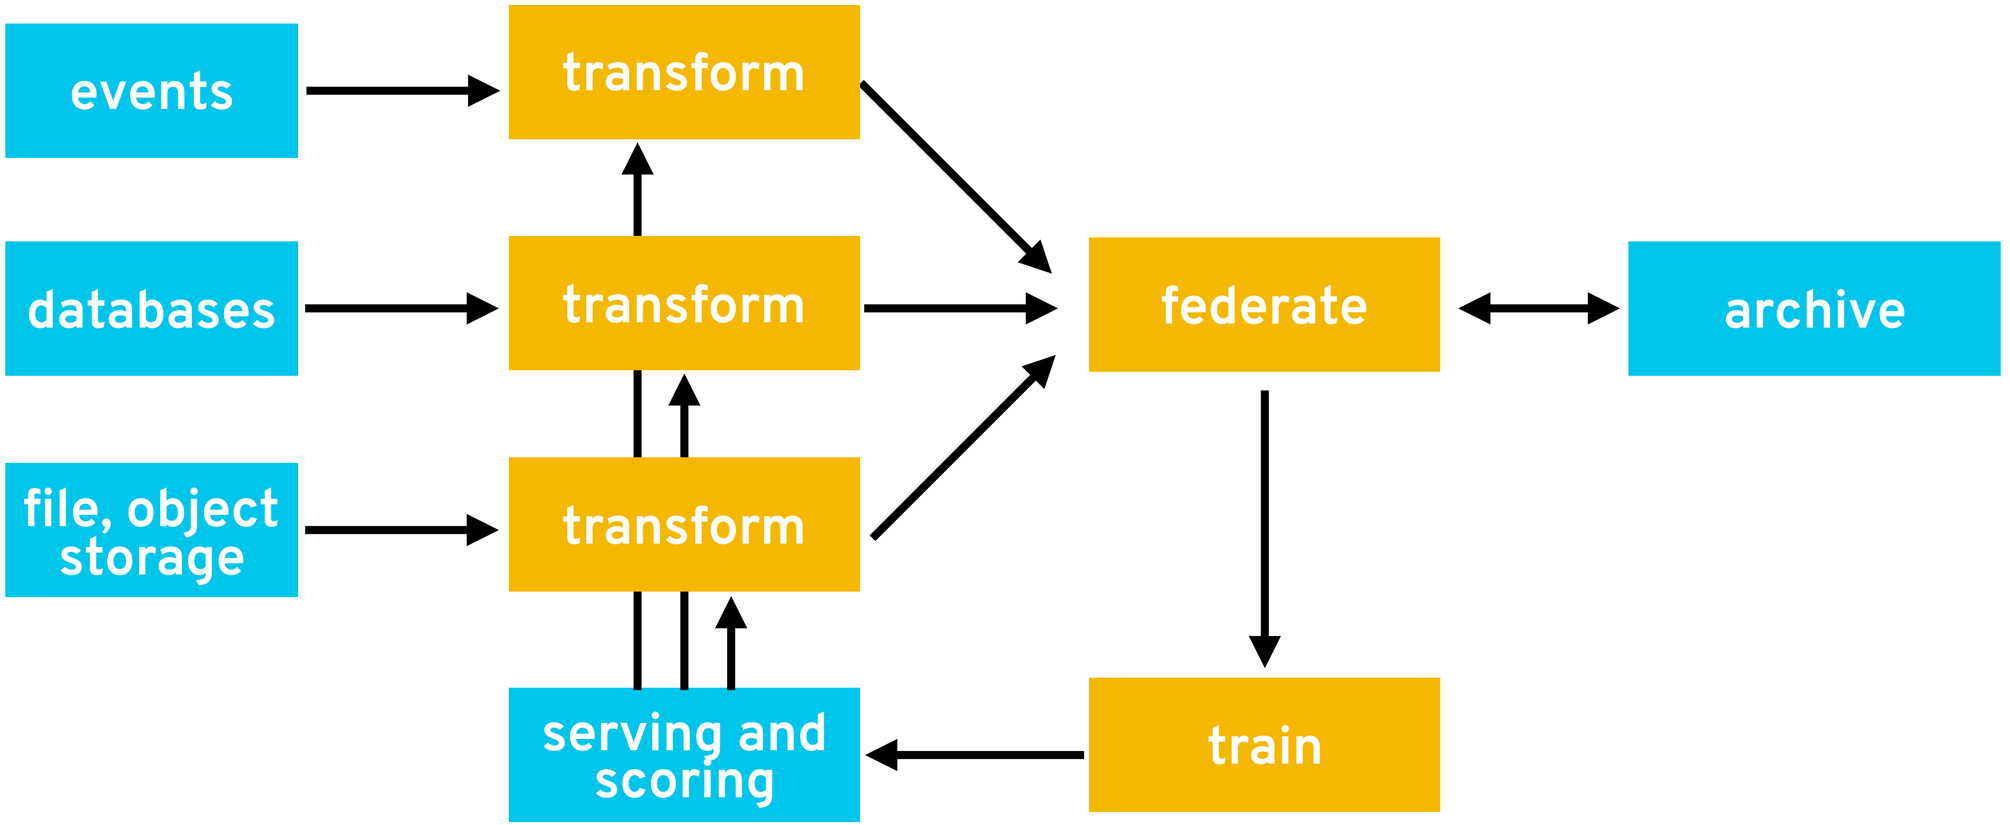
events
databases
file, object storage
transform
transform
transform
serving and scoring
federate
train
archive

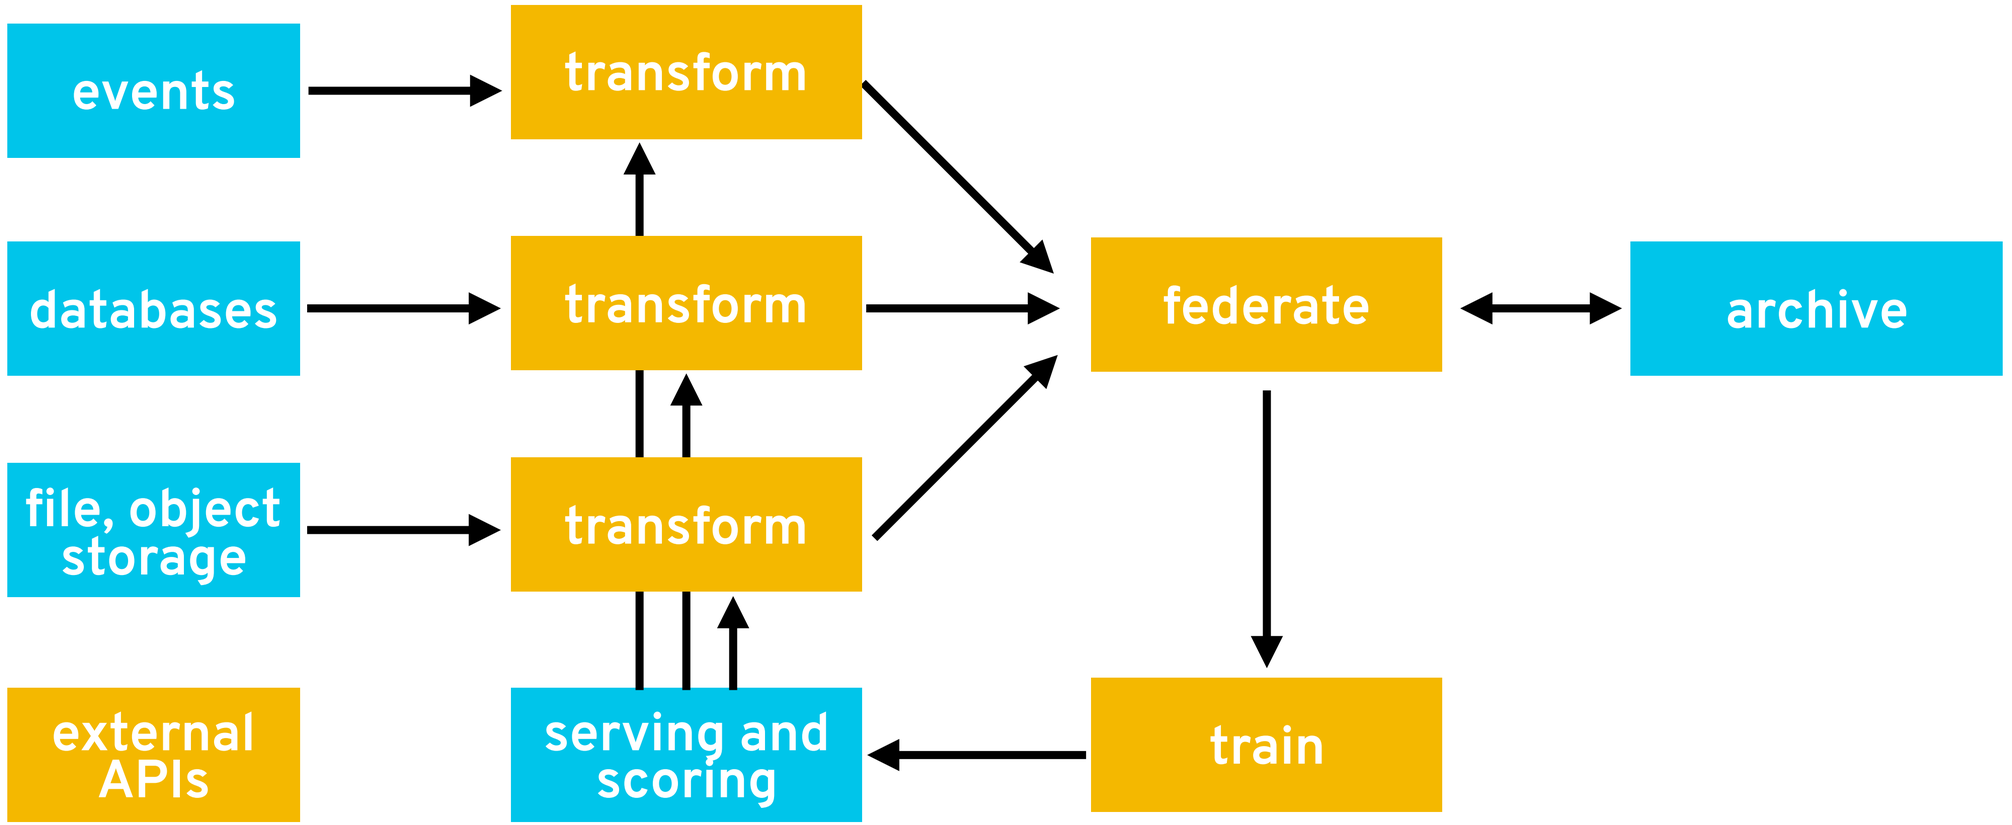
events
databases
file, object storage
external APIs
transform
transform
transform
serving and scoring
federate
train
archive

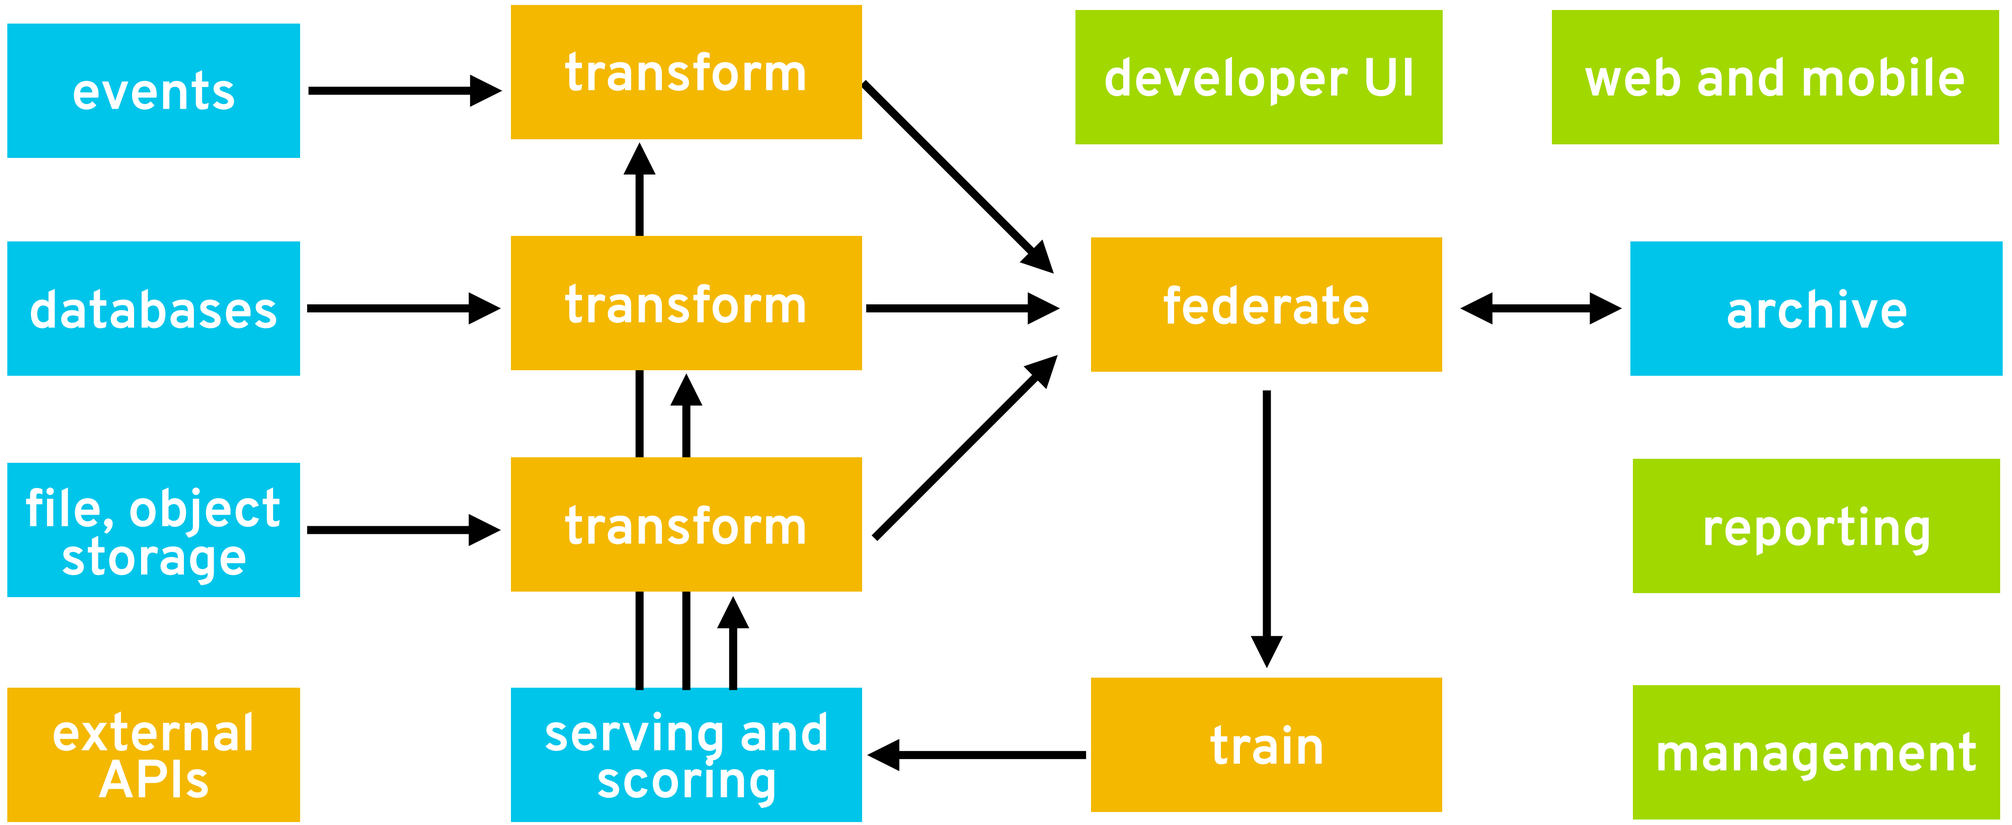
events
databases
file, object storage
external APIs
transform
transform
transform
serving and scoring
developer UI
federate
train
web and mobile
archive
reporting
management

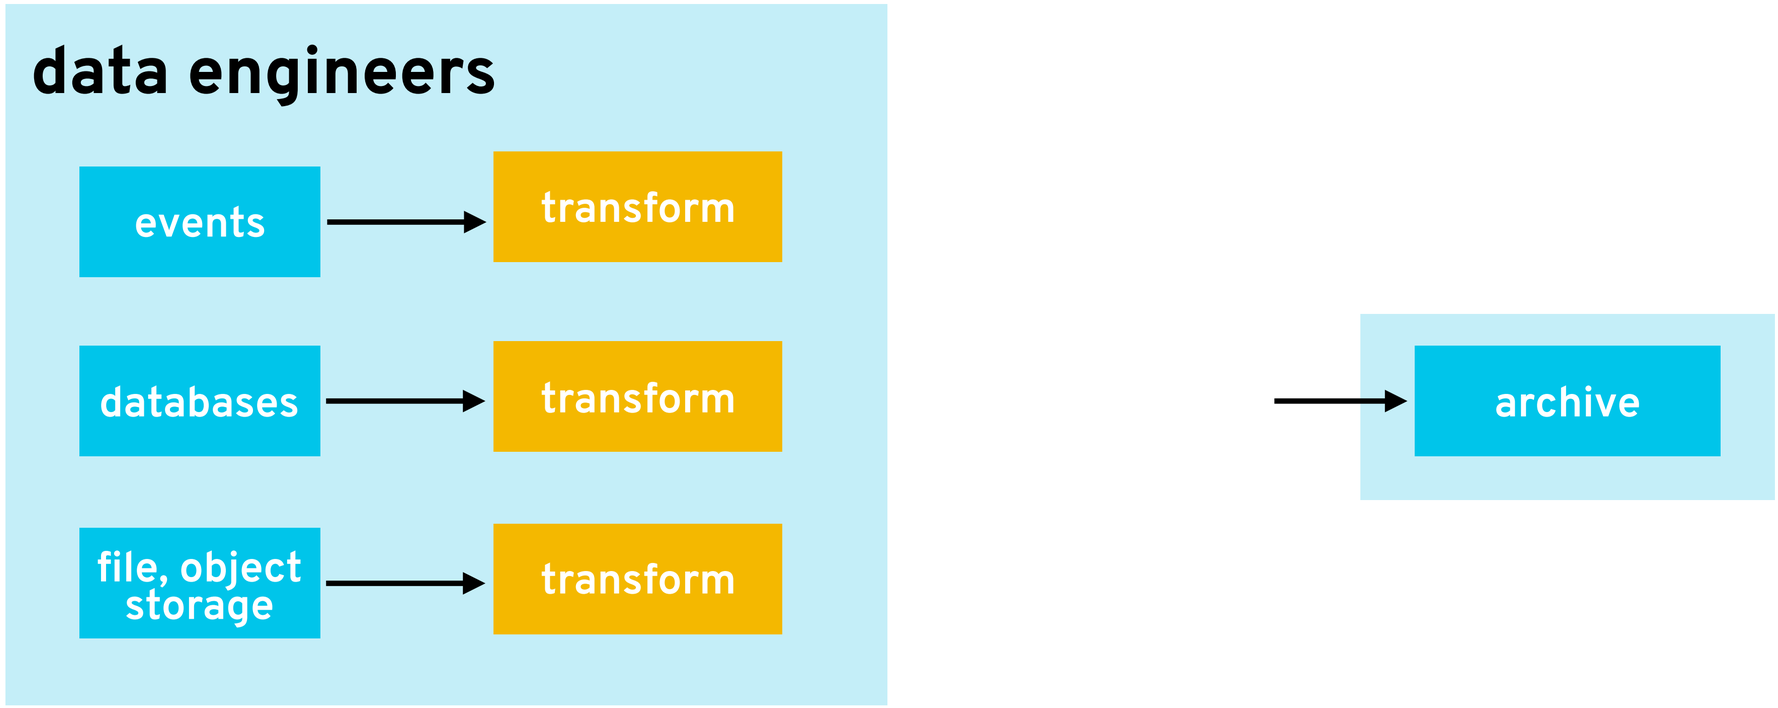
data engineers
events
transform
databases
transform
file, object storage
transform
archive

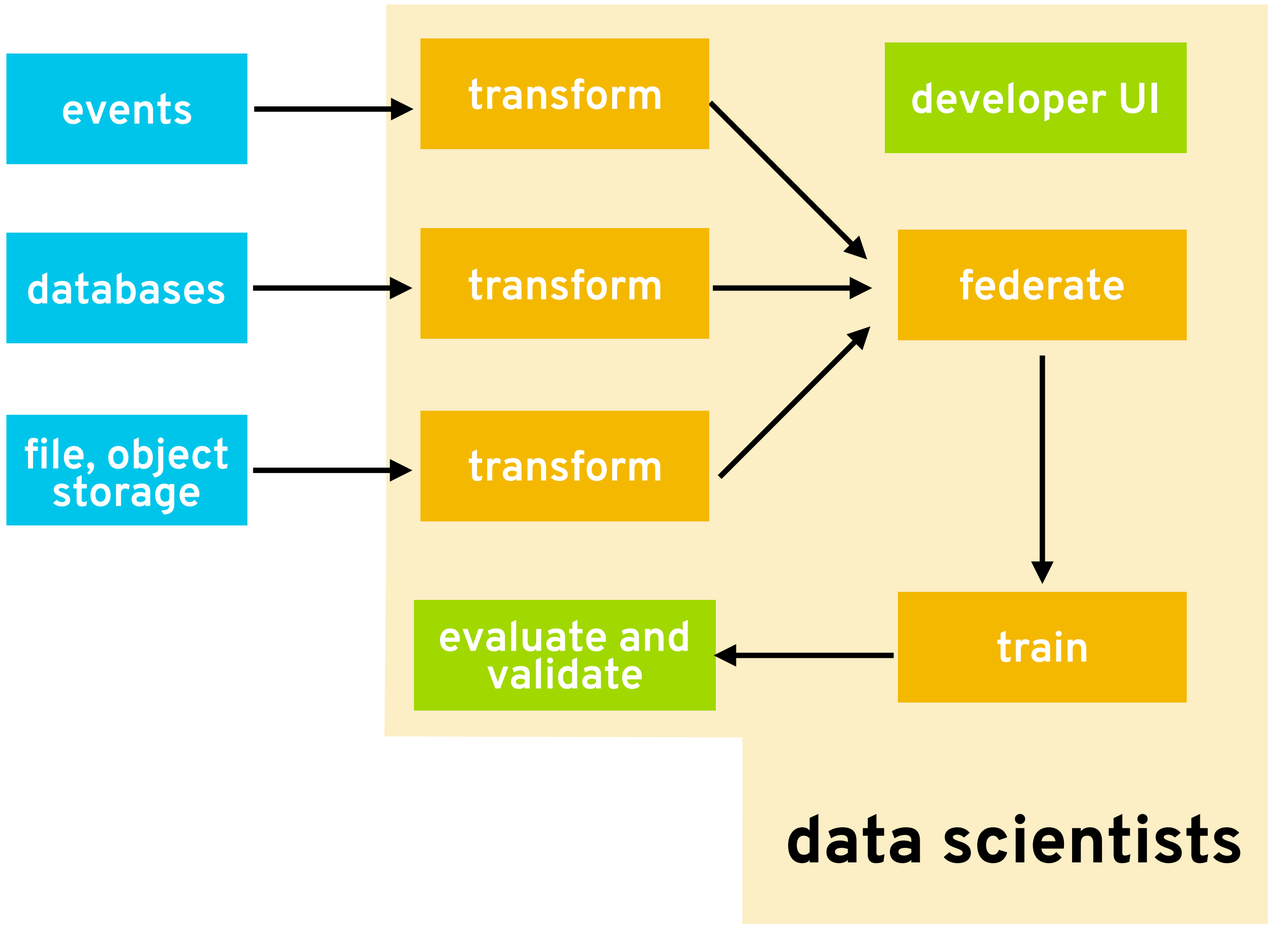
events
databases
file, object storage
transform
transform
transform
developer UI
federate
train
evaluate and validate
data scientists

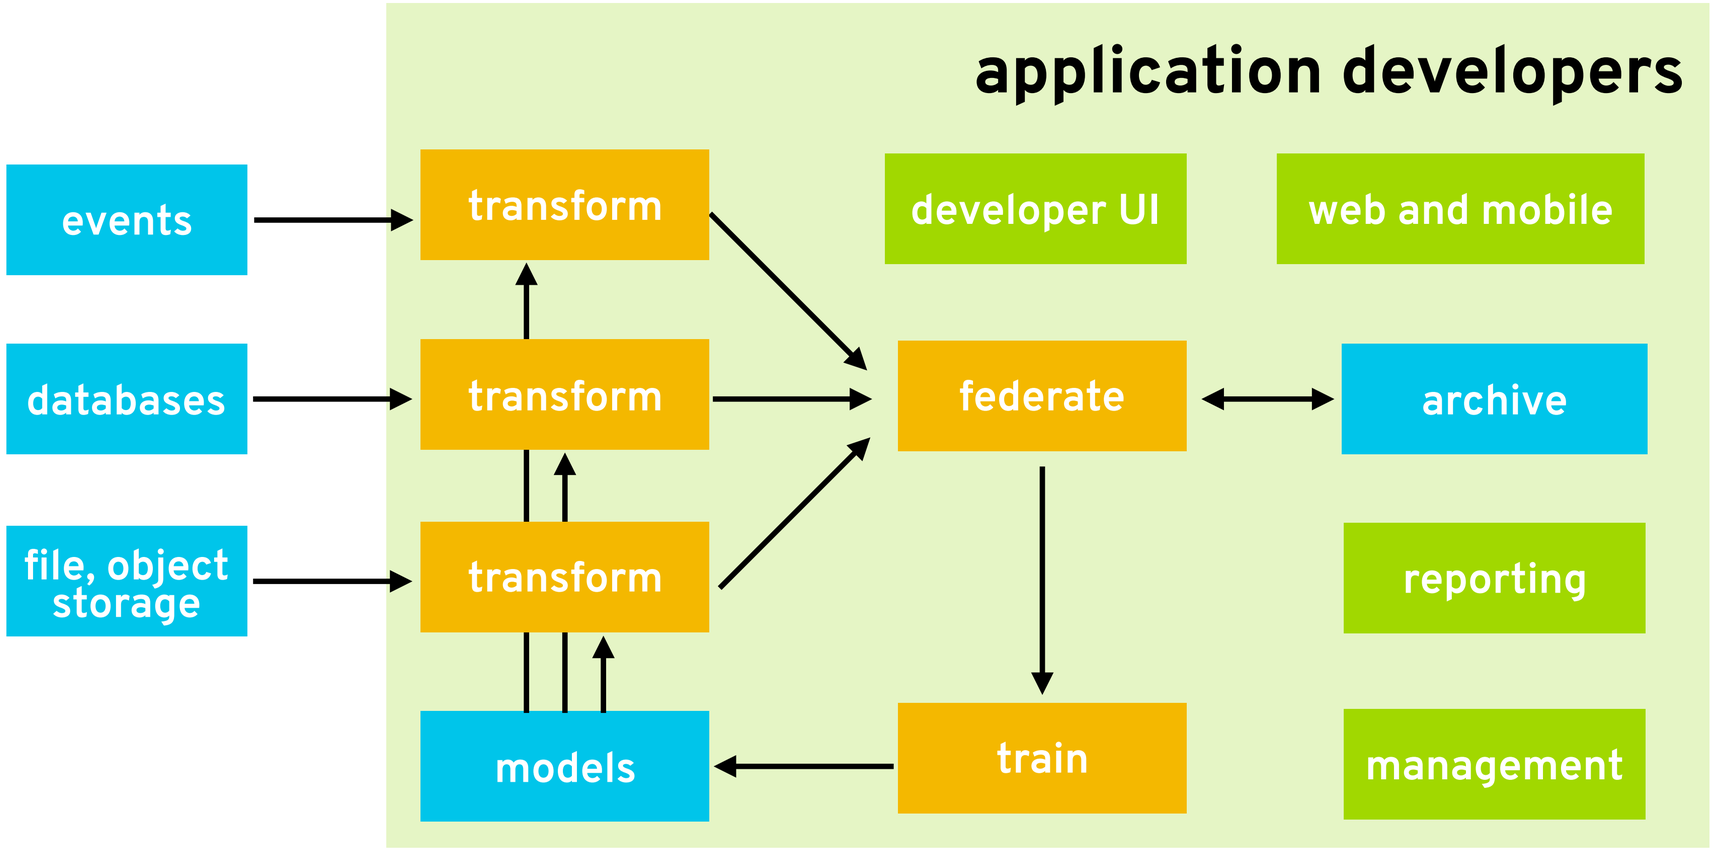
application developers
events
databases
file, object storage
transform
transform
transform
models
developer UI
web and mobile
federate
archive
reporting
train
management

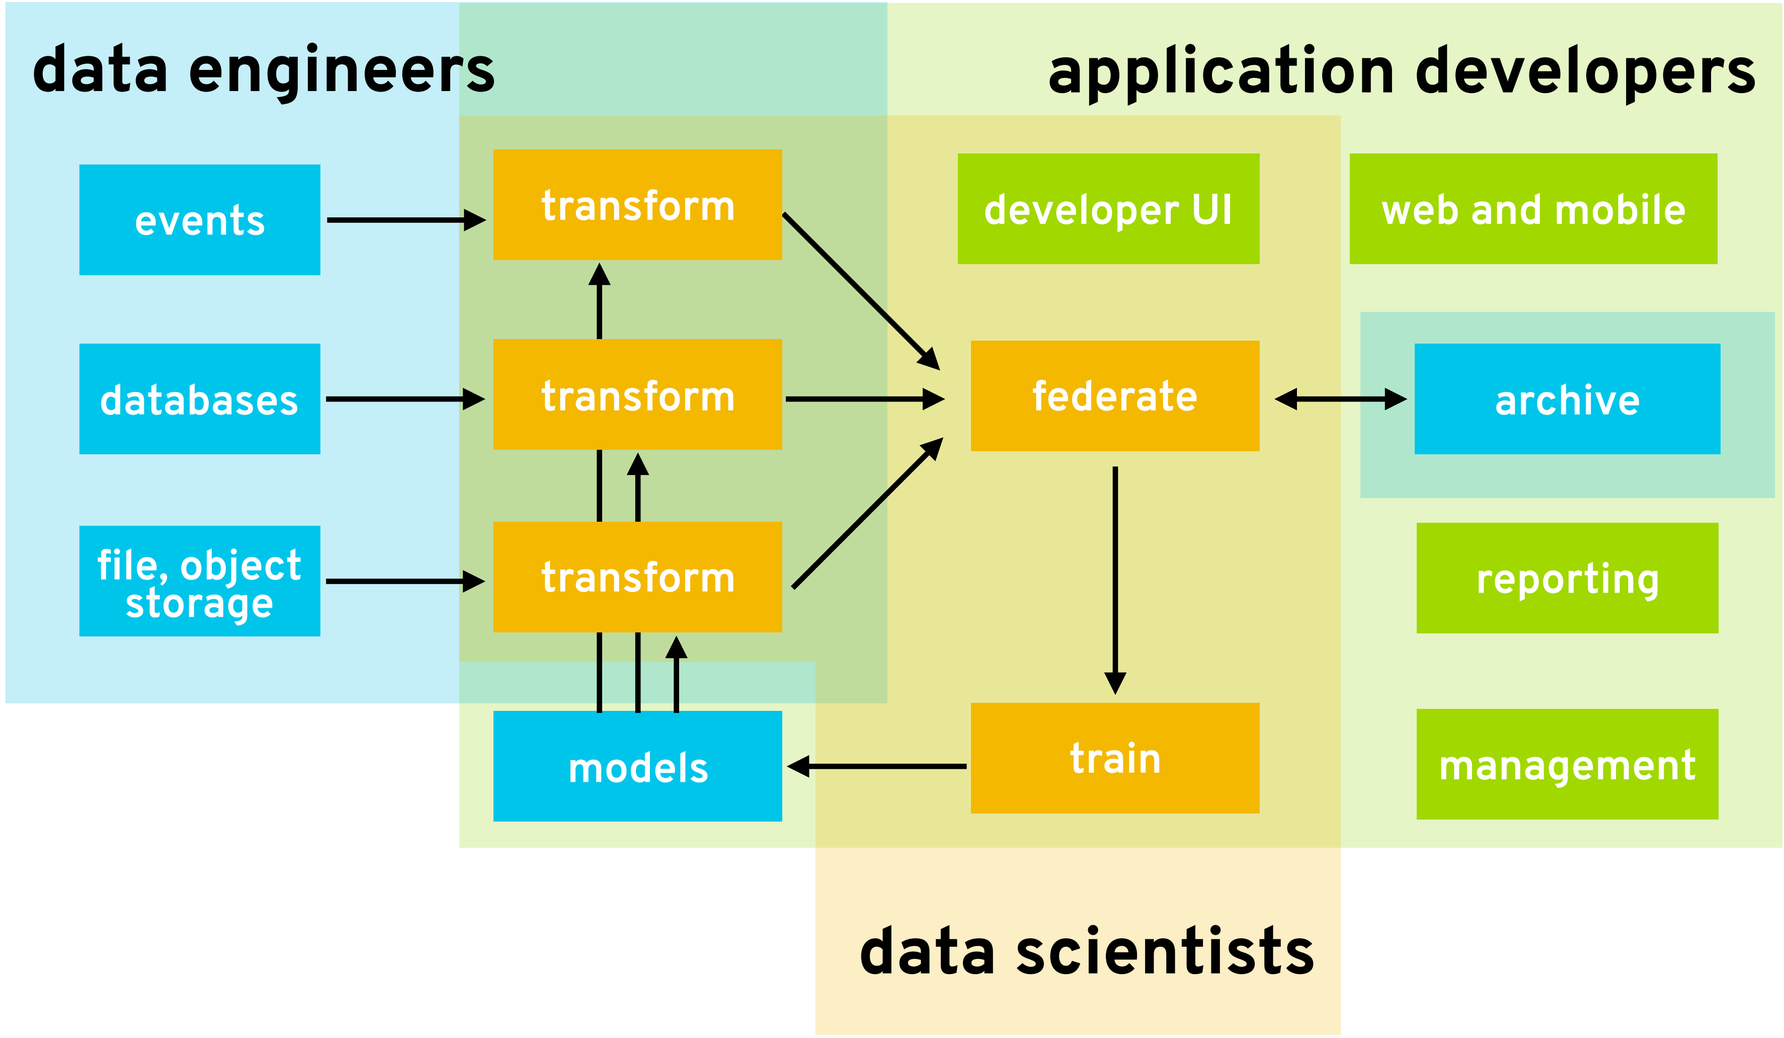
data engineers
application developers
events
databases
file, object storage
transform
transform
transform
models
developer UI
federate
train
web and mobile
archive
reporting
management
data scientists

Containers: great for application developers and data scientists

What is a container?

%

```
% pip install numpy
```

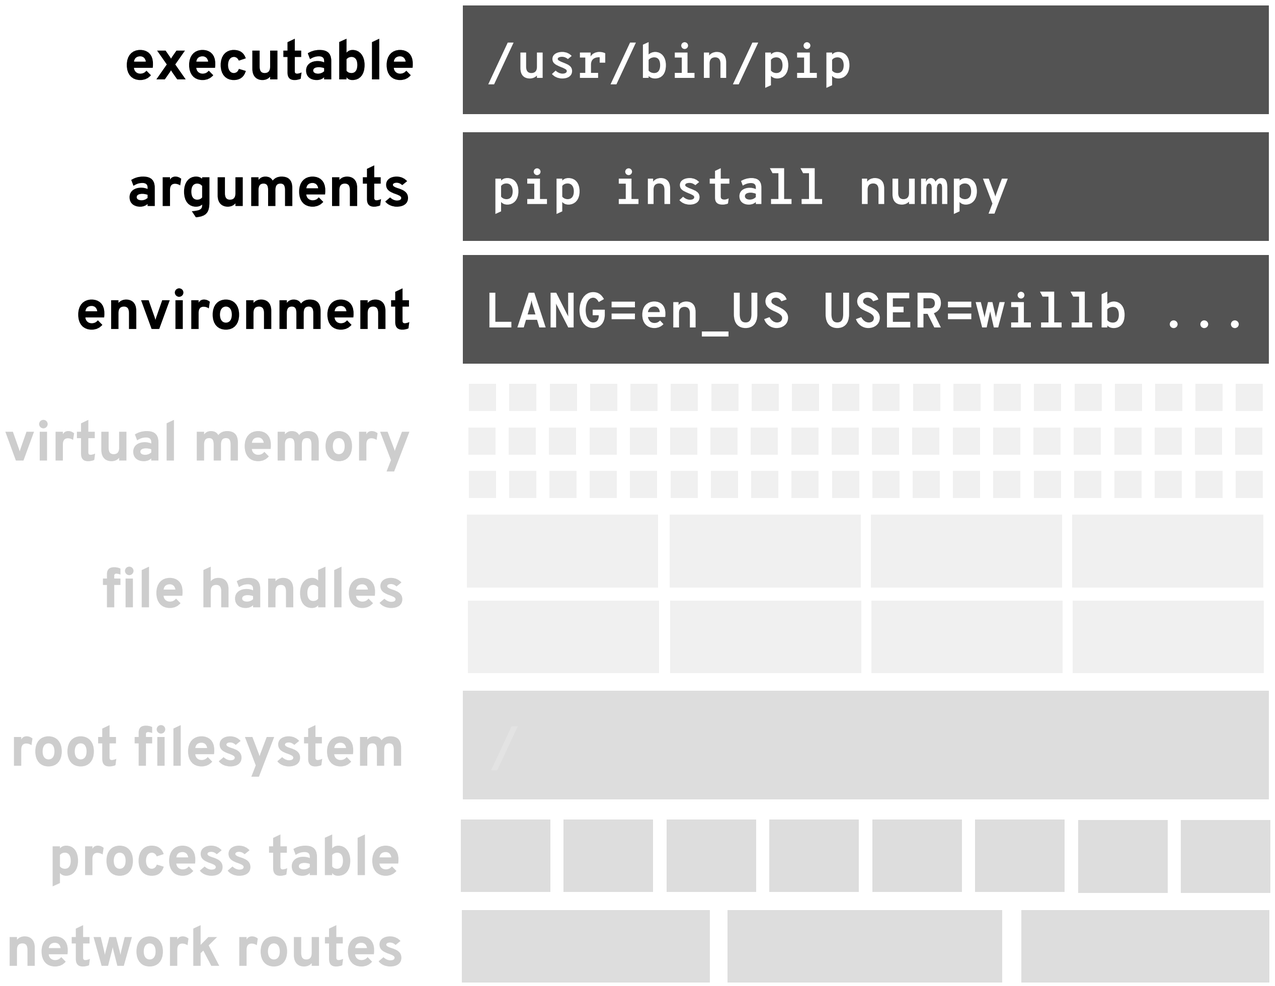
executable
/usr/bin/pip
arguments
pip install numpy
environment
LANG=en_US USER=willb ...
virtual memory
file handles
root filesystem
/
process table
network routes

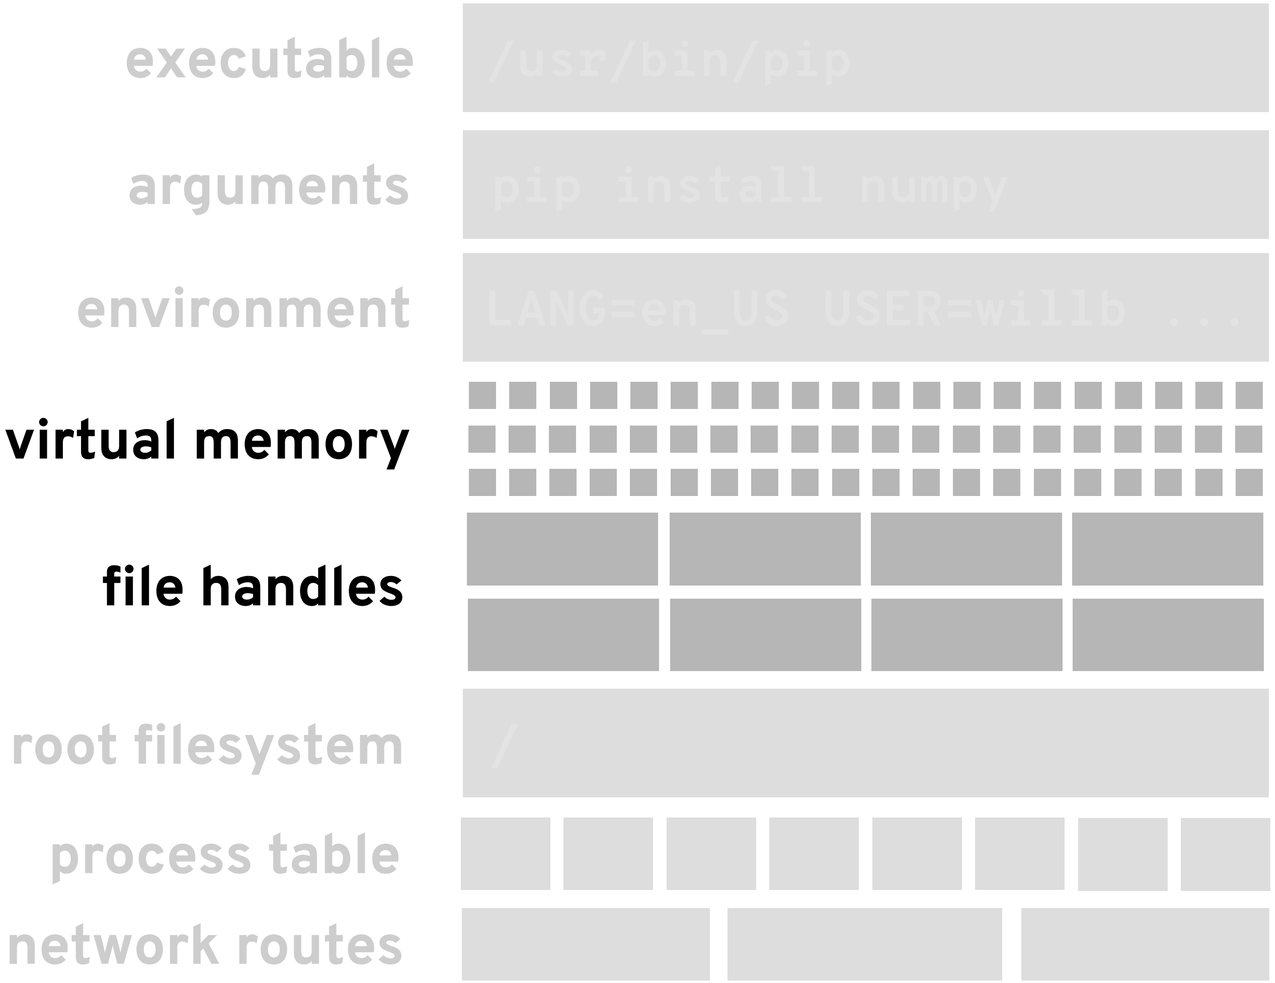
executable
/usr/bin/pip
arguments
pip install numpy
environment
LANG=en_US USER=willb ...
virtual memory
file handles
root filesystem
/
process table
network routes

executable /usr/bin/pip

arguments pip install numpy

root filesystem /

process table

network routes

executable /usr/bin/pip

arguments pip install numpy

root filesystem /

process table

network routes

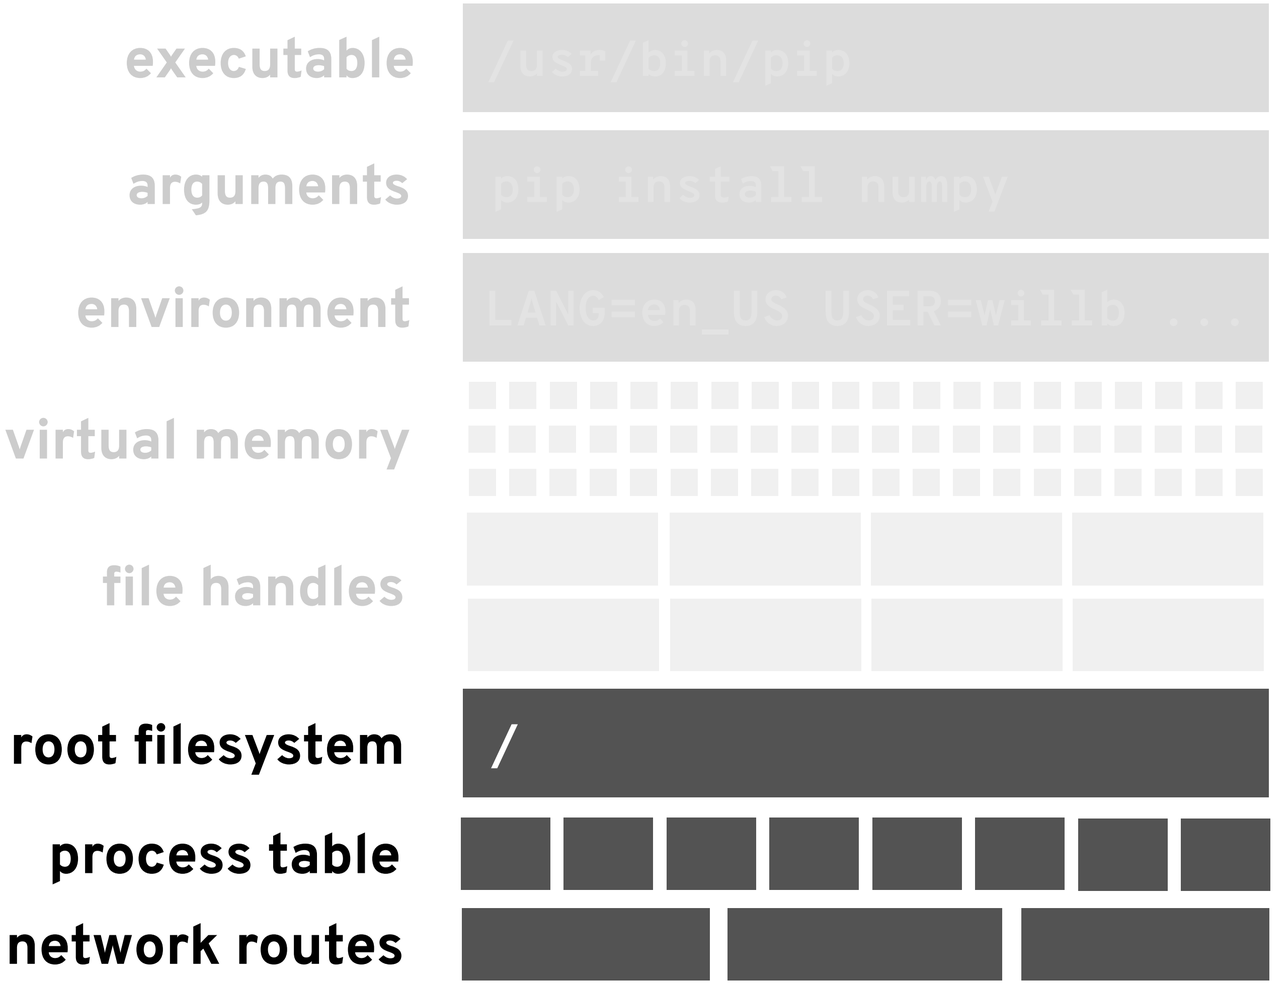
executable
/usr/bin/pip
arguments
pip install numpy
environment
LANG=en_US USER=willb ...
virtual memory
file handles
root filesystem
/
process table
network routes

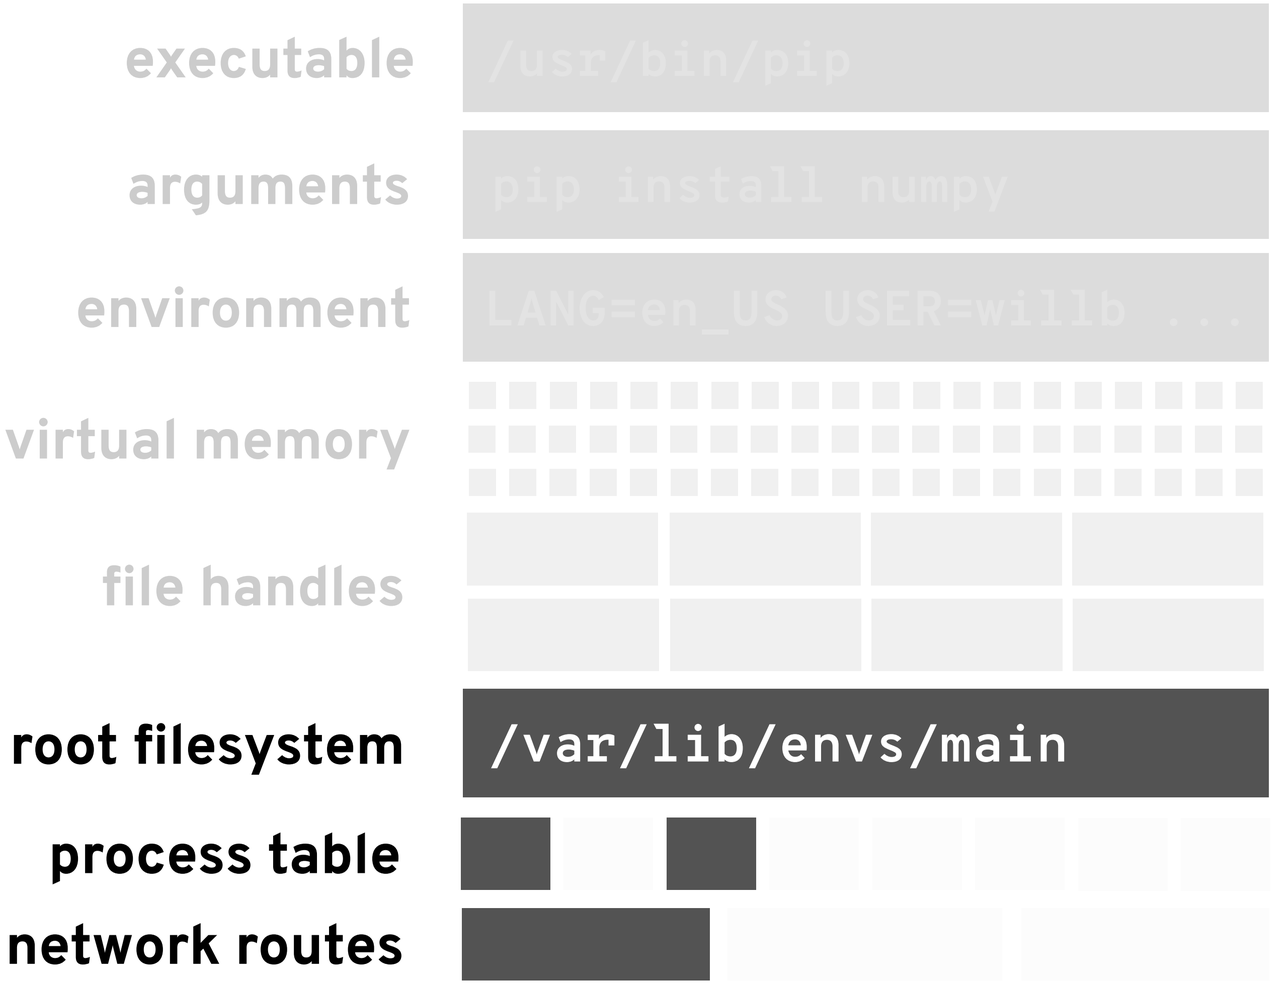
executable
/usr/bin/pip
arguments
pip install numpy
environment
LANG=en_US USER=willb ...
virtual memory
file handles
root filesystem
/var/lib/envs/main
process table
network routes

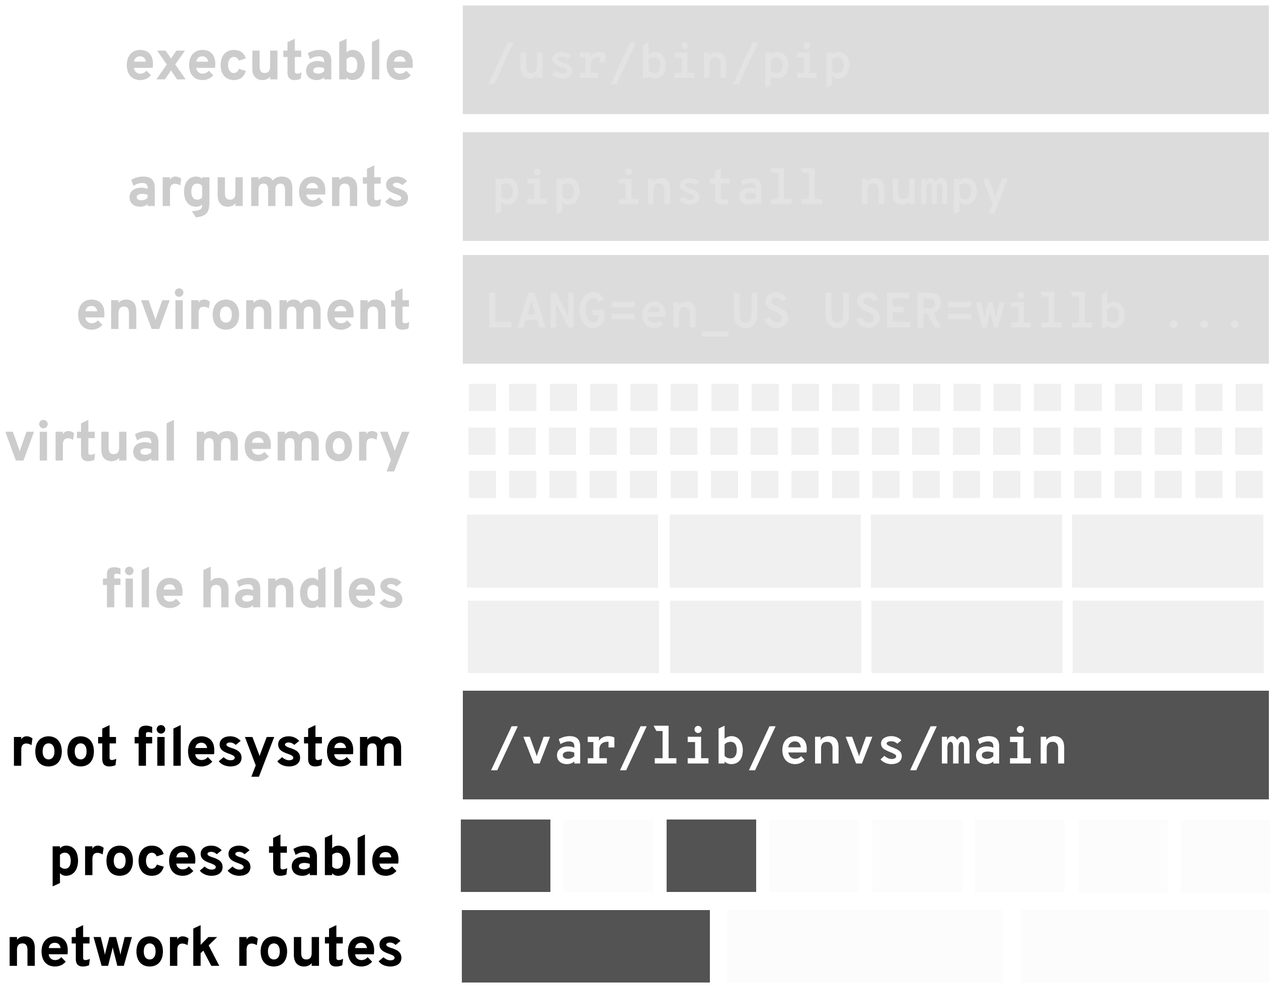
executable
/usr/bin/pip
arguments
pip install numpy
environment
LANG=en_US USER=willb ...
virtual memory
file handles
root filesystem
/var/lib/envs/main
process table
network routes


SPEED
LIMIT
55

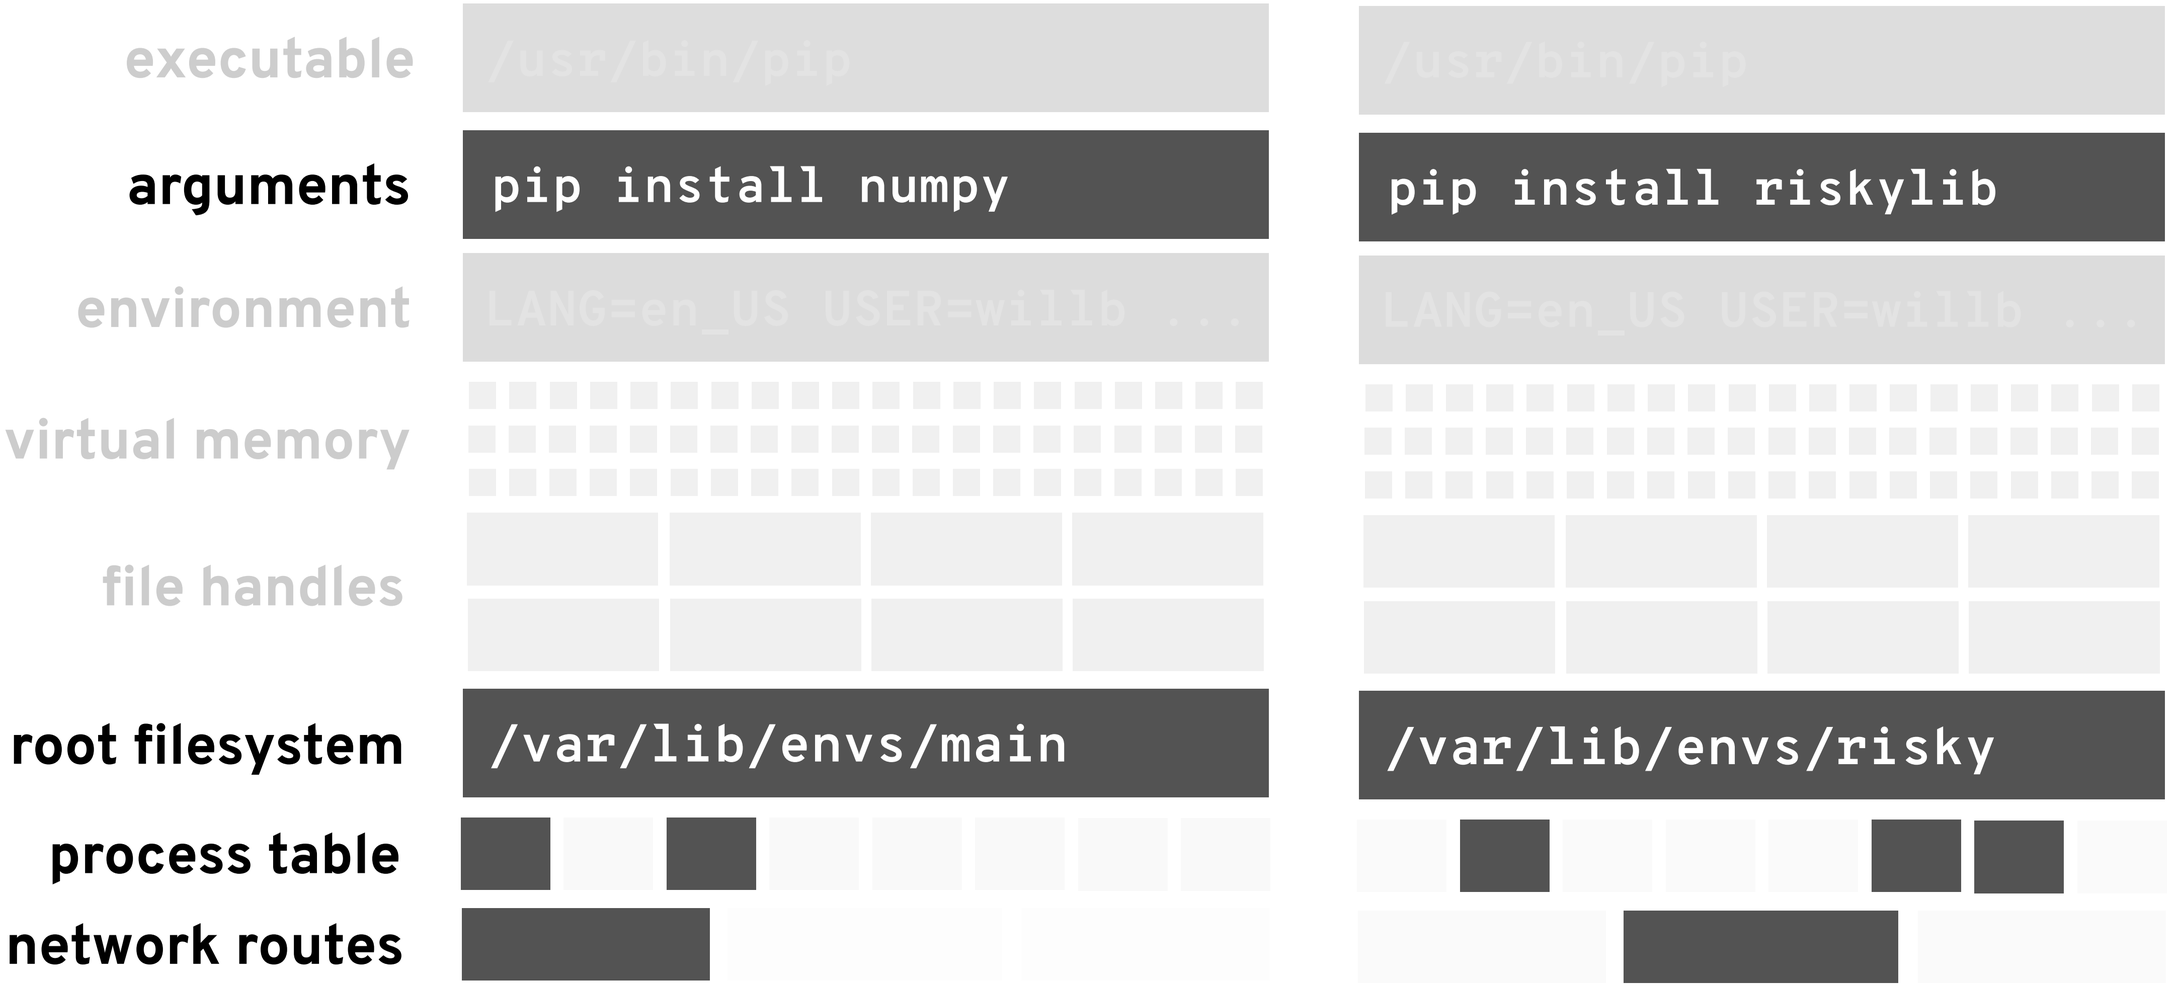
executable
/usr/bin/pip
/usr/bin/pip
arguments
pip install numpy
pip install riskylib
environment
LANG=en_US USER=willb ...
LANG=en_US USER=willb ...
virtual memory
file handles
root filesystem
/var/lib/envs/main
/var/lib/envs/risky
process table
network routes

Immutable images

base image

Immutable images

Immutable images

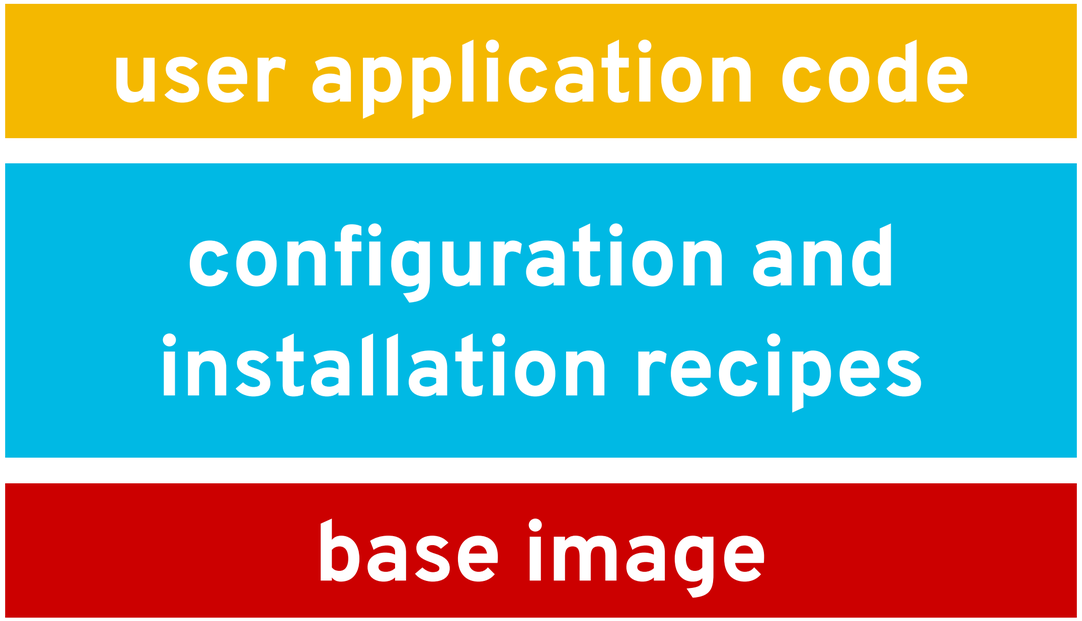

Immutable images

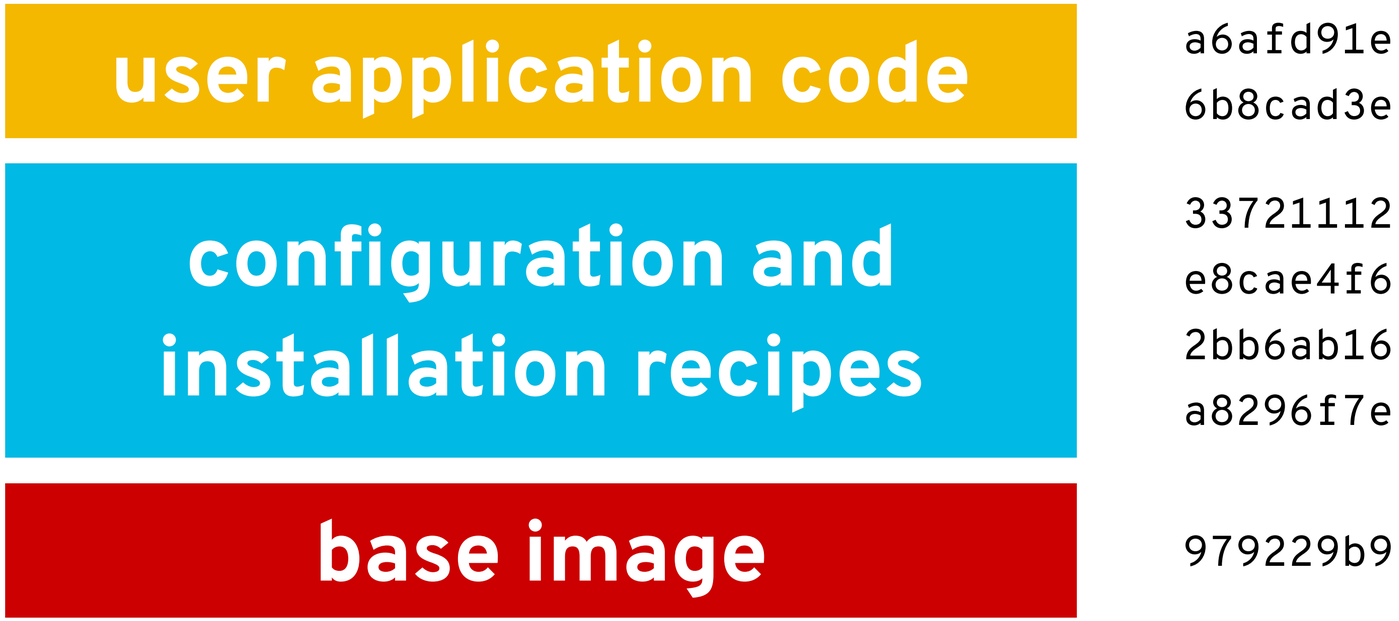

Stateless microservices

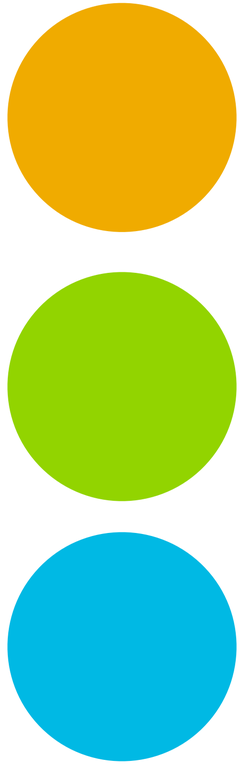

Stateless microservices

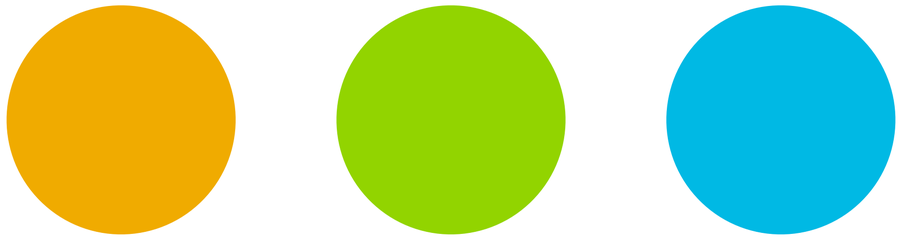

Stateless microservices

Stateless microservices

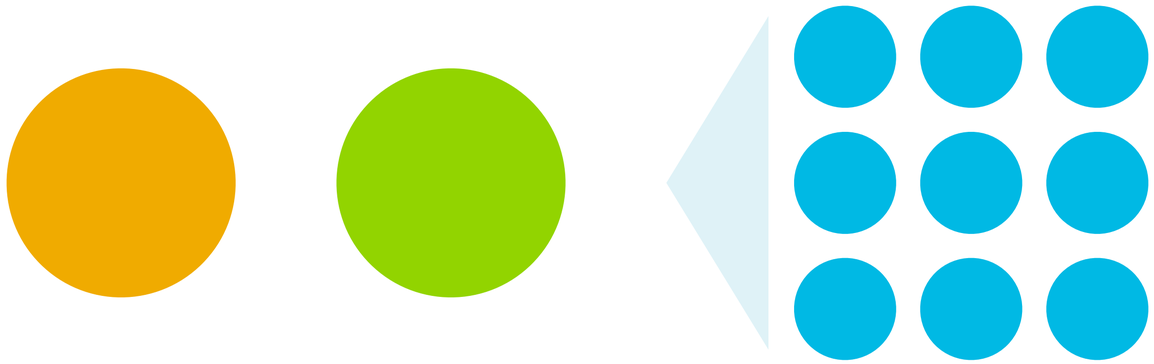

Declarative app configuration

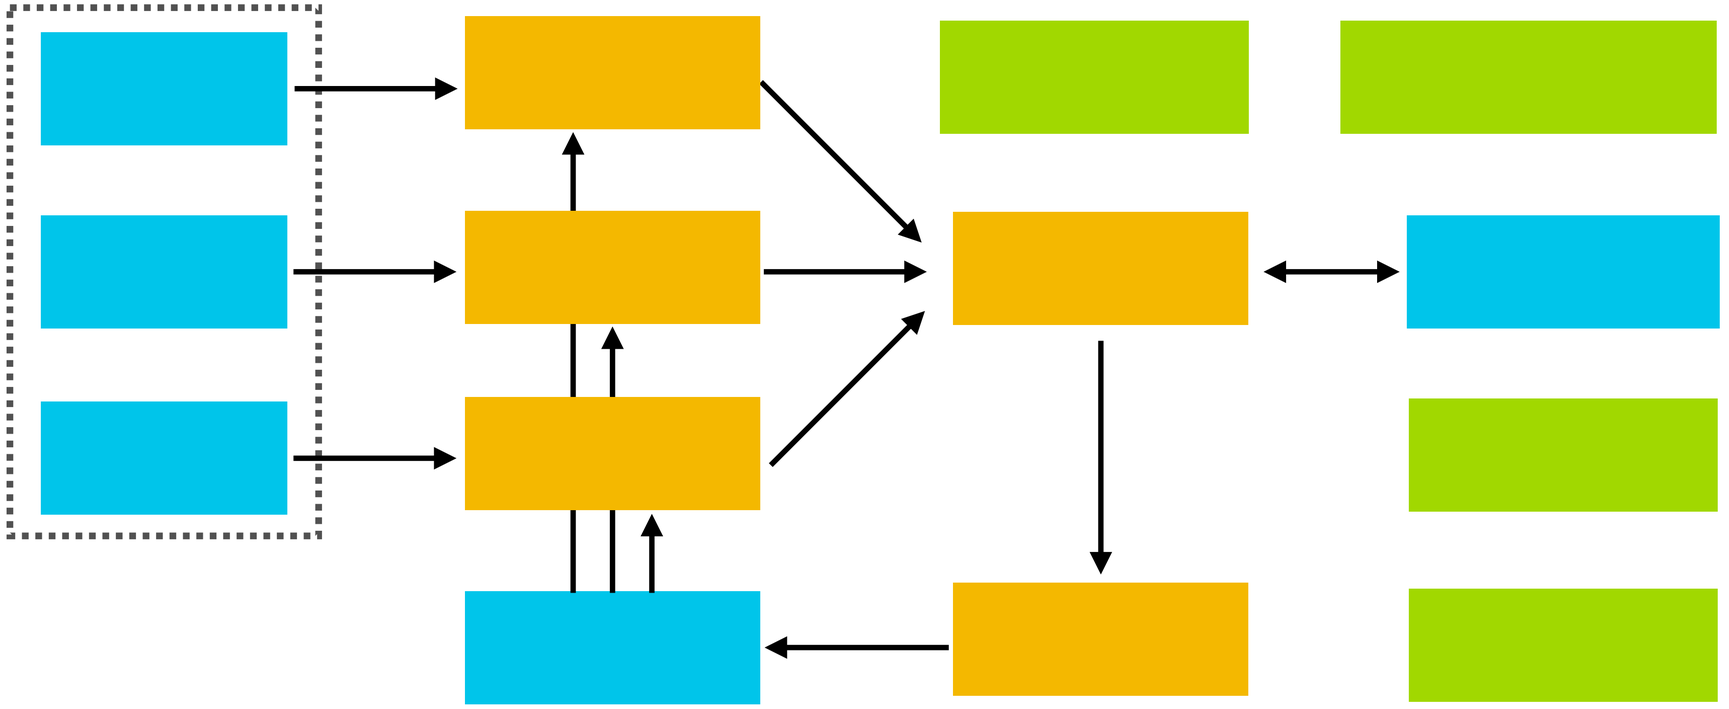

Integration and deployment

Integration and deployment

Integration and deployment

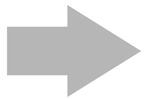

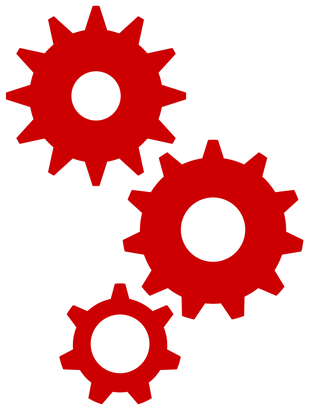

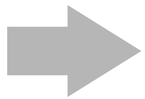

application code

configuration and installation recipes

base image

Integration and deployment

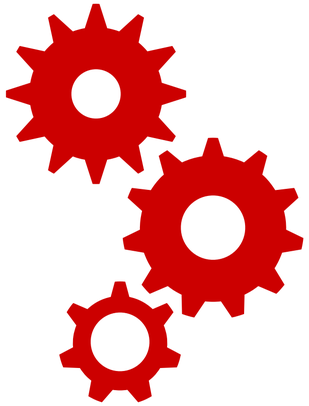

application code

configuration and installation recipes

base image

Integration and deployment

application code

configuration and installation recipes

base image

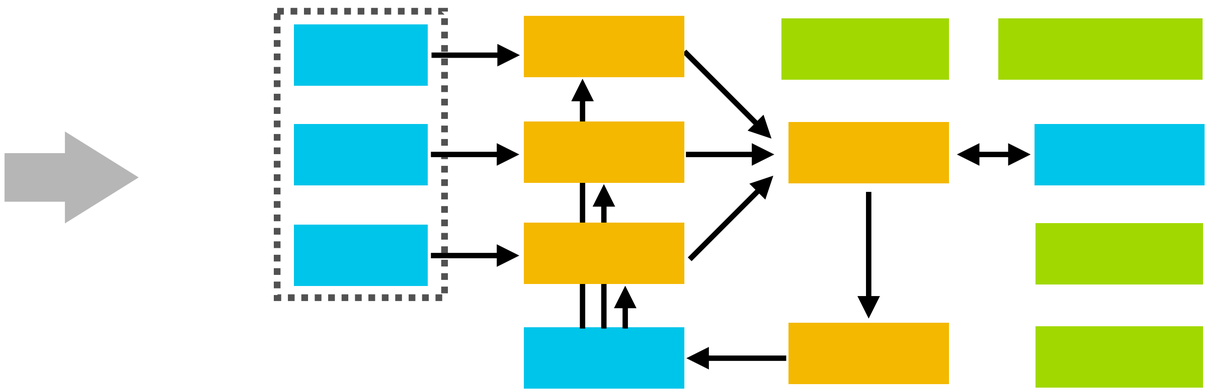

Flexible service routing

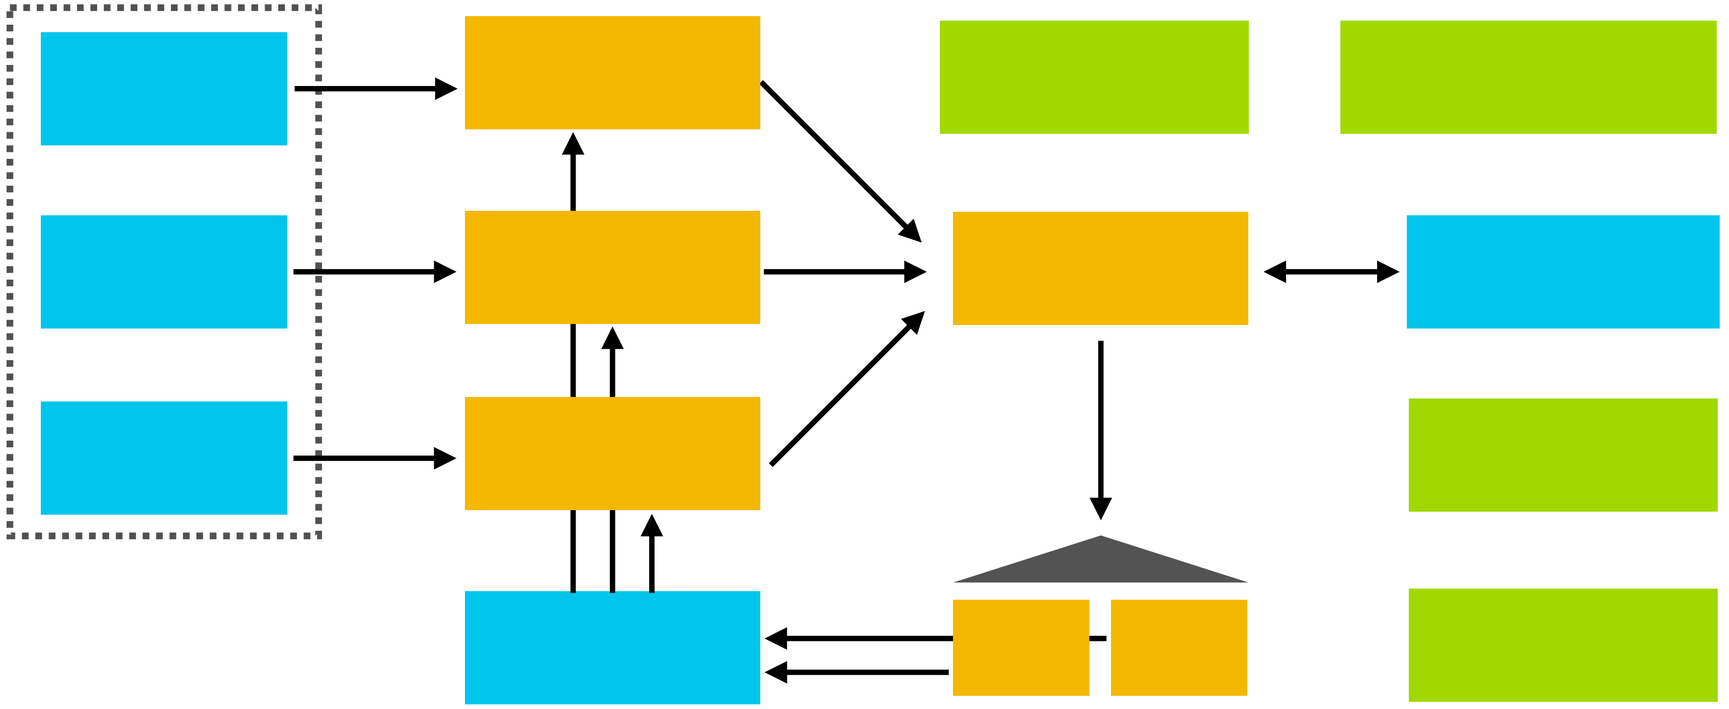

Flexible service routing

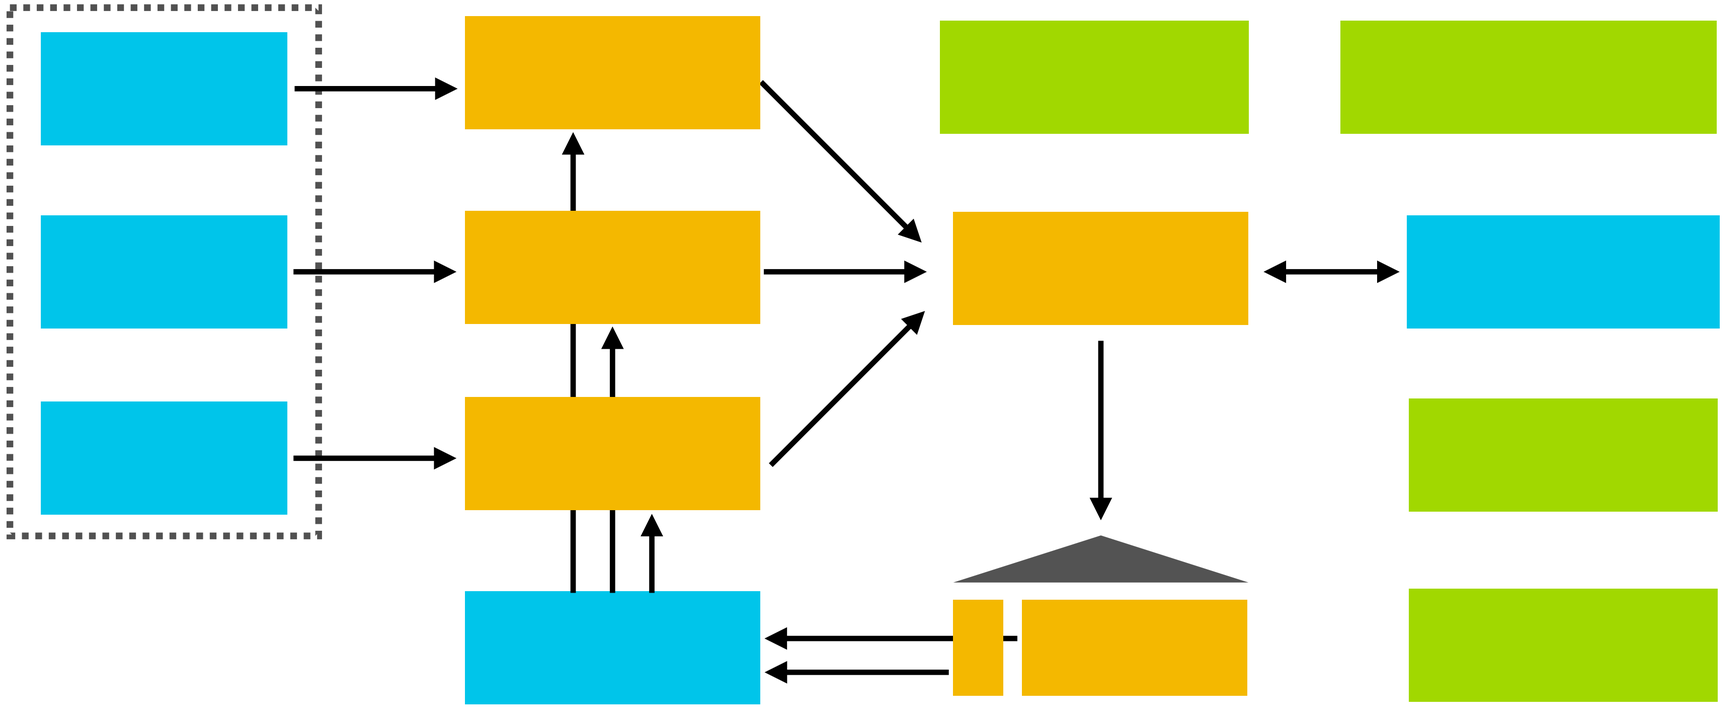

How can we build *intelligent applications* in containers?

Apache Spark

Apache Spark

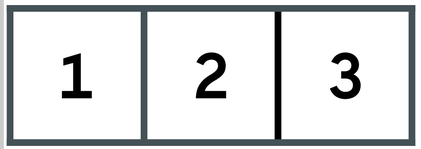

Apache Spark

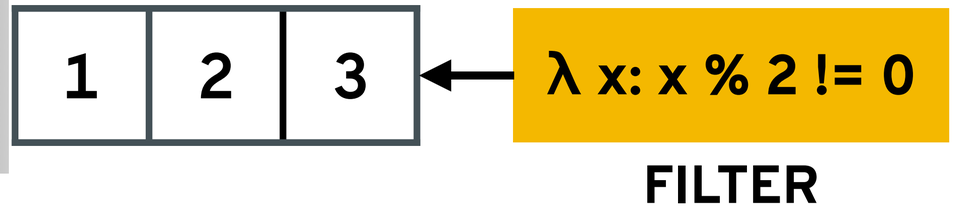

Apache Spark

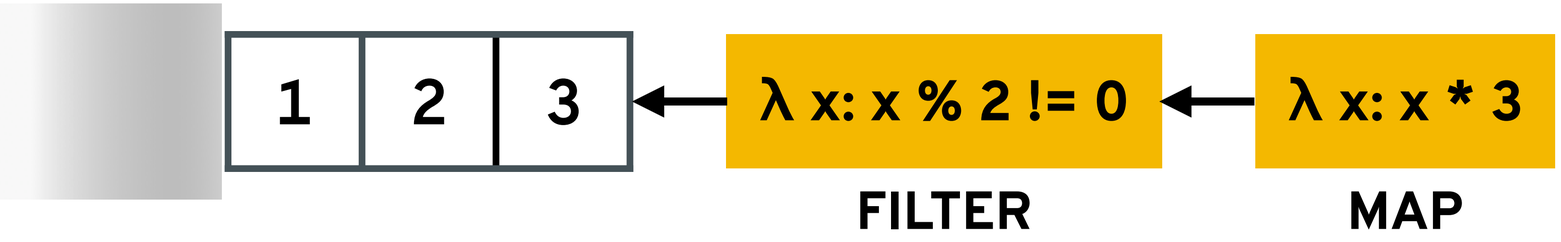

Apache Spark

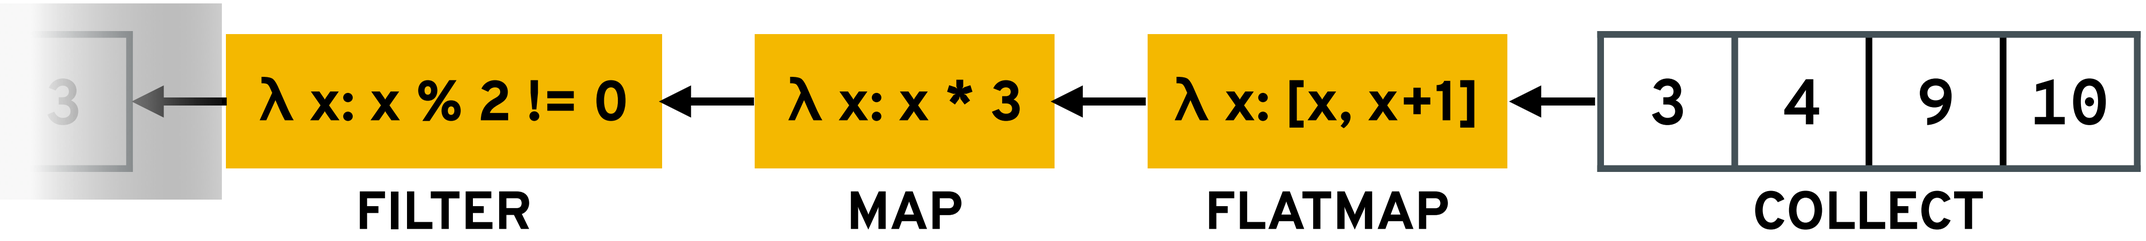
3
λ x: x % 2 != 0
λ x: x * 3
λ x: [x, x+1]
3
4
9
10
FILTER
MAP
FLATMAP
COLLECT

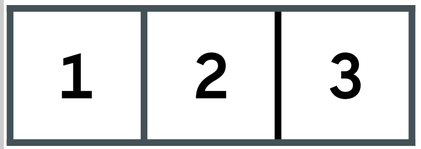
1
2
3

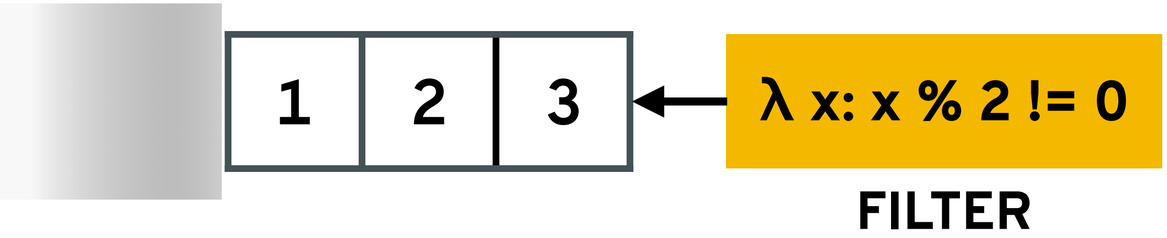
1
2
3
λ x: x % 2 != 0
FILTER

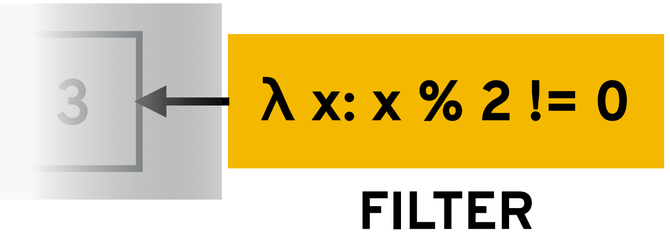
3
λ x: x % 2 != 0
FILTER

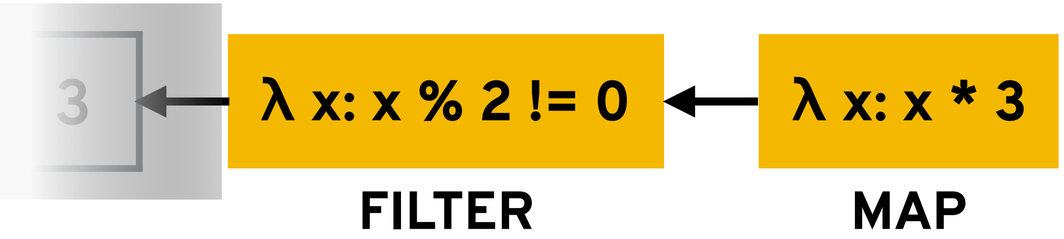
3
λ x: x % 2 != 0
λ x: x * 3
FILTER
MAP

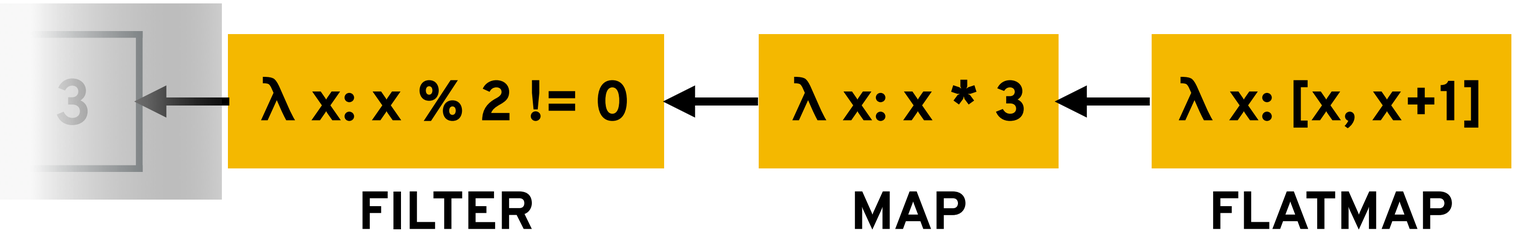
3
λ x: x % 2 != 0
λ x: x * 3
λ x: [x, x+1]
FILTER
MAP
FLATMAP

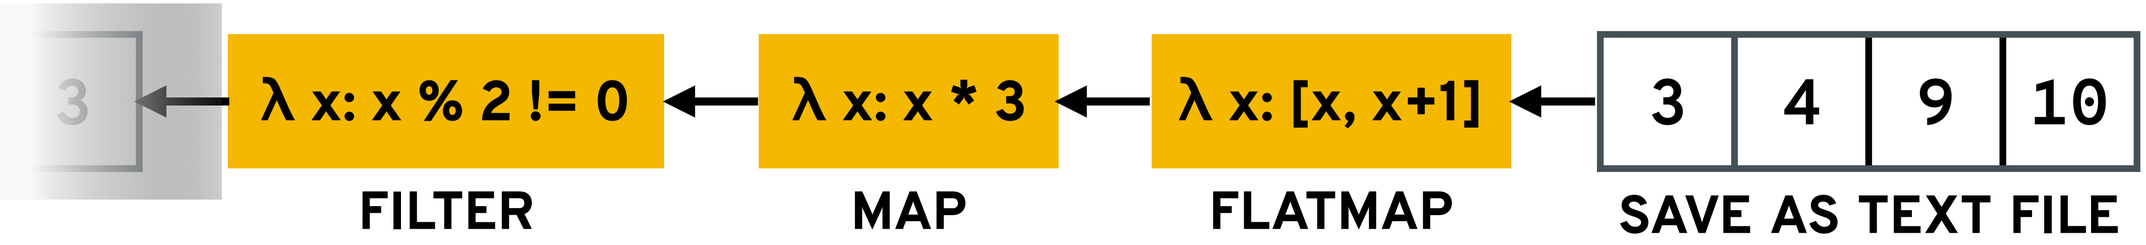
3
λ x: x % 2 != 0
FILTER
λ x: x * 3
MAP
λ x: [x, x+1]
FLATMAP
3
4
9
10
SAVE AS TEXT FILE

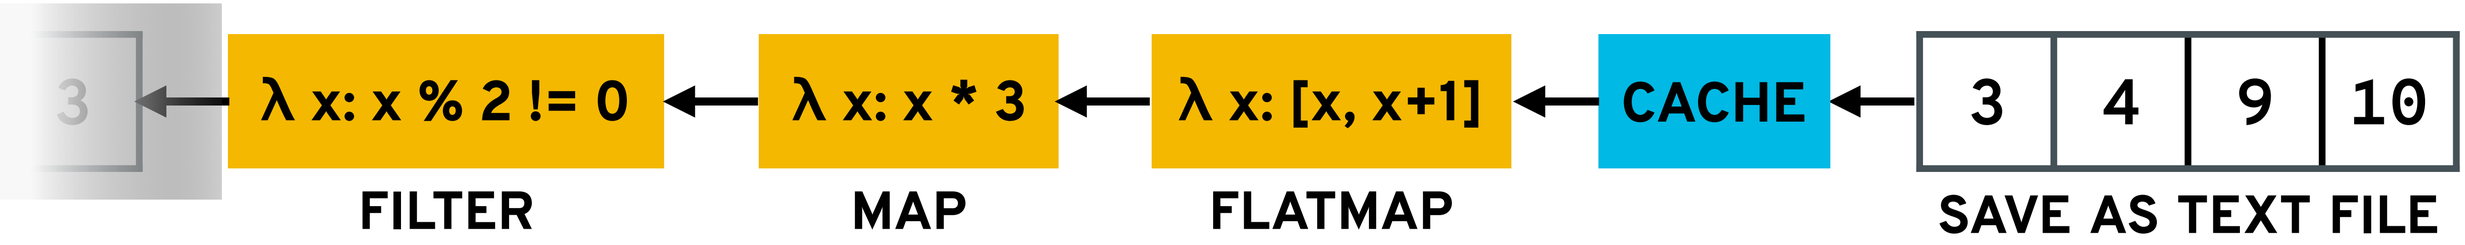
3
λ x: x % 2 != 0
FILTER
λ x: x * 3
MAP
λ x: [x, x+1]
FLATMAP
CACHE
3
4
9
10
SAVE AS TEXT FILE

Spark core

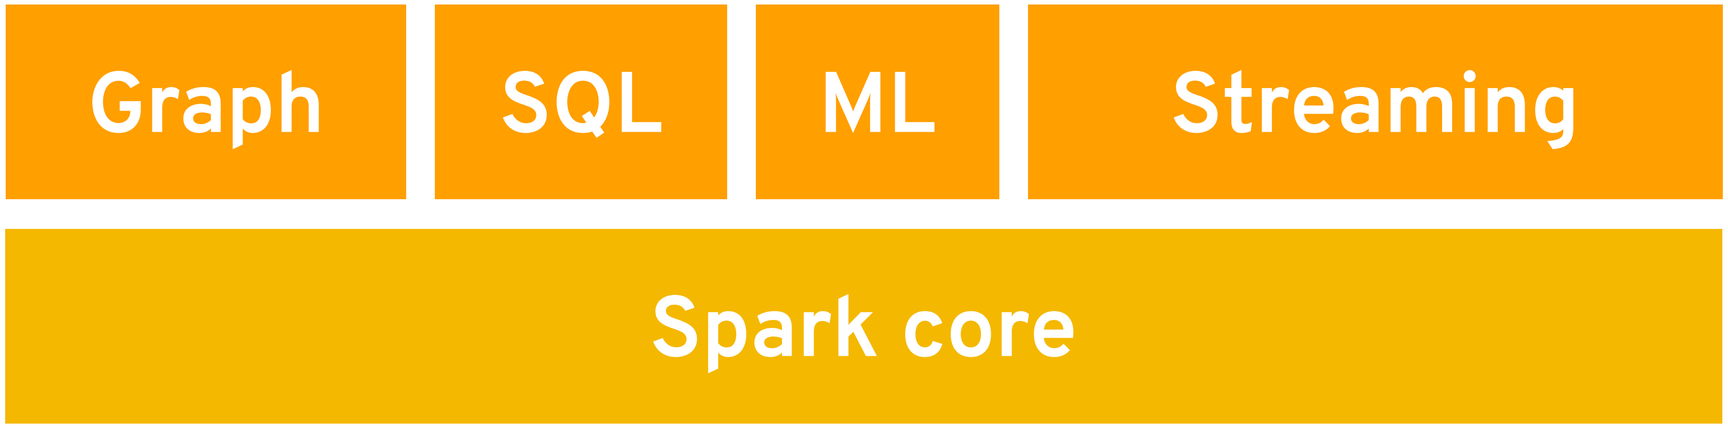
Graph
SQL
ML
Streaming
Spark core

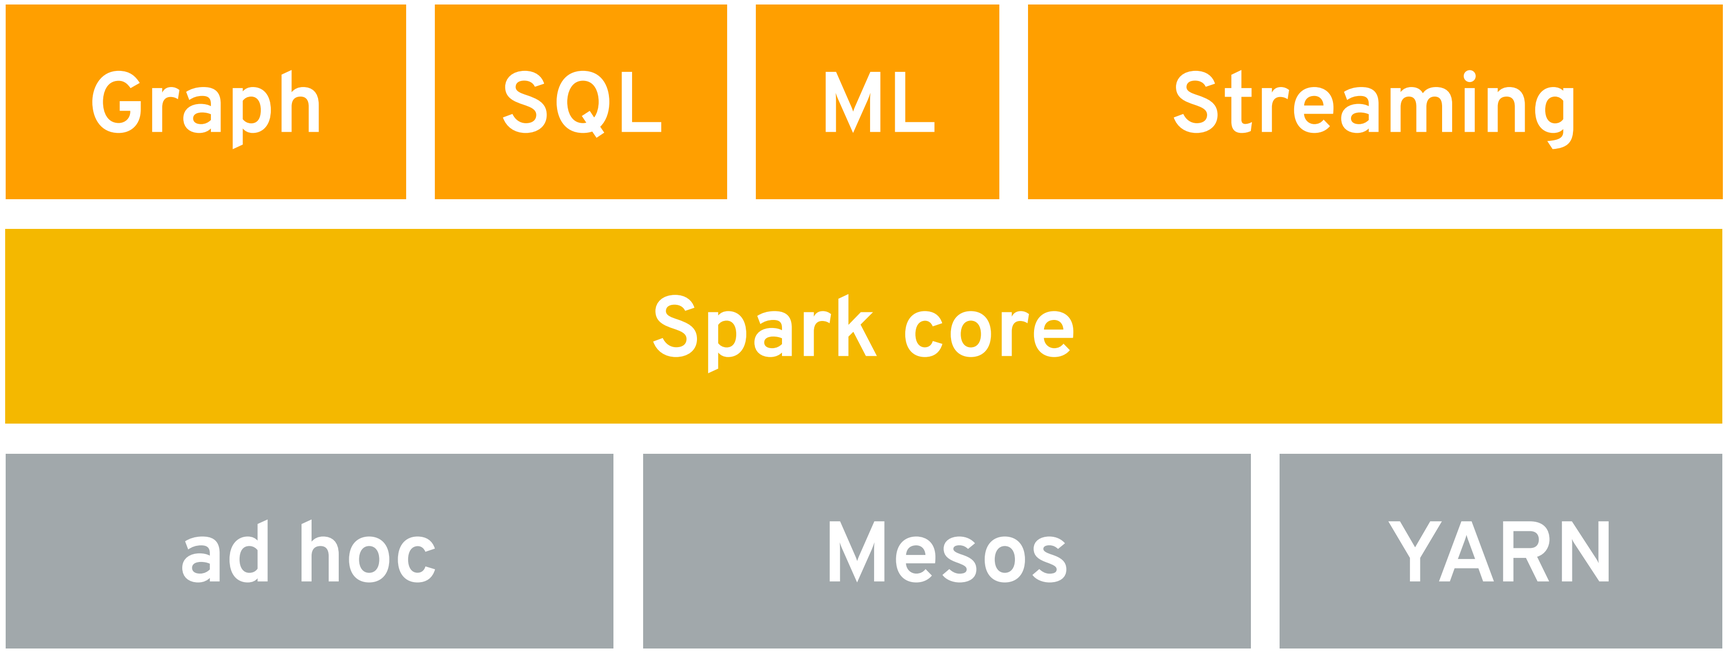
Graph
SQL
ML
Streaming
Spark core
ad hoc
Mesos
YARN

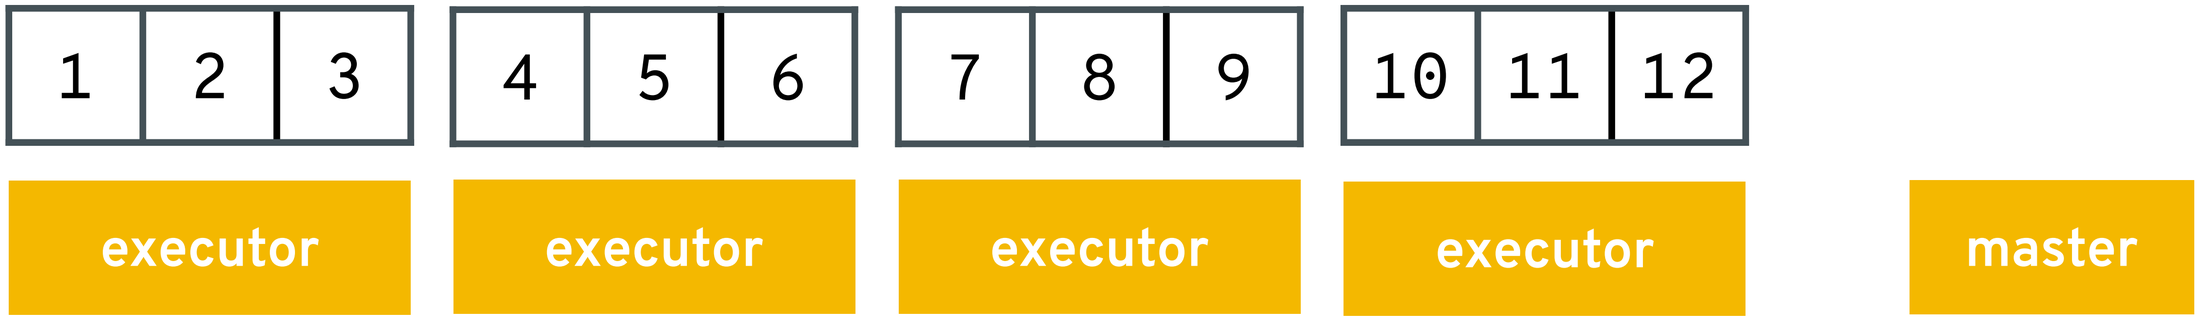
1
2
3
4
5
6
7
8
9
10
11
12
executor
executor
executor
executor
master

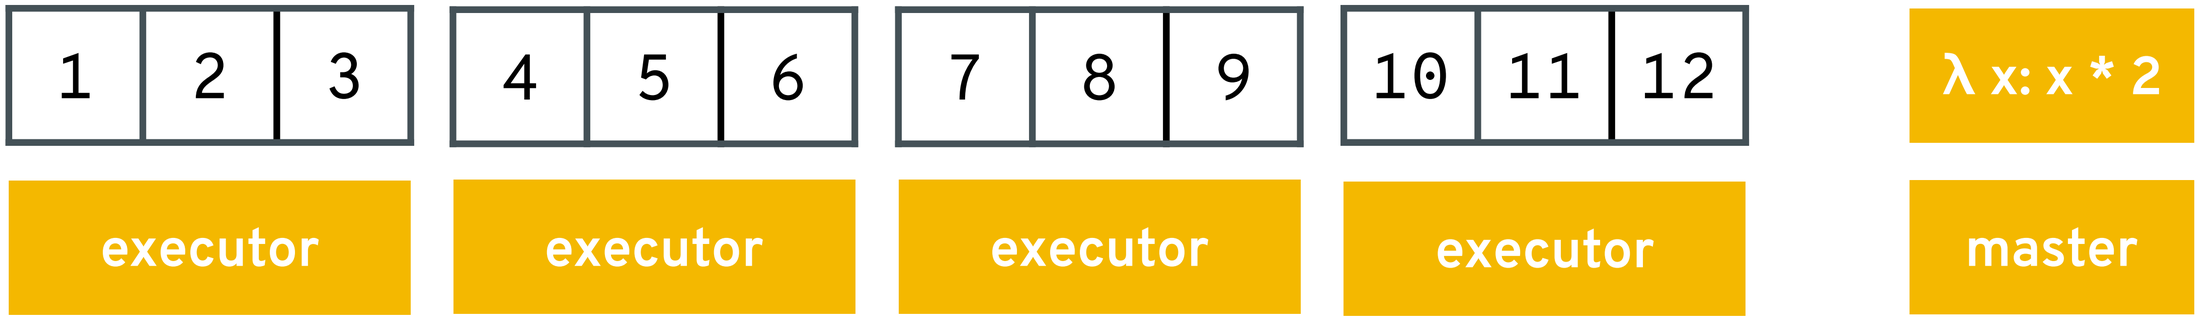
1
2
3
4
5
6
7
8
9
10
11
12
λ x: x * 2
executor
executor
executor
executor
master

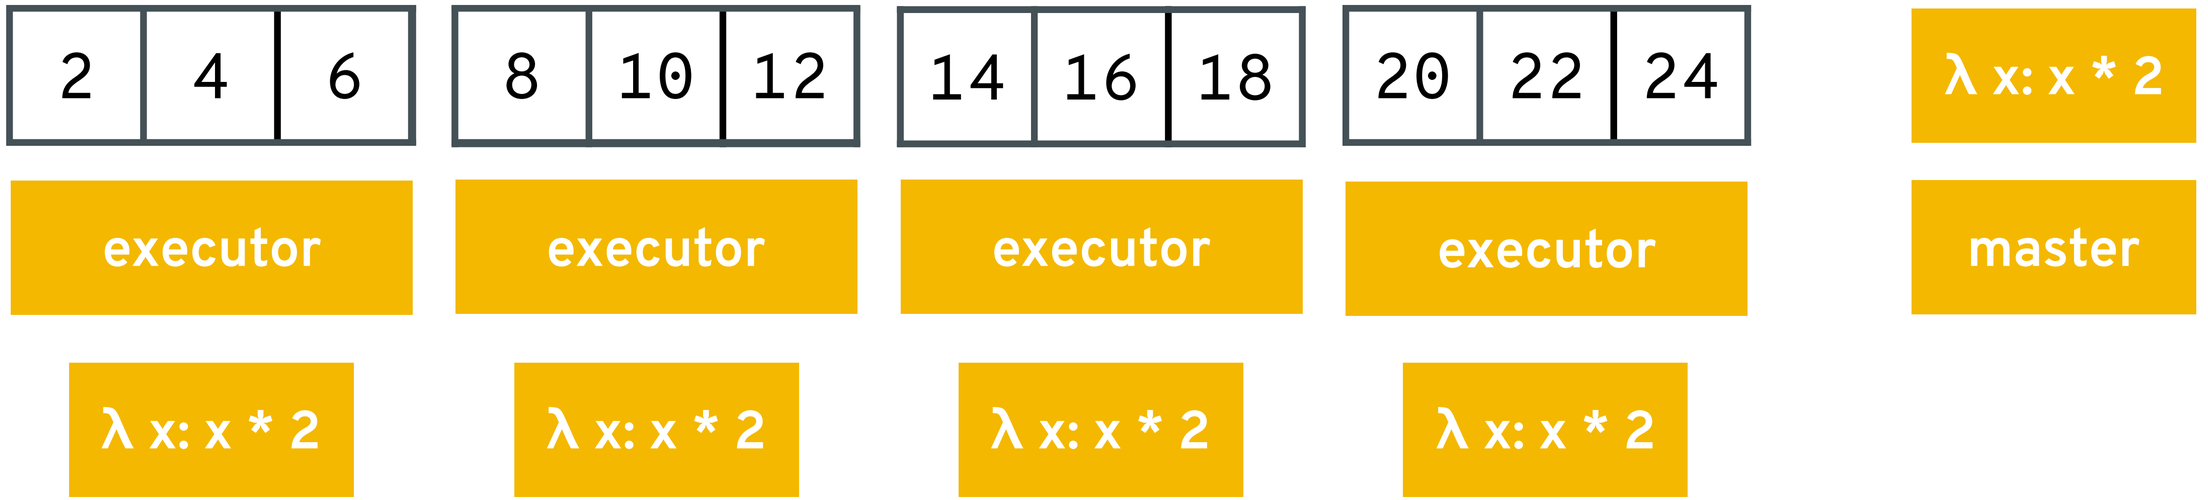
2
4
6
8
10
12
14
16
18
20
22
24
λ x: x * 2
executor
executor
executor
executor
master
λ x: x * 2
λ x: x * 2
λ x: x * 2
λ x: x * 2


2
4
6
8
10
12
20
22
24
λ x: x * 2
executor
executor
executor
master
λ x: x * 2
λ x: x * 2
λ x: x * 2

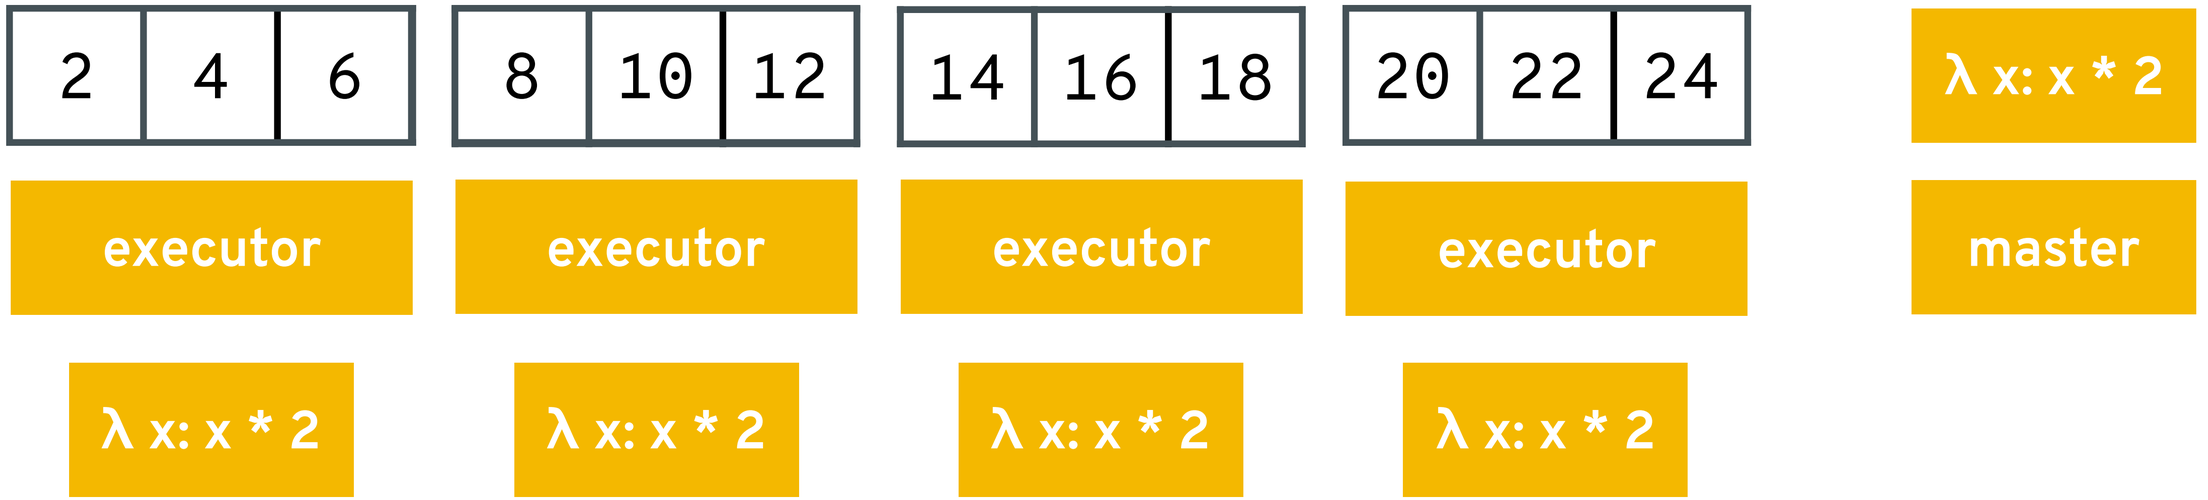
2
4
6
8
10
12
14
16
18
20
22
24
λ x: x * 2
executor
executor
executor
executor
master
λ x: x * 2
λ x: x * 2
λ x: x * 2
λ x: x * 2

Streaming data

Goal: use the **same abstraction** for batch and "streaming" (micro-batch) data by **dividing a stream** into **many small RDDs.**

input stream

Streaming data

Goal: use the **same abstraction** for batch and "streaming" (micro-batch) data by **dividing a stream** into **many small RDDs.**

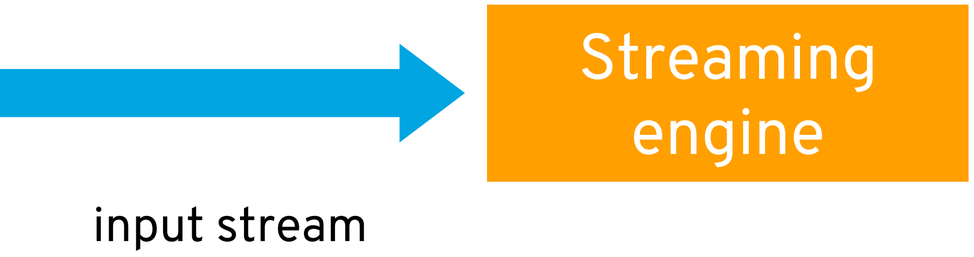

Streaming data

Goal: use the **same abstraction** for batch and "streaming" (micro-batch) data by **dividing a stream** into **many small RDDs.**

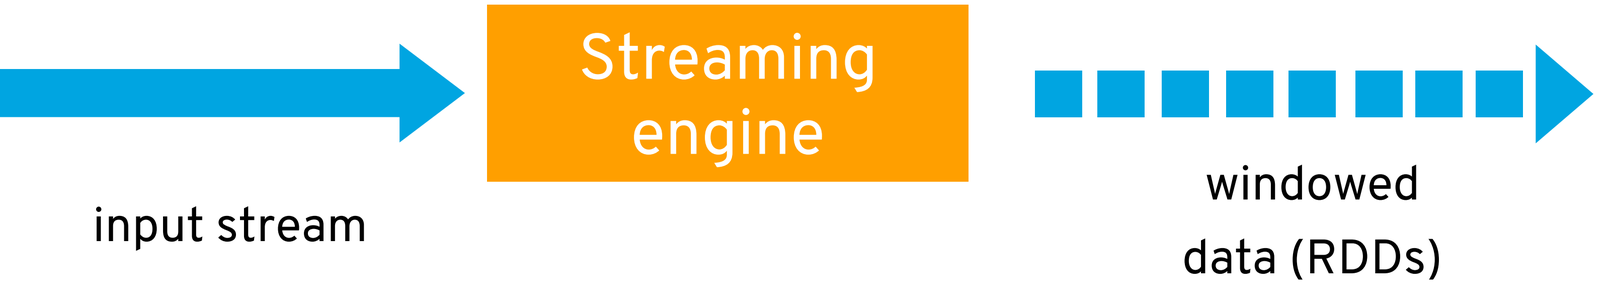

Streaming data

Goal: use the **same abstraction** for batch and "streaming" (micro-batch) data by **dividing a stream** into **many small RDDs.**

Streaming data

Goal: use the **same abstraction** for batch and "streaming" (micro-batch) data by **dividing a stream** into **many small RDDs.**

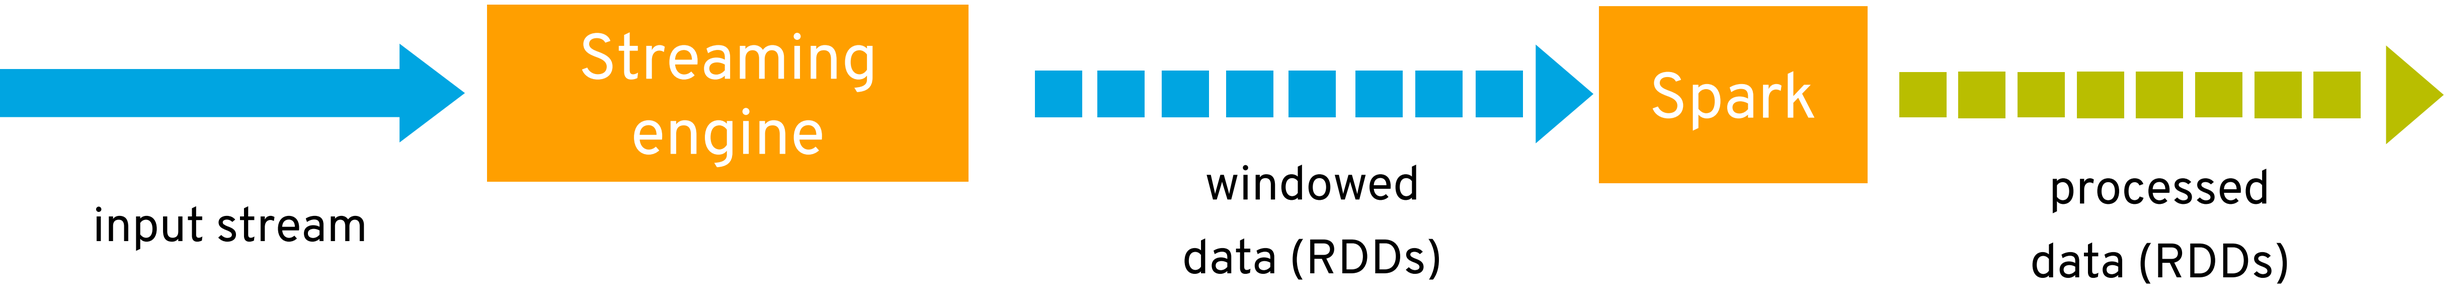

Structured queries

The capacity to run arbitrary code in RDDs is powerful but comes with **an important tradeoff**: Spark can't rearrange RDD programs to improve their performance.

Writing Spark programs with a **query DSL** allows Spark to generate **optimized execution plans**.

Structured query in Spark

SQL interface (unchecked syntax or semantics)

```
SELECT word, COUNT(*) FROM words GROUP BY word
```

Data frame interface (semantics checked at run-time)

```
words.groupBy('word').count()
```

Dataset interface (mostly checked at compile-time)

Query planning

```
SELECT * FROM A, B WHERE
  A.ID = B.ID AND
  uncommon(A.X) AND
  extremelyRare(B.Y)
```

A naïve plan

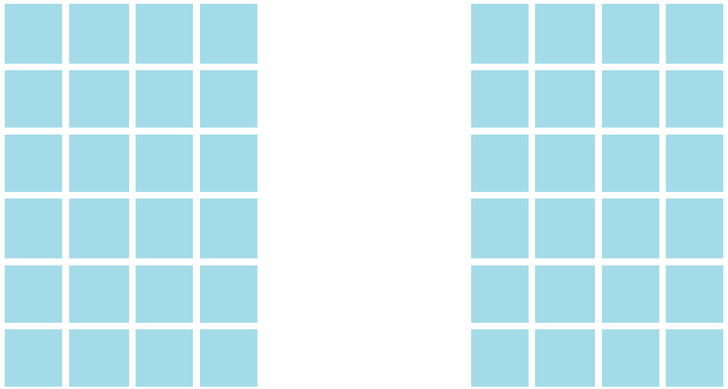

A naïve plan

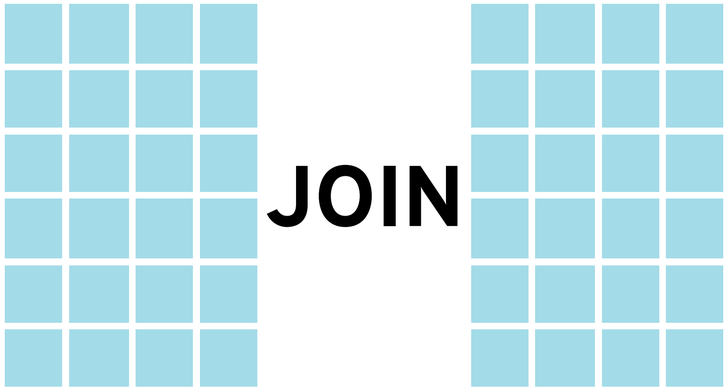

A naïve plan

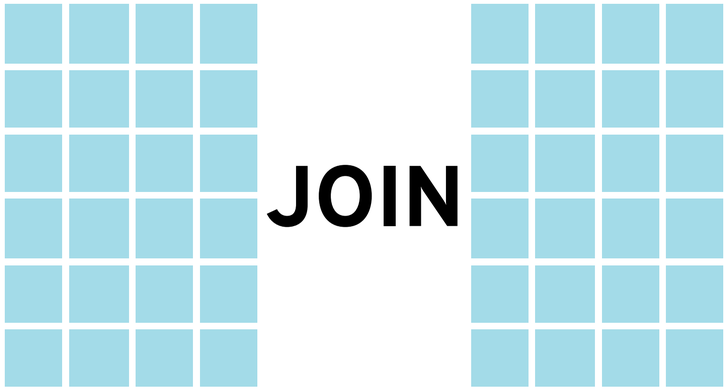

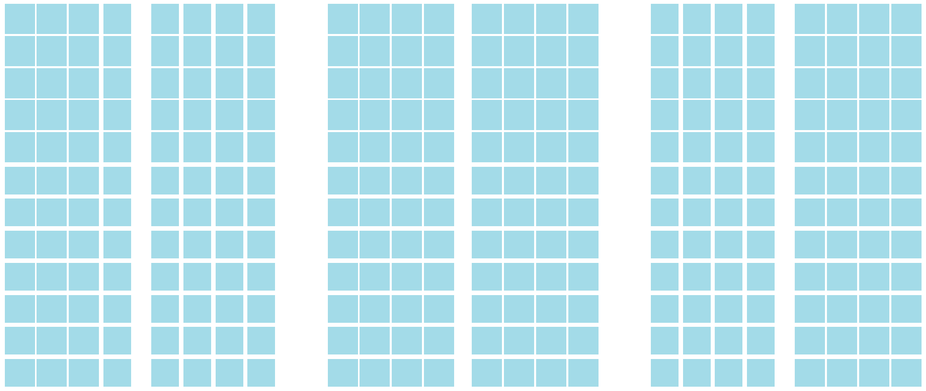

A naïve plan

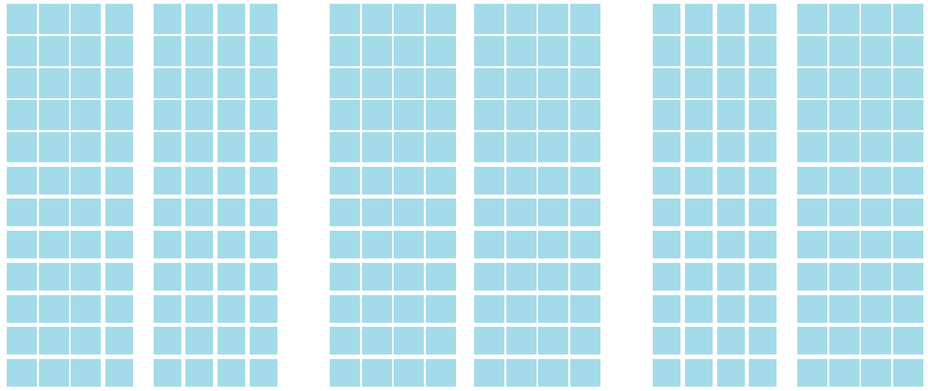

A naïve plan

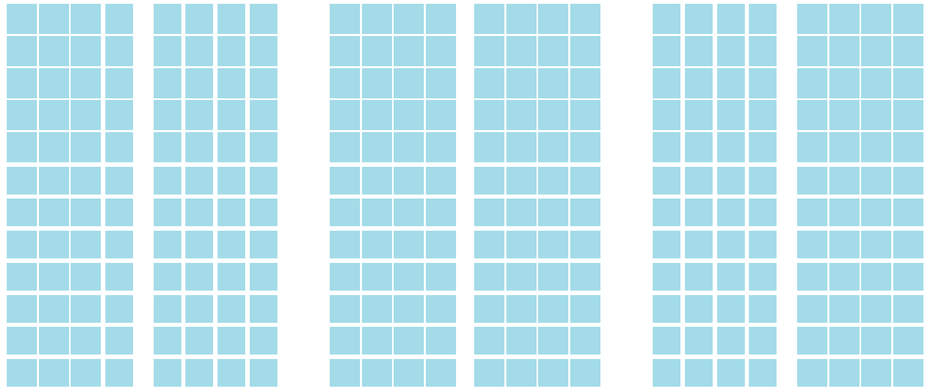

FILTER

A naïve plan

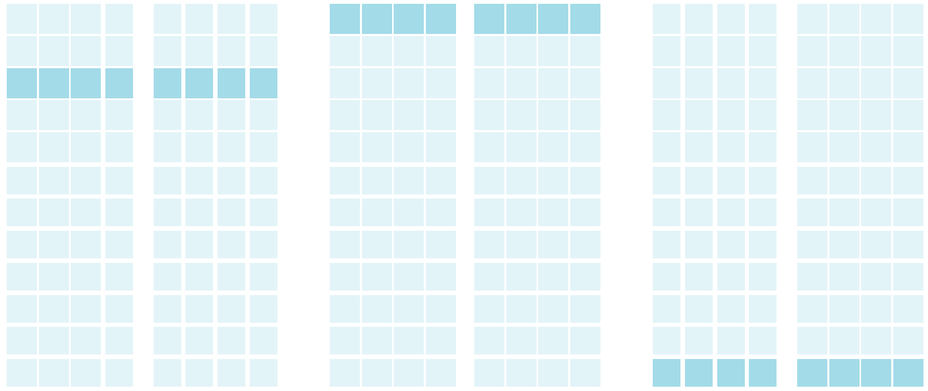

FILTER

An optimized plan

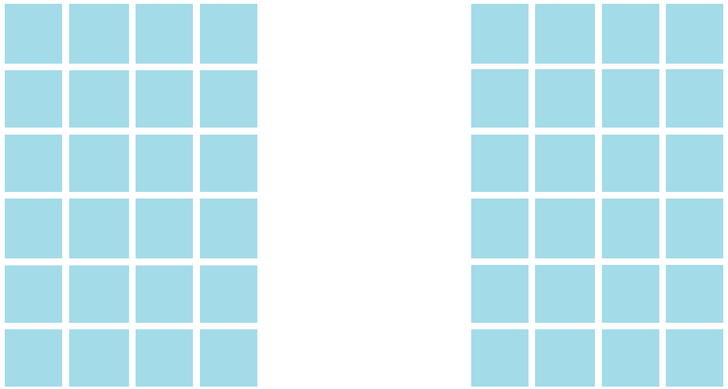

An optimized plan

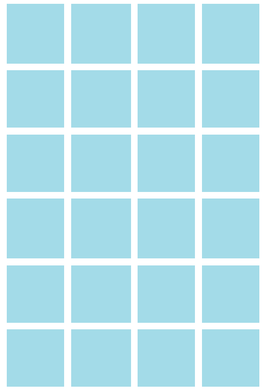

FILTER

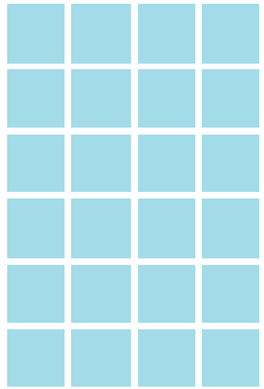

FILTER

An optimized plan

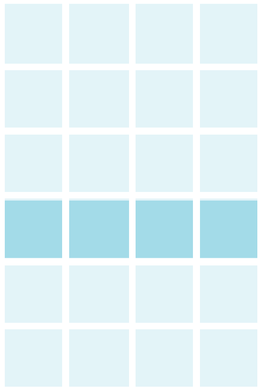

FILTER

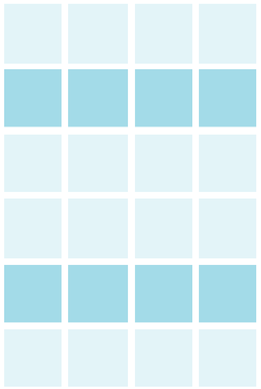

FILTER

An optimized plan

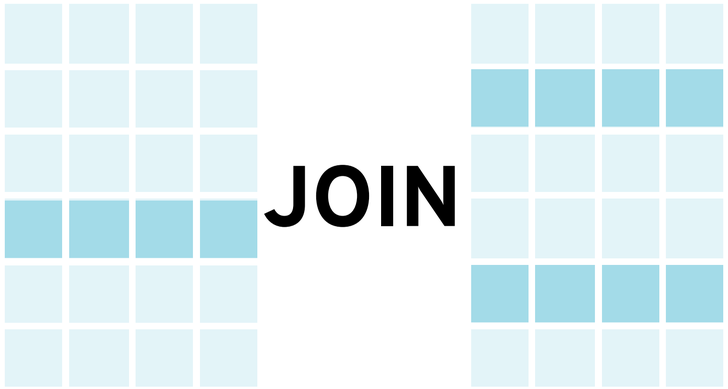

An optimized plan

Structured streaming combines **stream processing** with **query planning** for **high-performance analytics on events**!

The Kappa architecture

events

The Kappa architecture

The Kappa architecture

The Kappa architecture

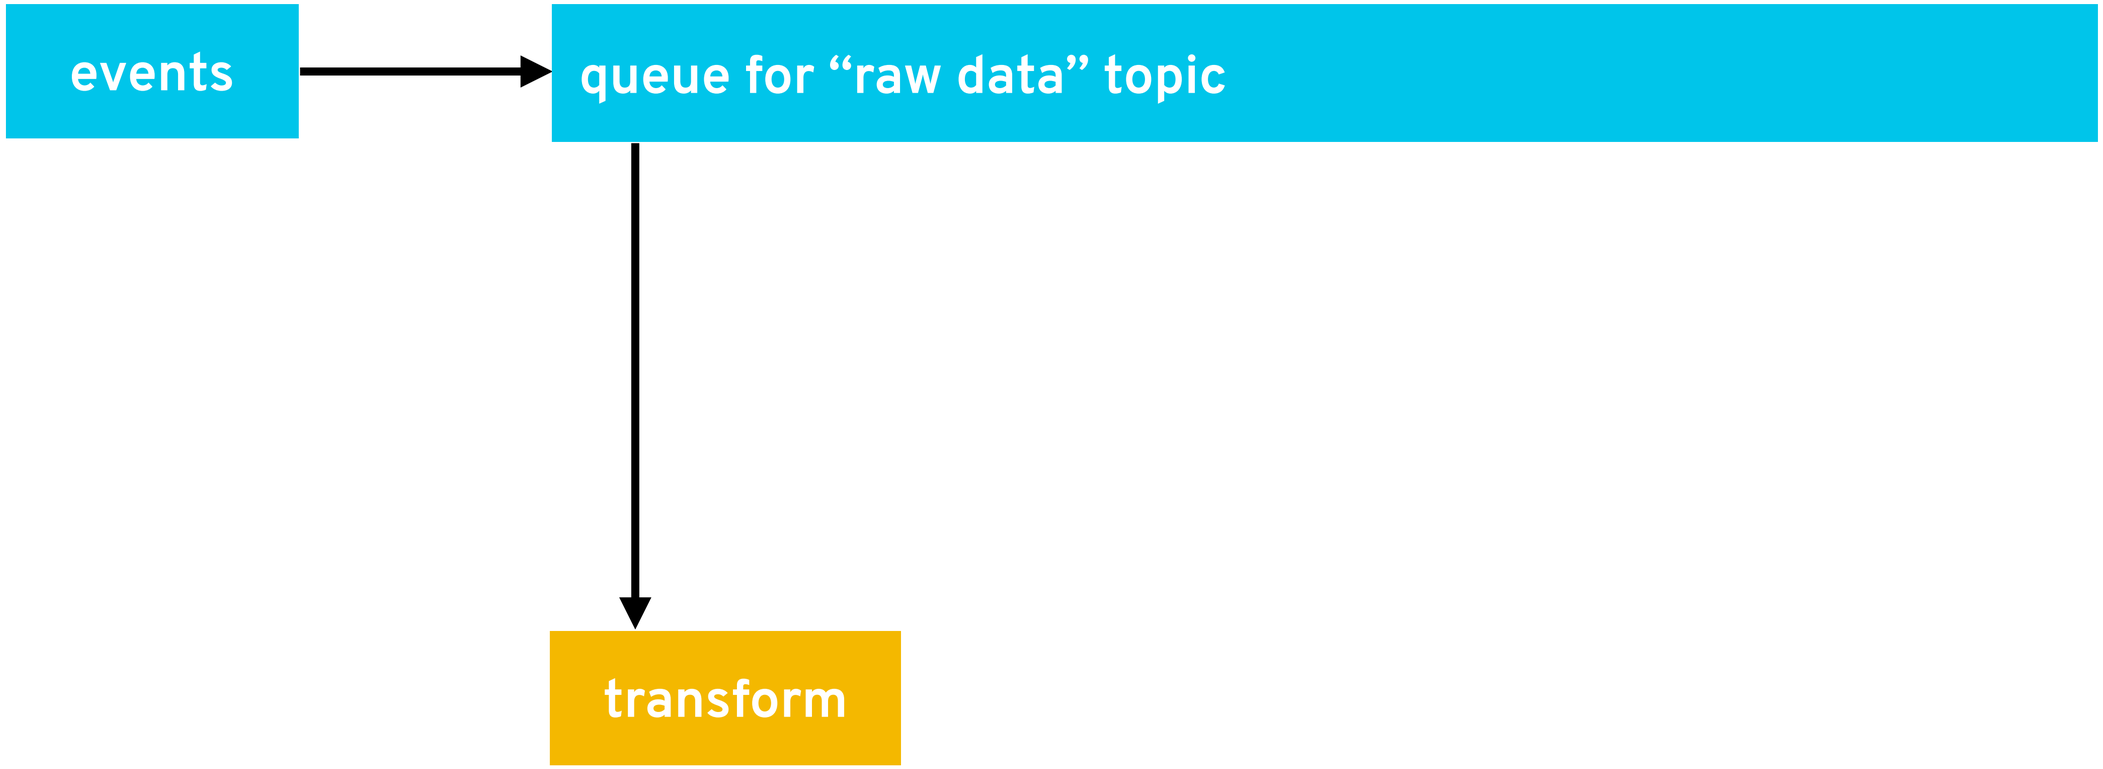

The Kappa architecture

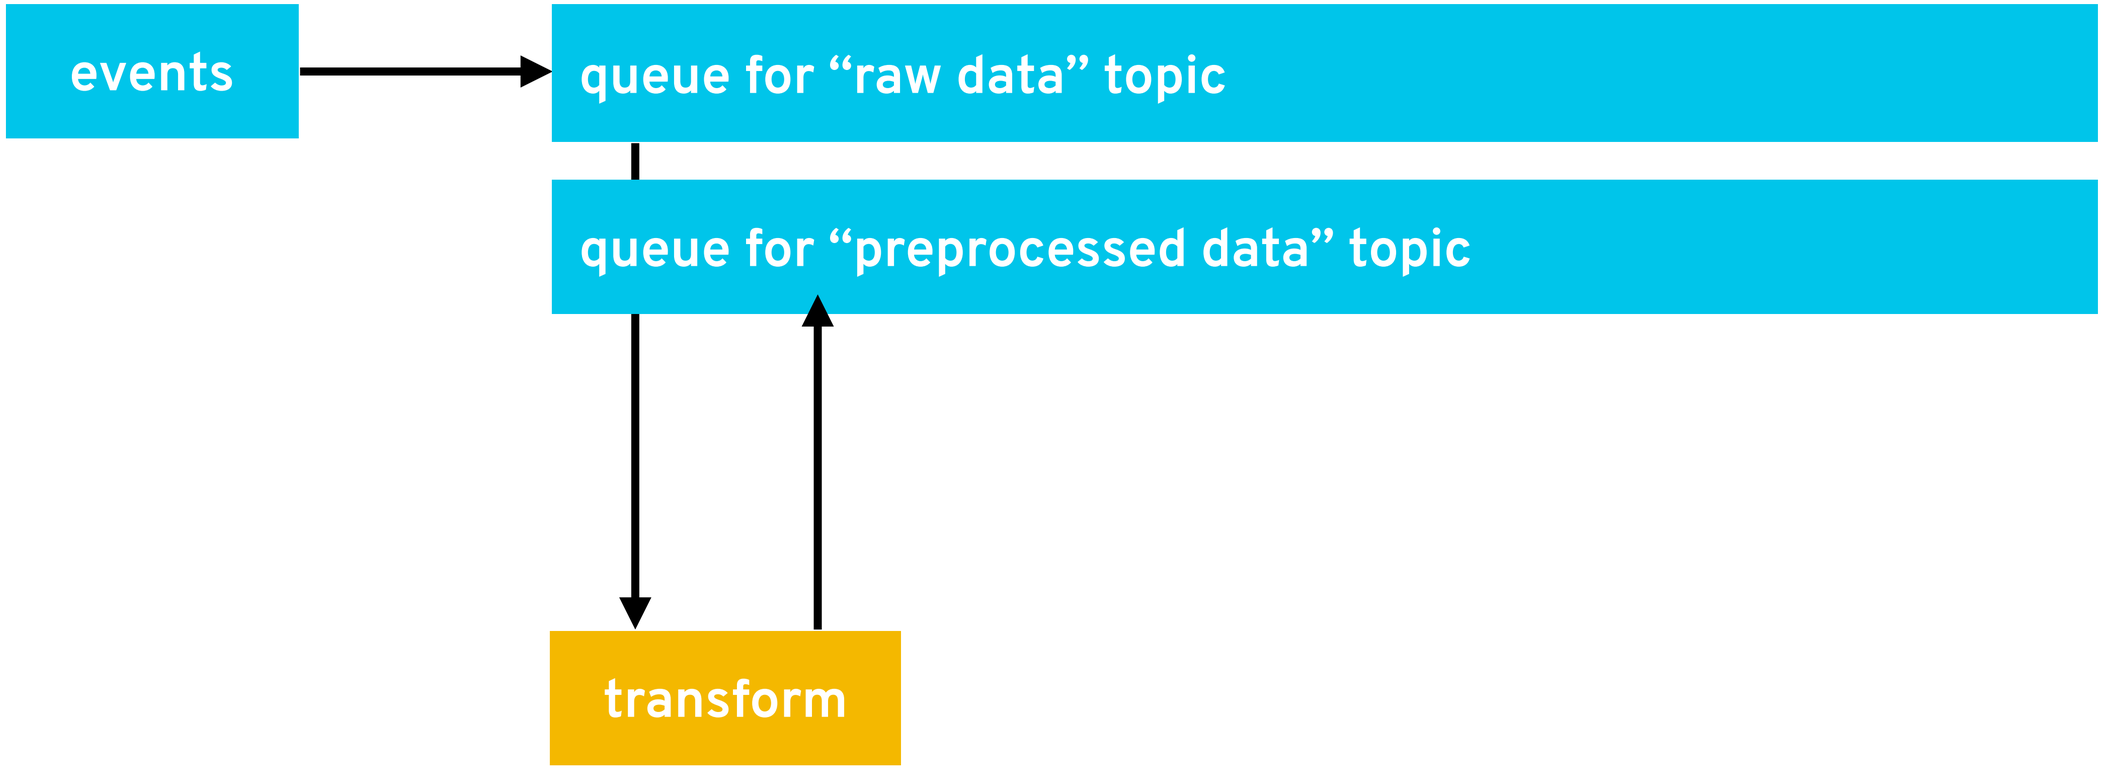

The Kappa architecture

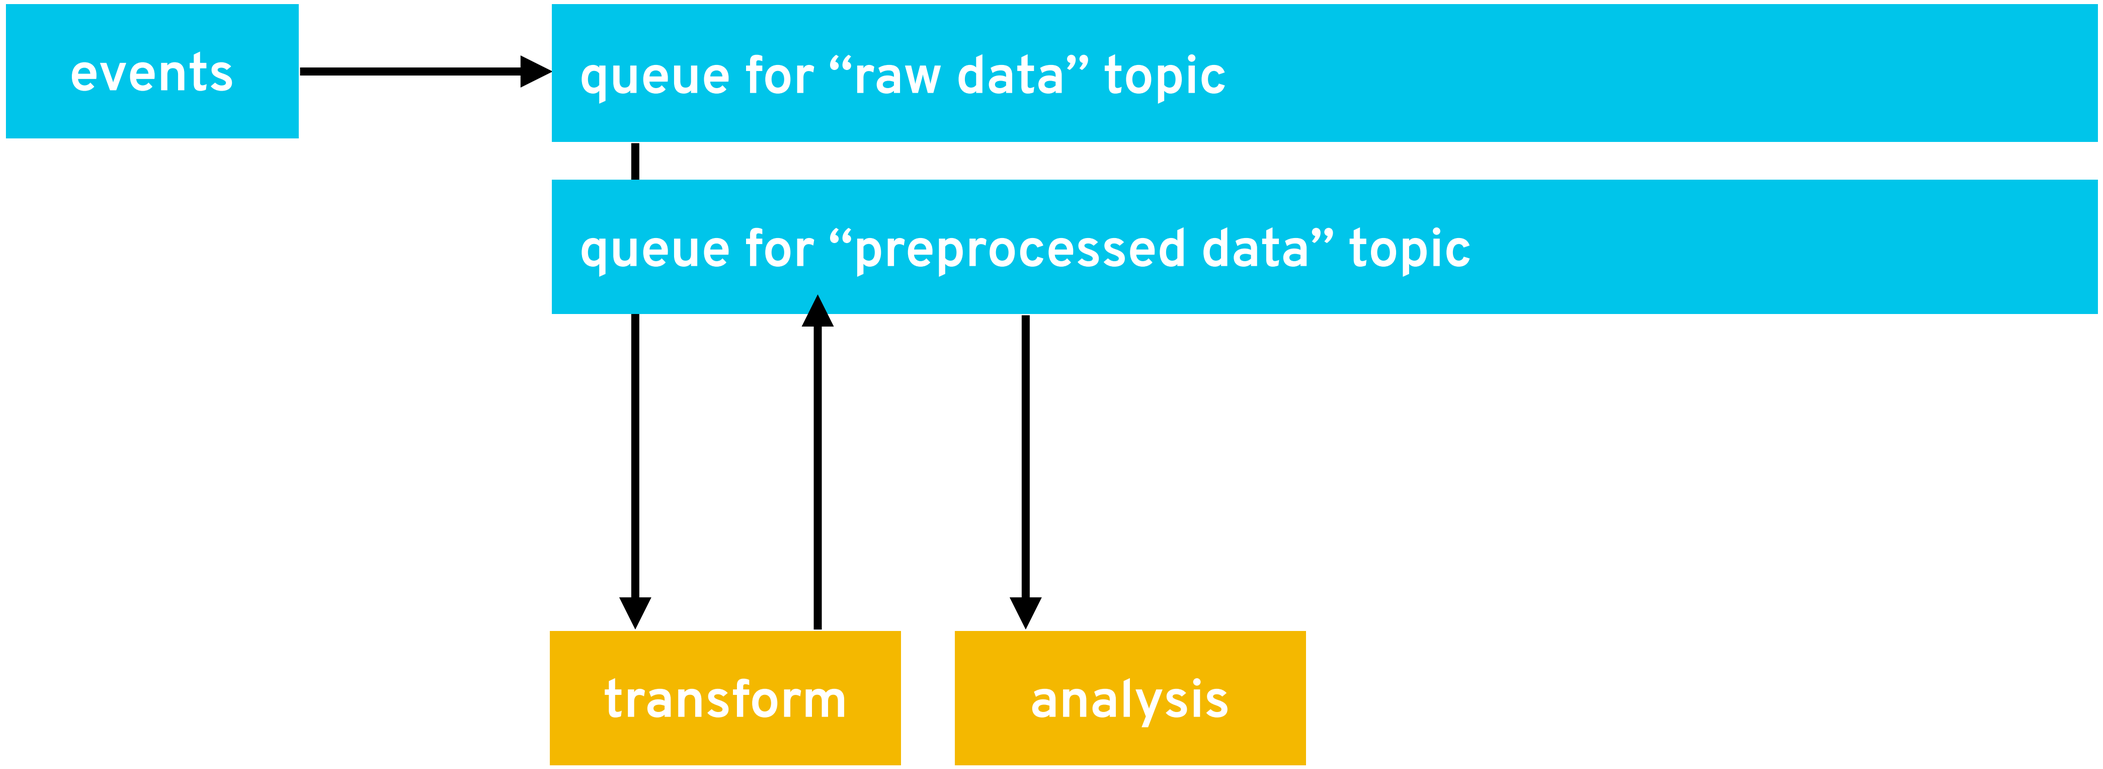

The Kappa architecture

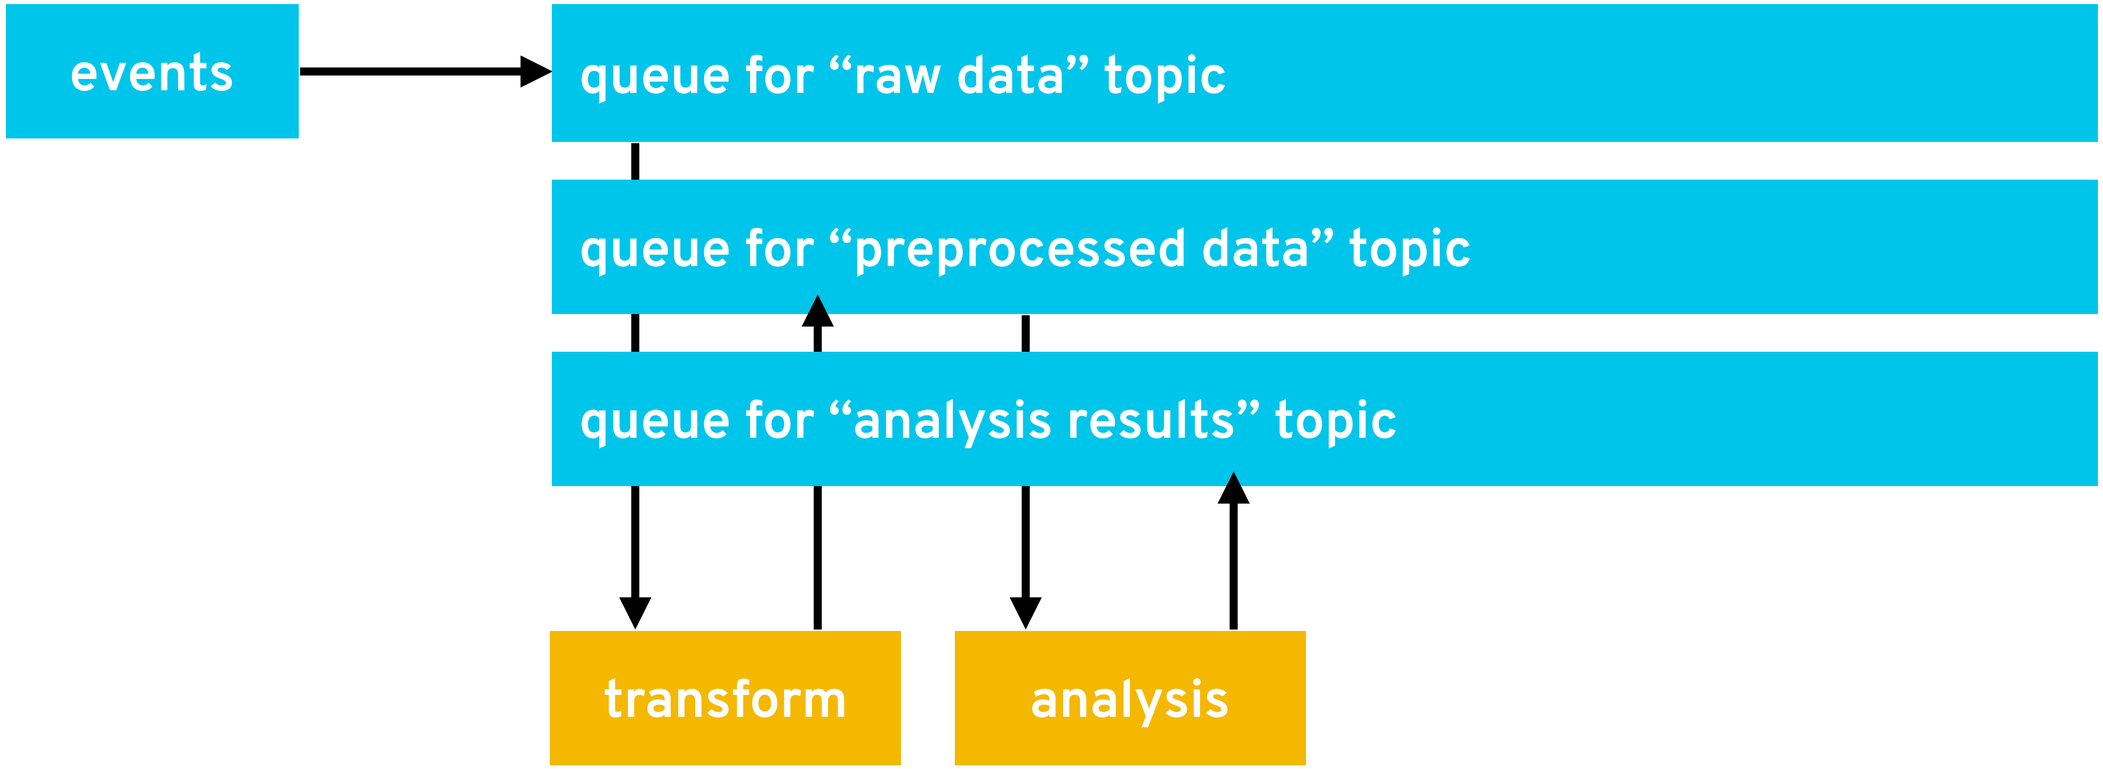

The Kappa architecture

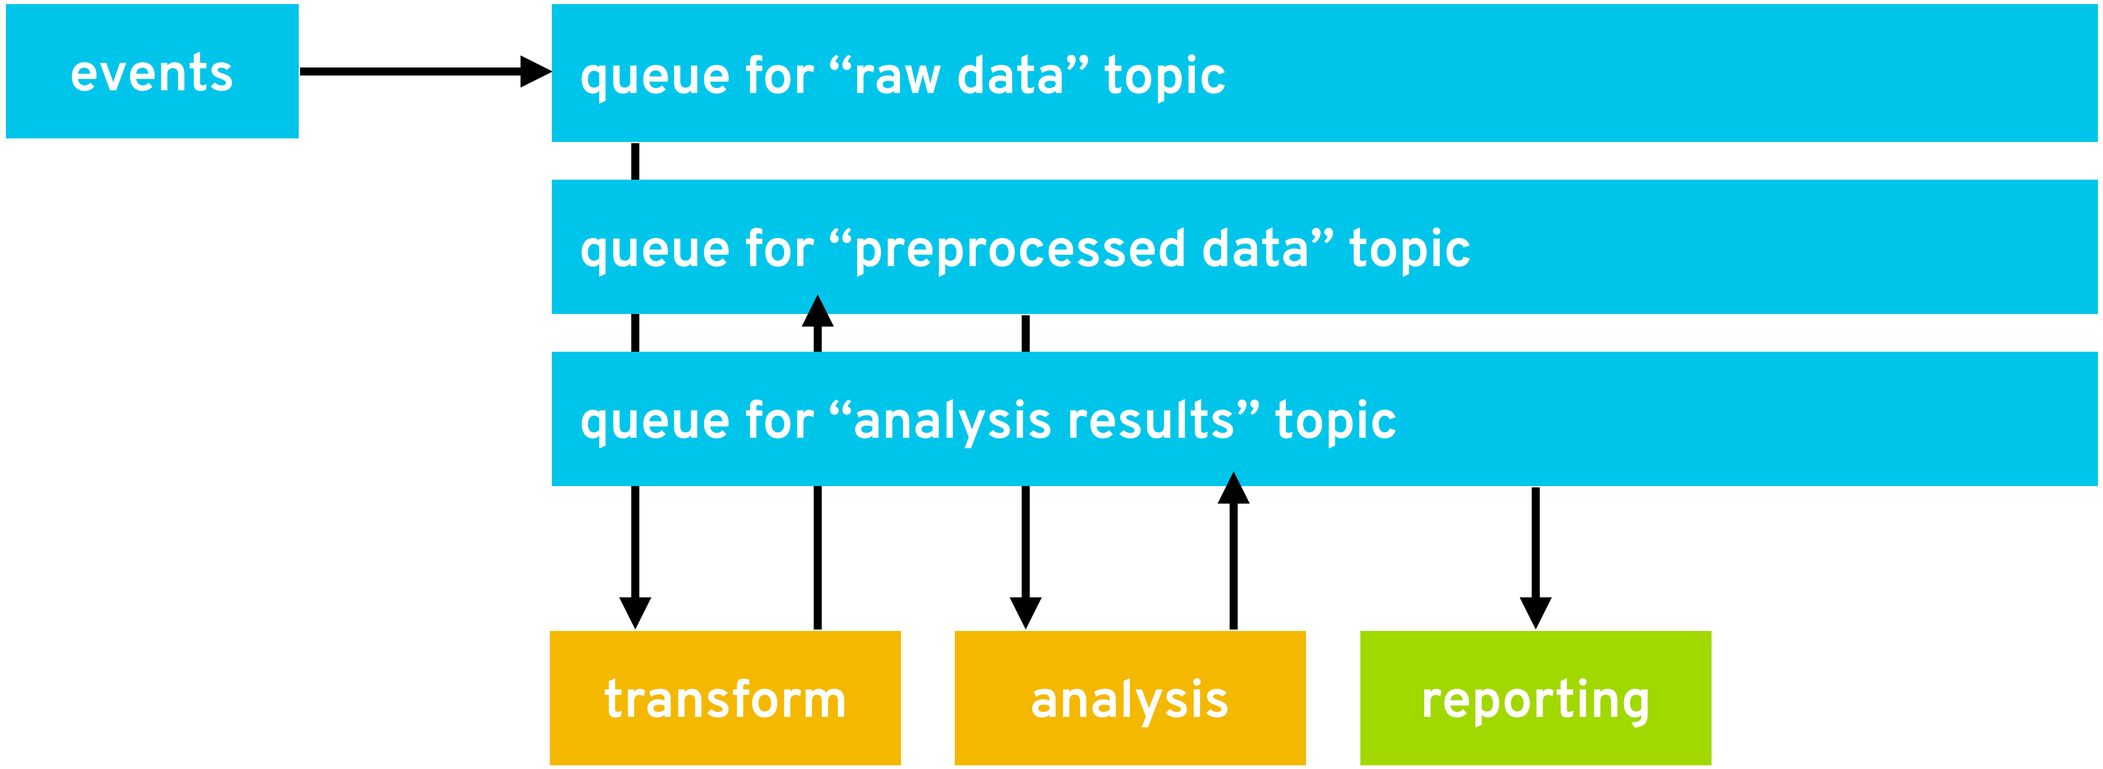

The Kappa architecture

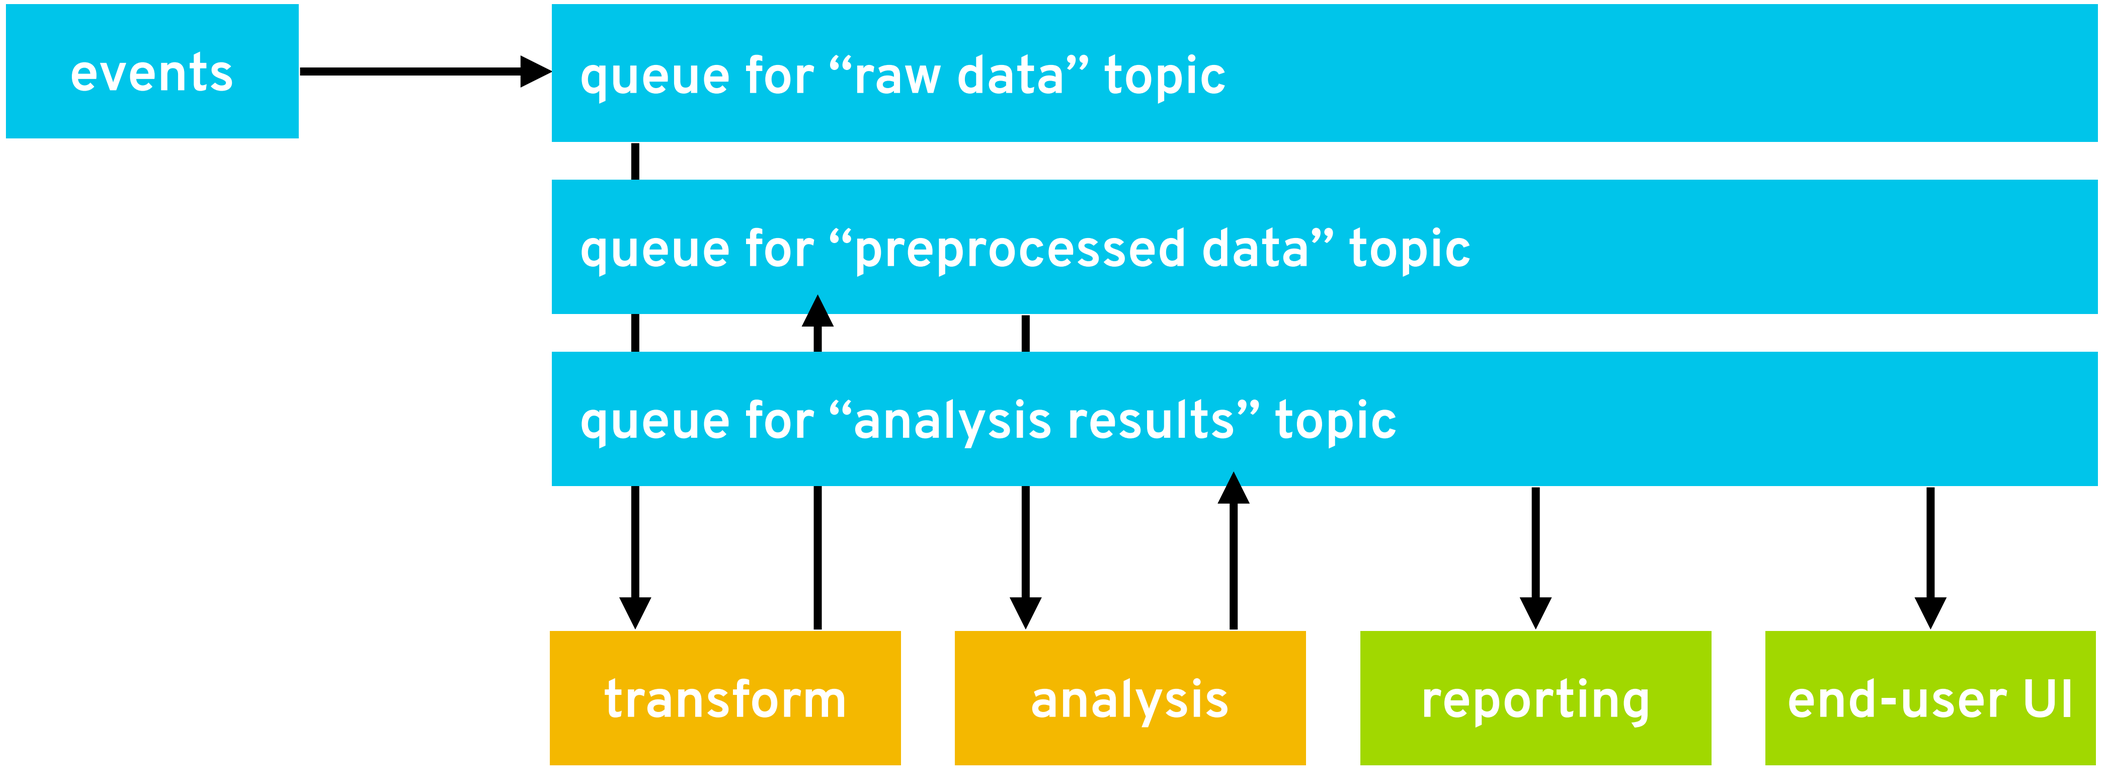

fedmsg bus

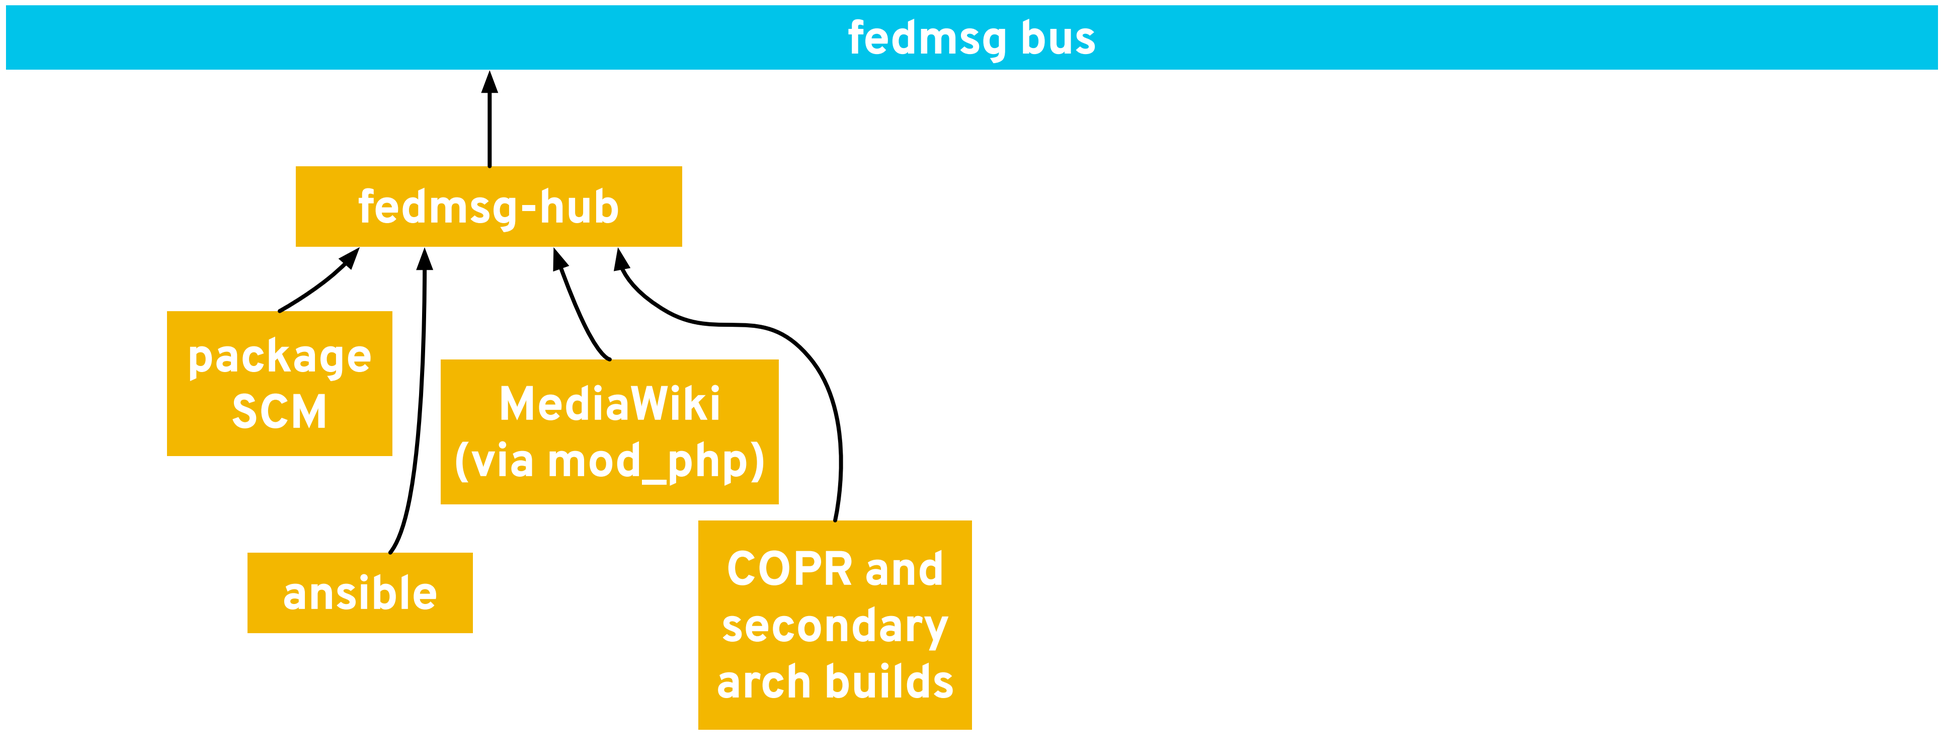
fedmsg bus
fedmsg-hub
package SCM
ansible
MediaWiki (via mod_php)
COPR and secondary arch builds

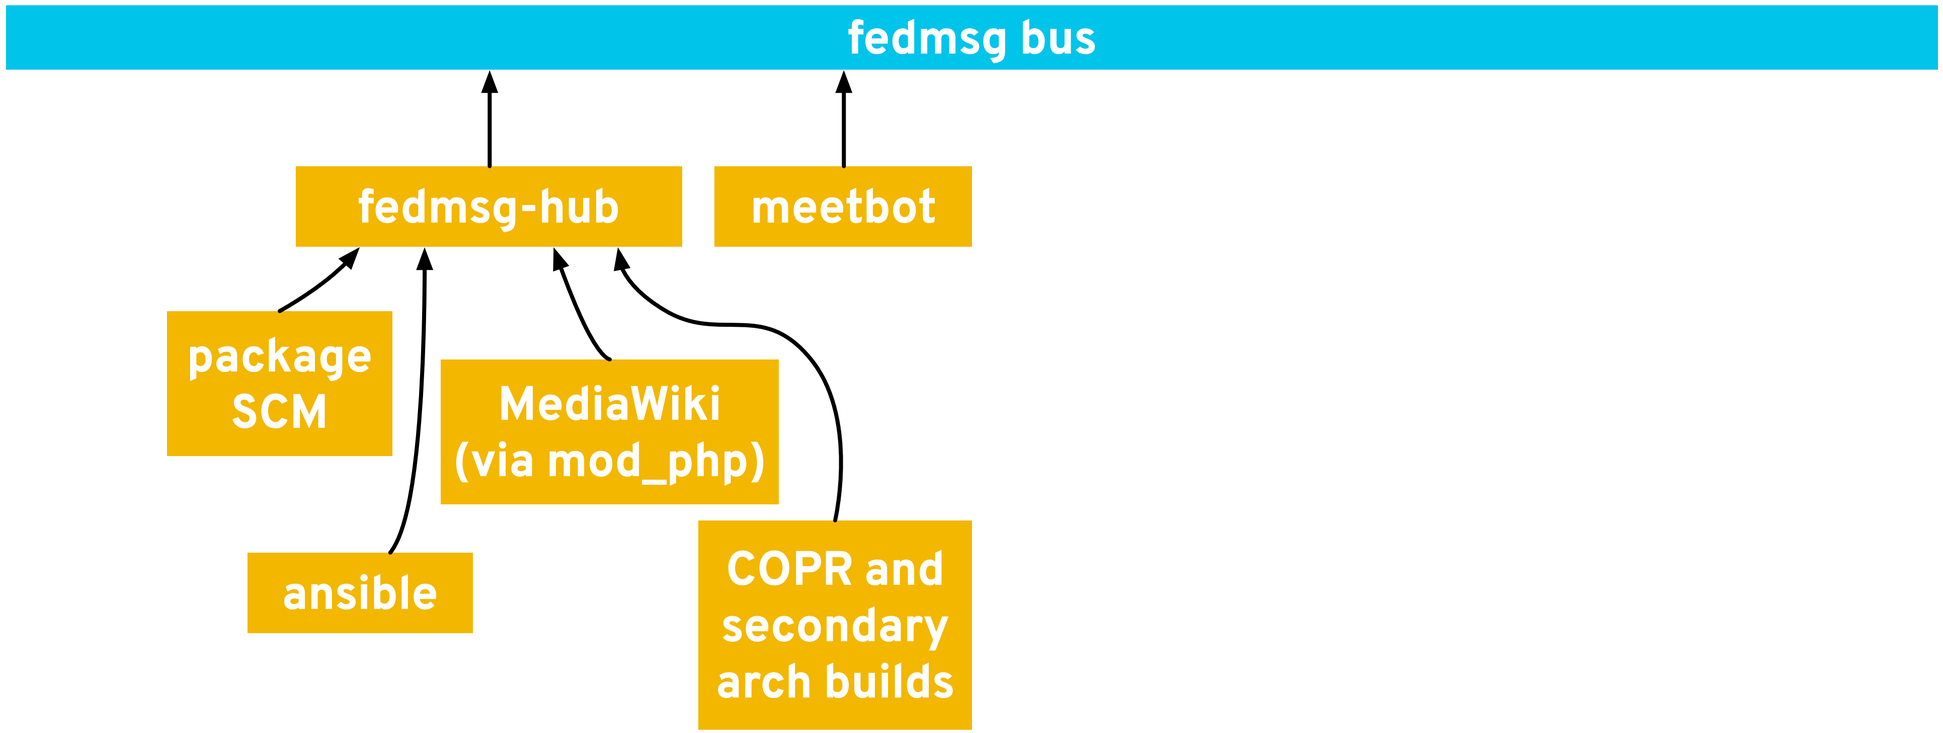
fedmsg bus
fedmsg-hub
meetbot
package SCM
MediaWiki (via mod_php)
ansible
COPR and secondary arch builds

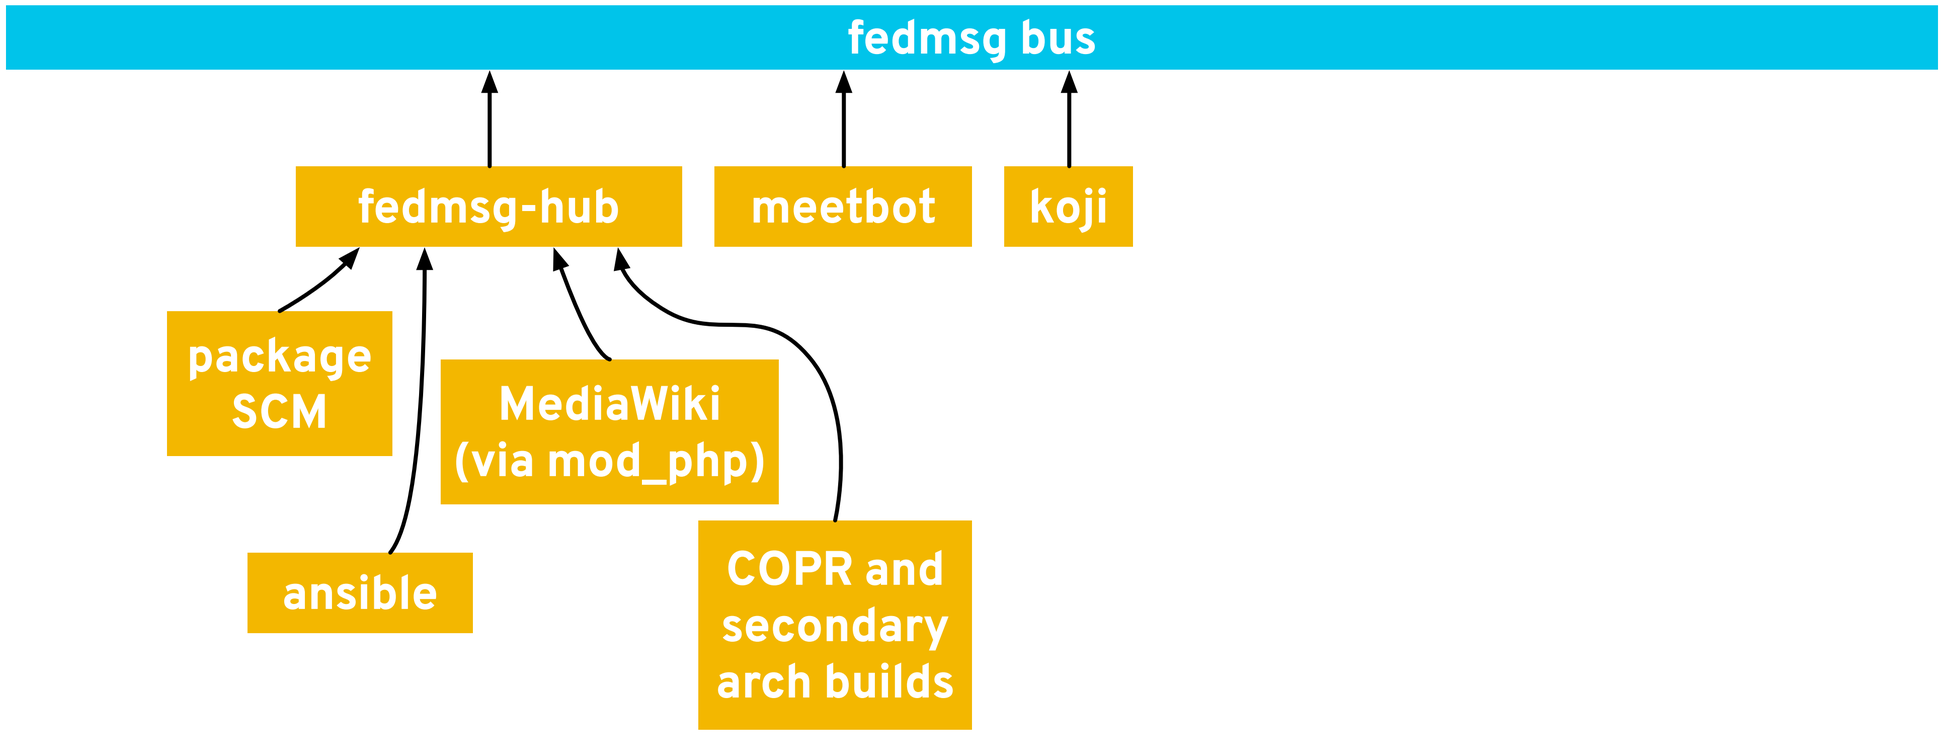
fedmsg bus
fedmsg-hub
meetbot
koji
package SCM
MediaWiki (via mod_php)
ansible
COPR and secondary arch builds

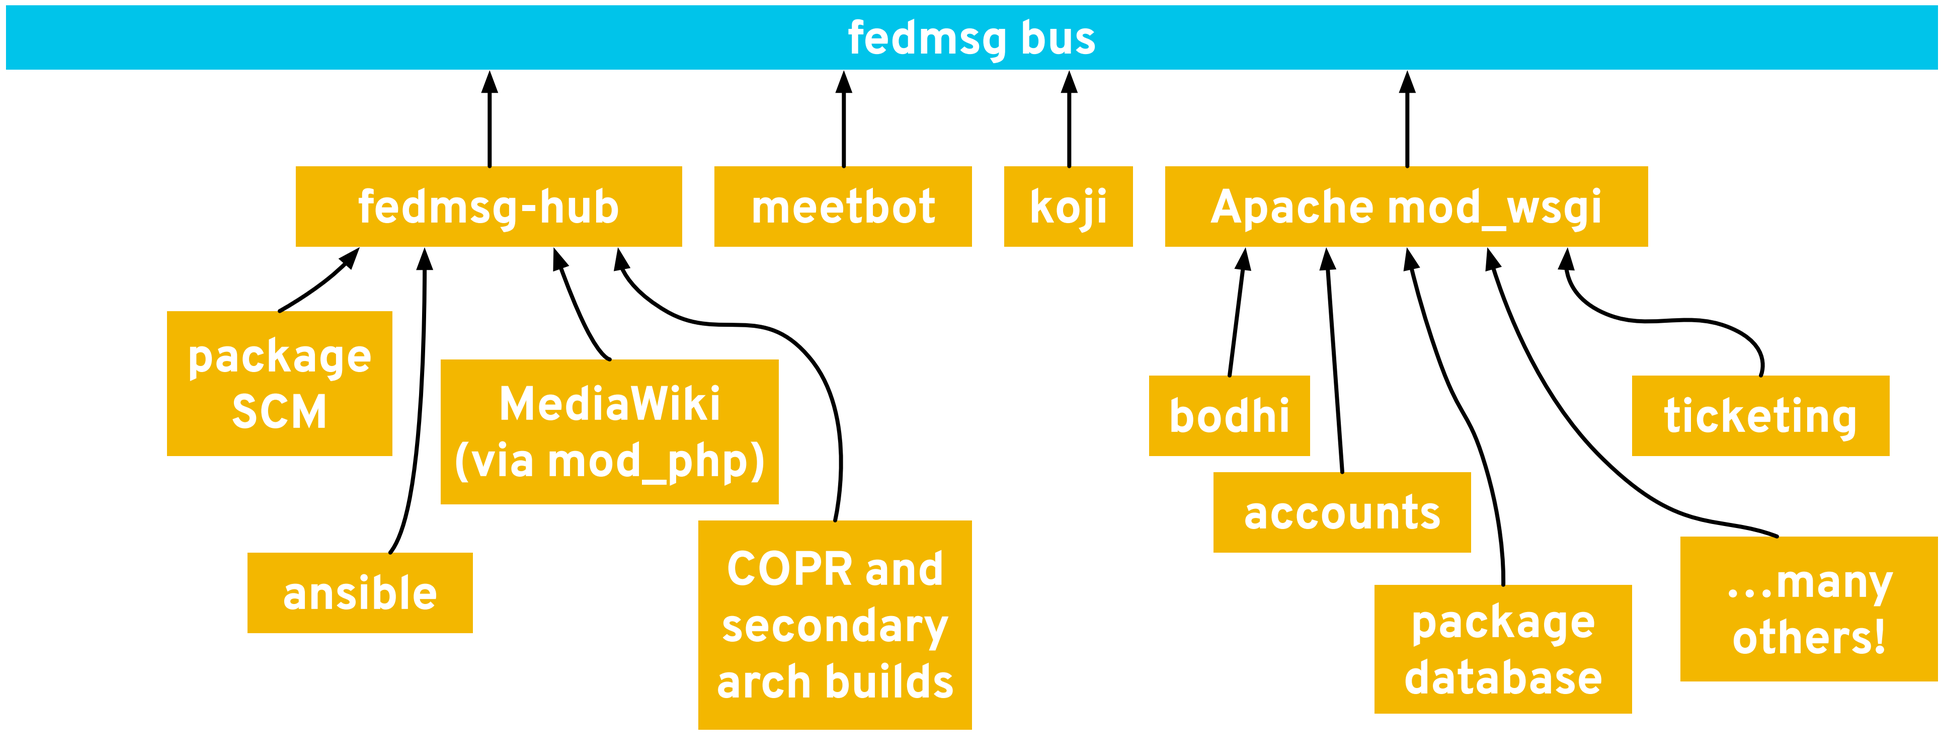
fedmsg bus
fedmsg-hub
meetbot
koji
Apache mod_wsgi
package SCM
MediaWiki (via mod_php)
ansible
COPR and secondary arch builds
bodhi
accounts
package database
ticketing
...many others!

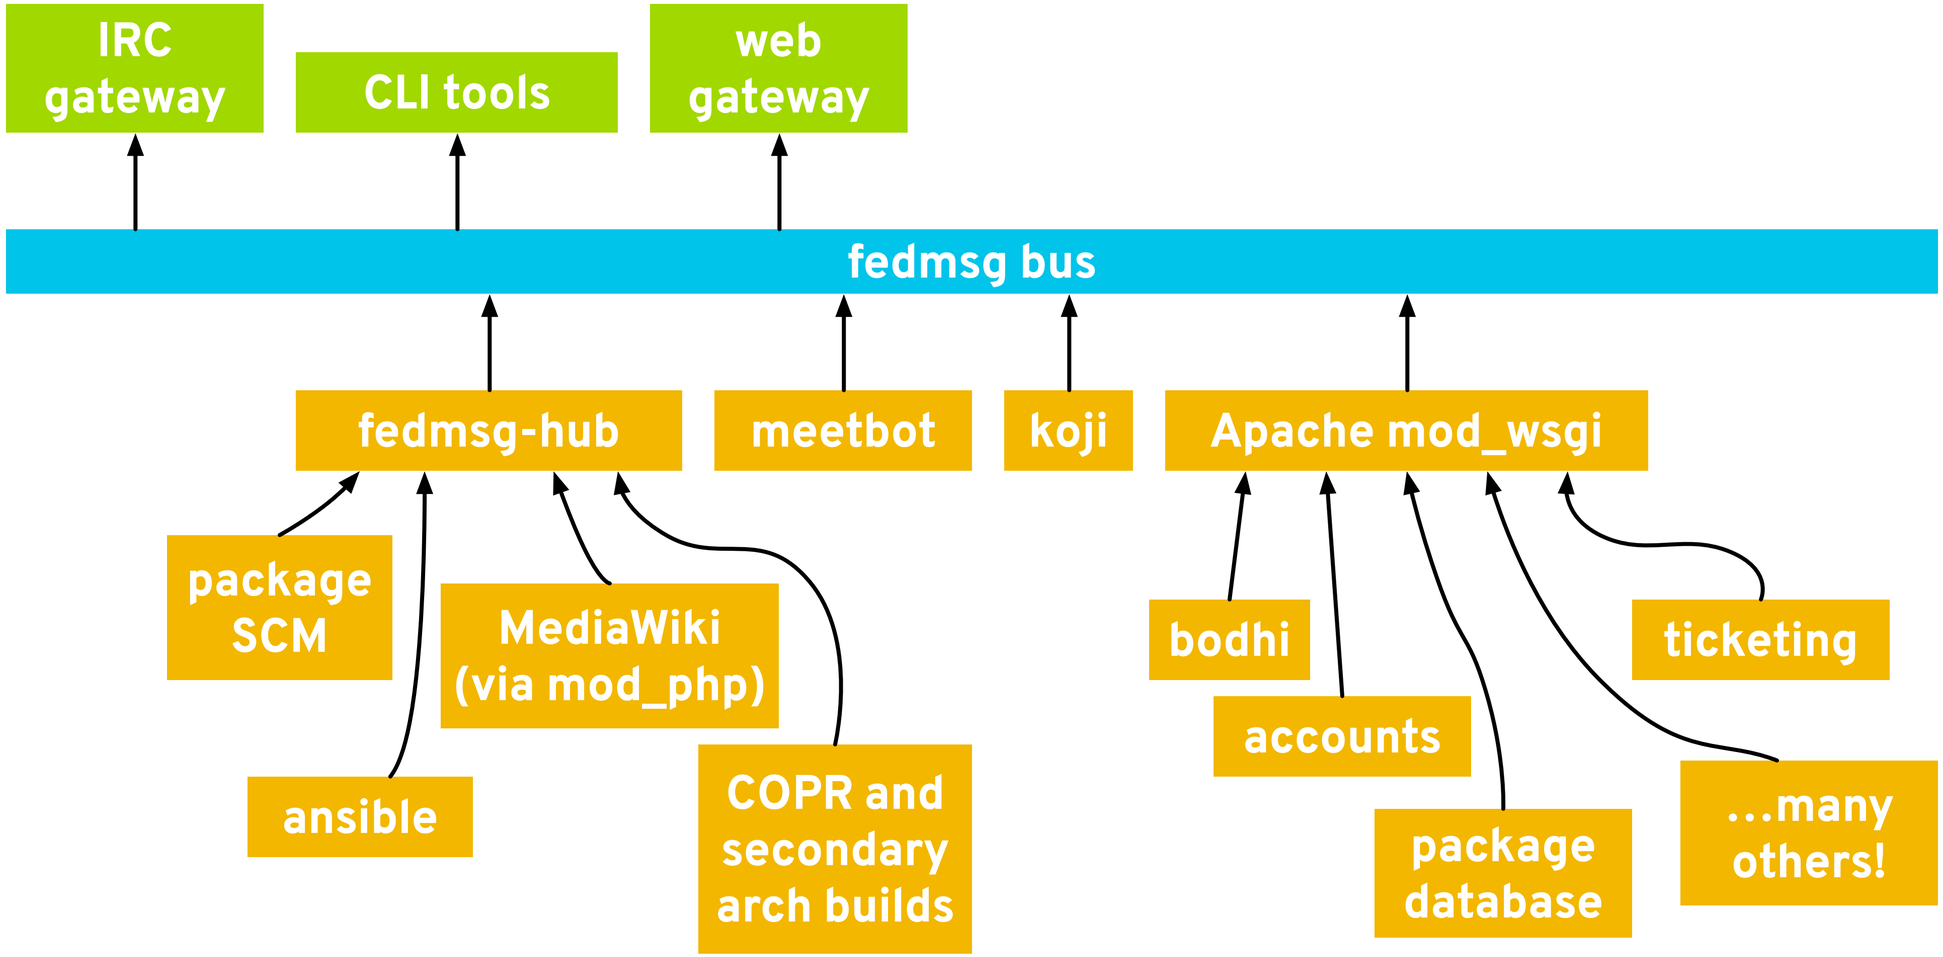
IRC gateway
CLI tools
web gateway
fedmsg bus
fedmsg-hub
meetbot
koji
Apache mod_wsgi
package SCM
MediaWiki (via mod_php)
ansible
COPR and secondary arch builds
bodhi
accounts
package database
ticketing
...many others!

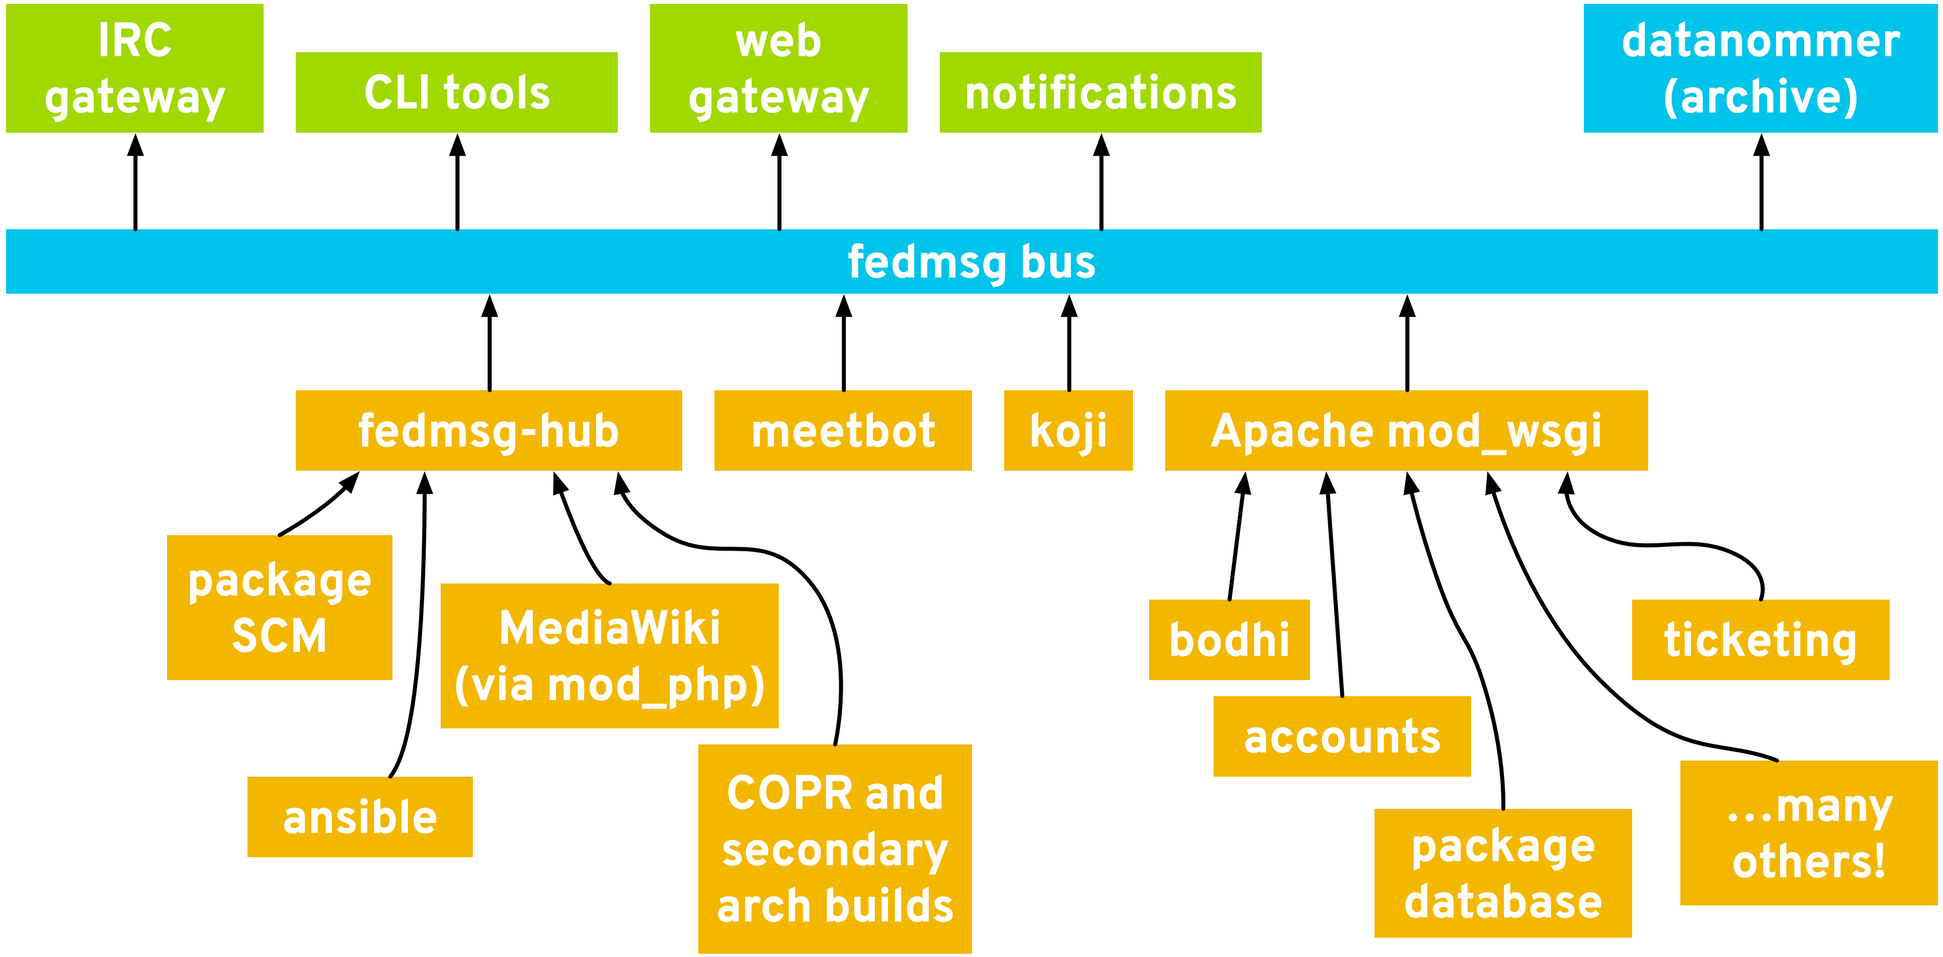
IRC gateway
CLI tools
web gateway
notifications
datanommer (archive)
fedmsg bus
fedmsg-hub
meetbot
koji
Apache mod_wsgi
package SCM
MediaWiki (via mod_php)
ansible
COPR and secondary arch builds
bodhi
accounts
package database
ticketing
...many others!

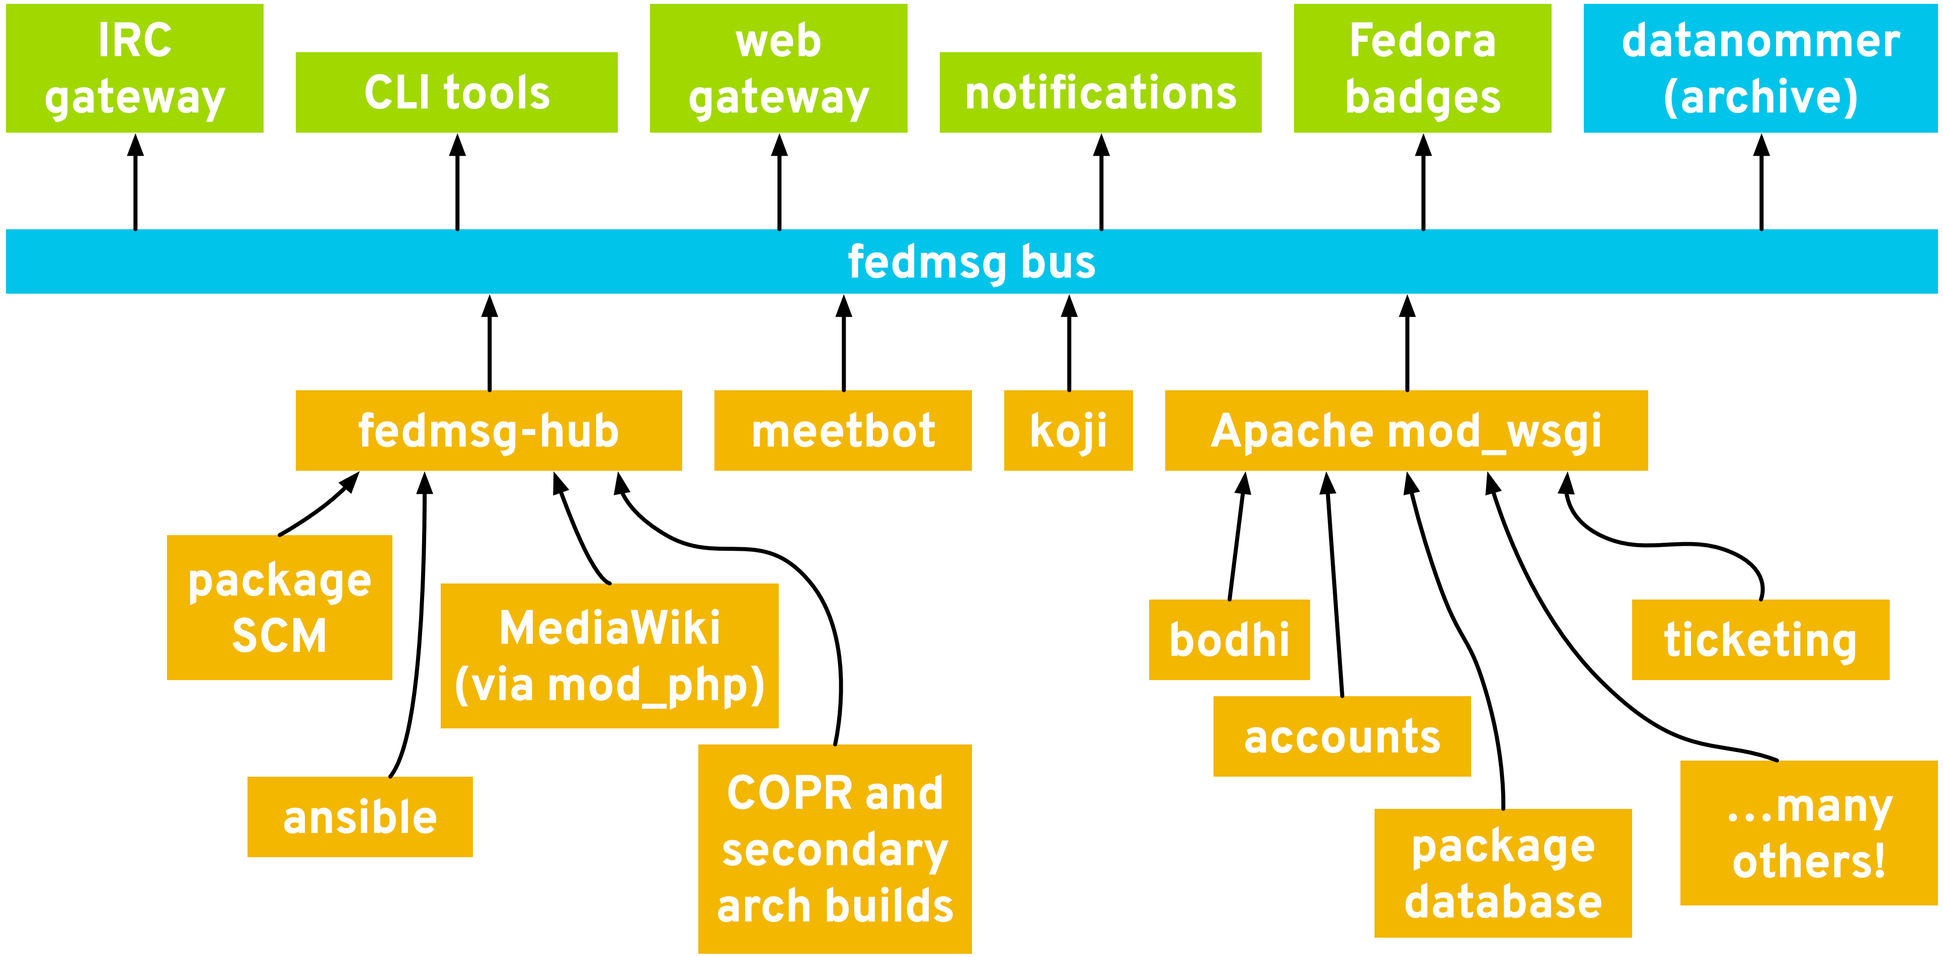
IRC gateway
CLI tools
web gateway
notifications
Fedora badges
datanommer (archive)
fedmsg bus
fedmsg-hub
meetbot
koji
Apache mod_wsgi
package SCM
MediaWiki (via mod_php)
ansible
COPR and secondary arch builds
bodhi
accounts
package database
ticketing
...many others!

Apache Kafka

Apache Kafka

Apache Kafka

Multitenant compute clusters

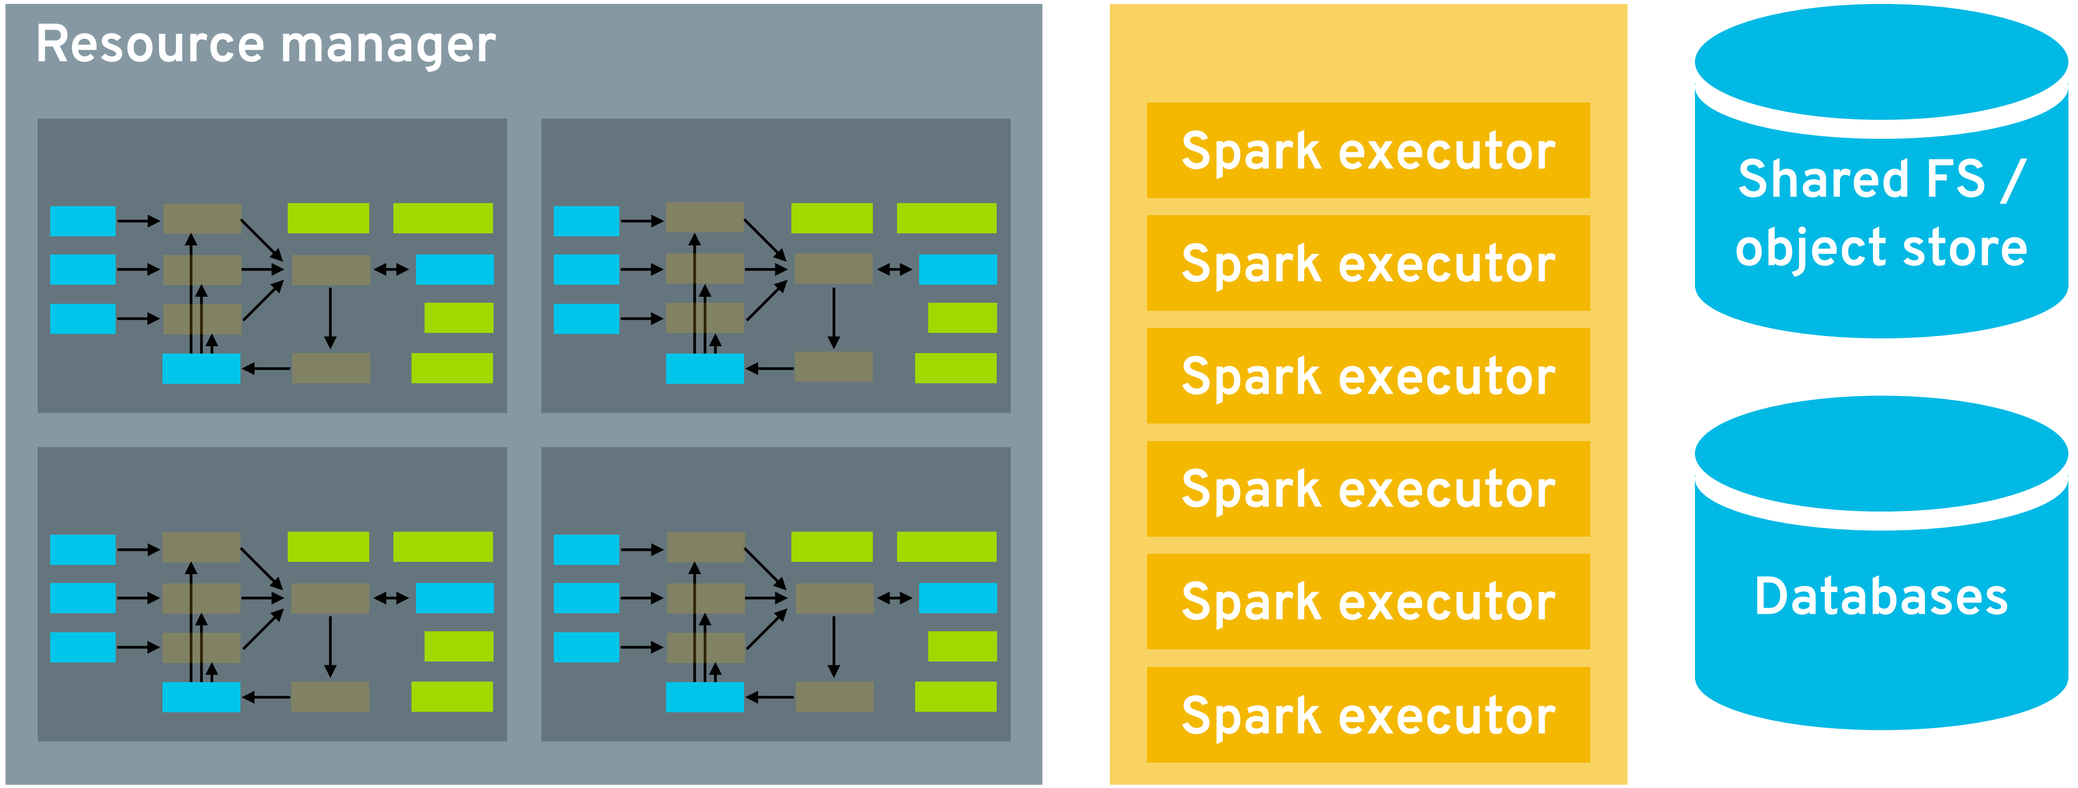

Multitenant compute clusters

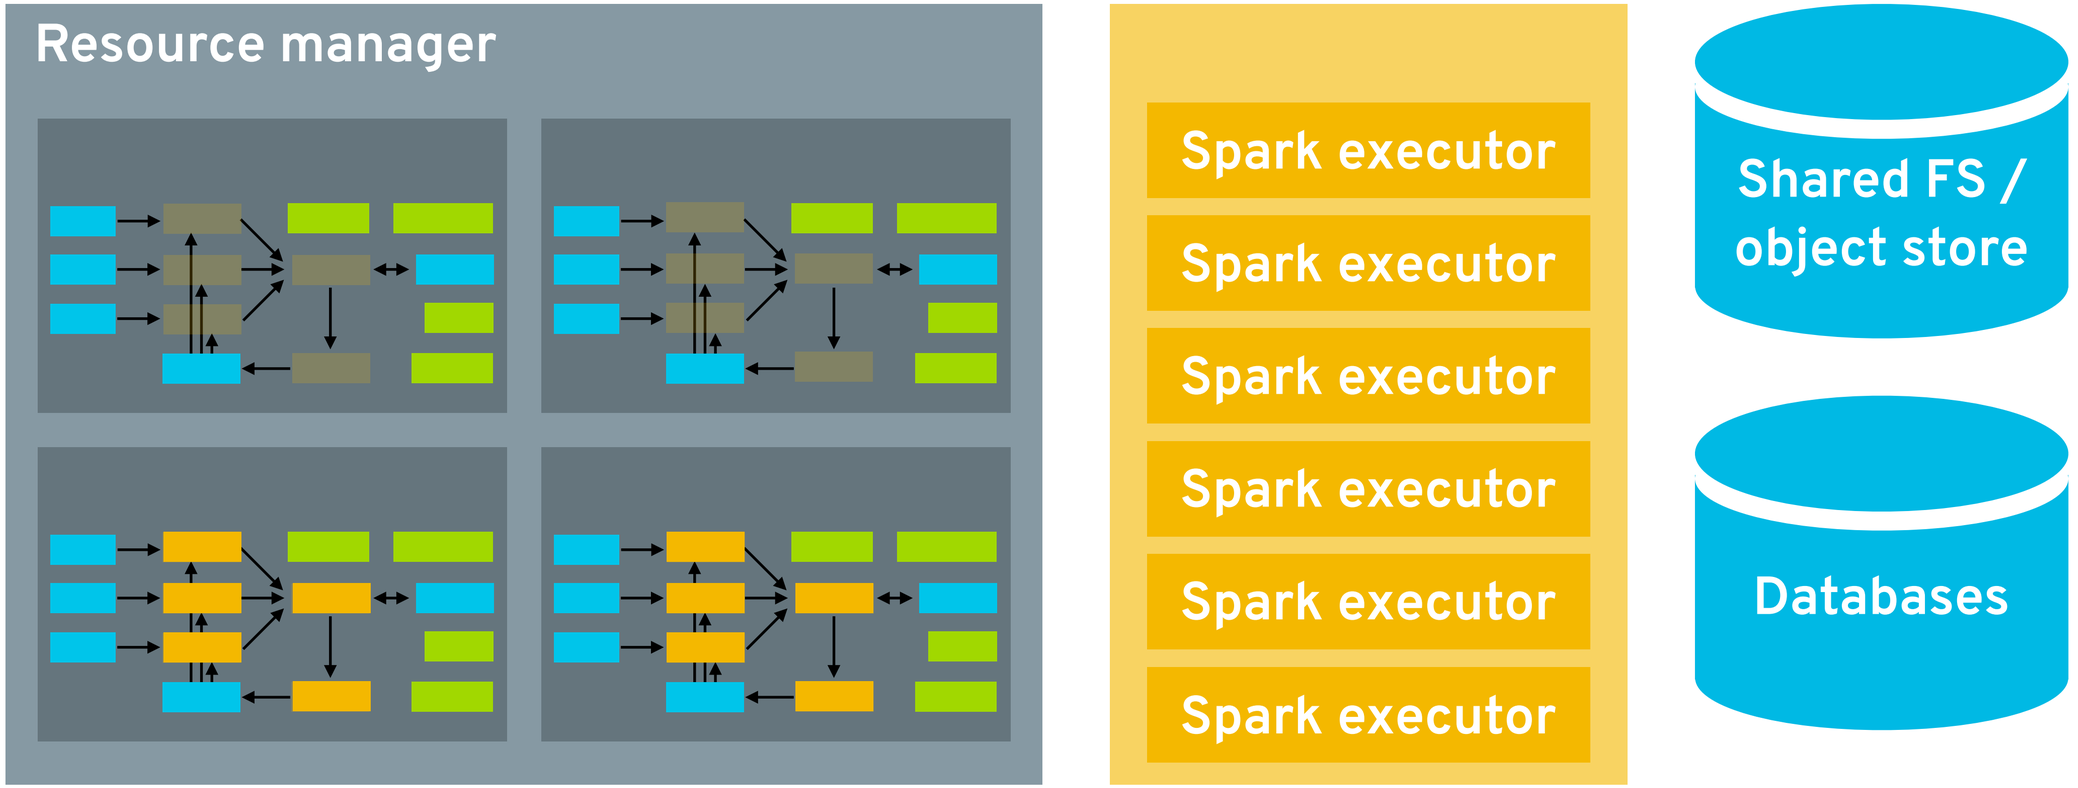

One cluster per application

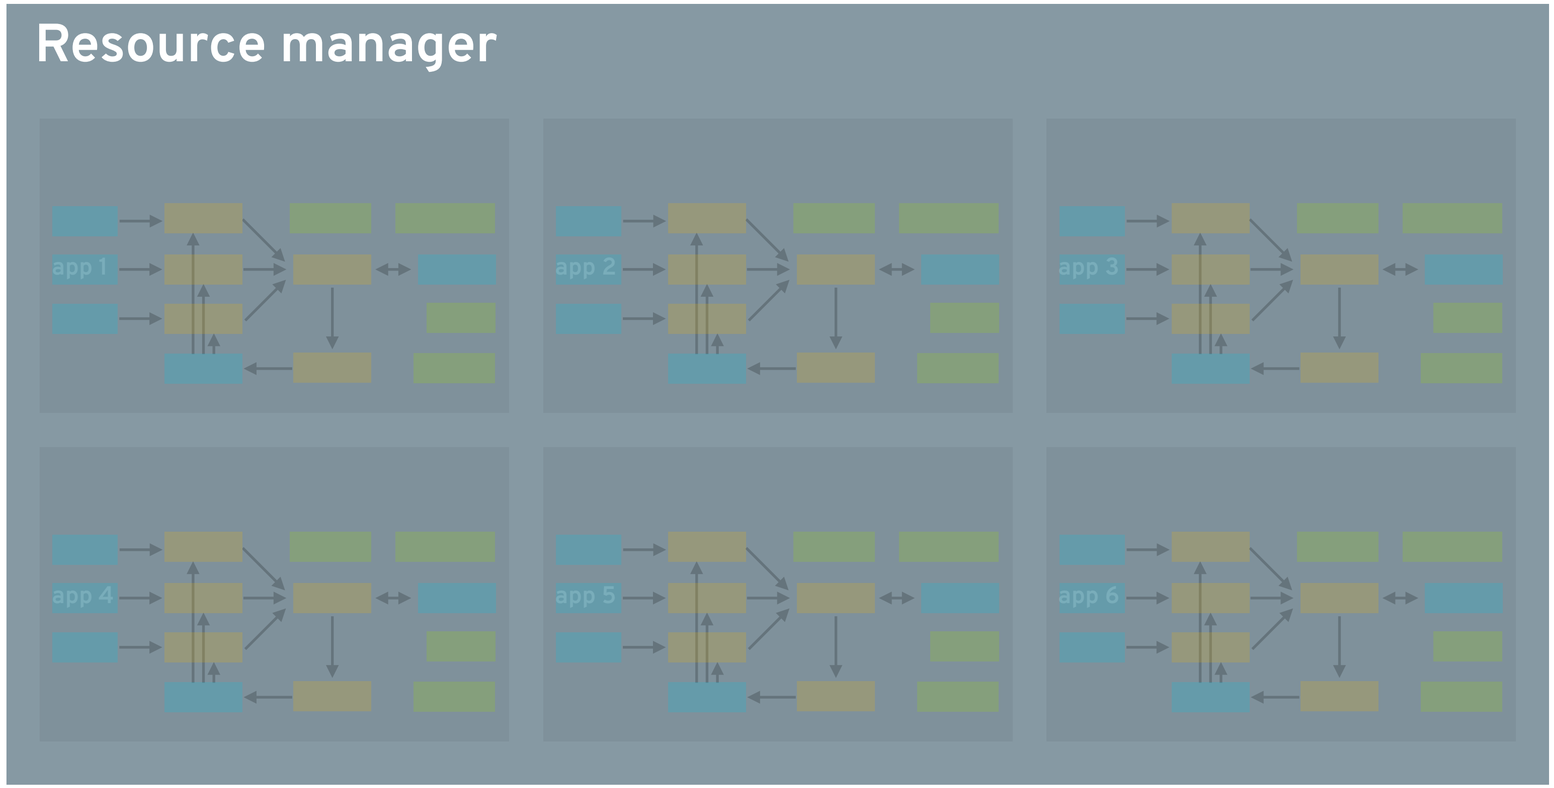

Shared FS / object store

Databases

One cluster per application

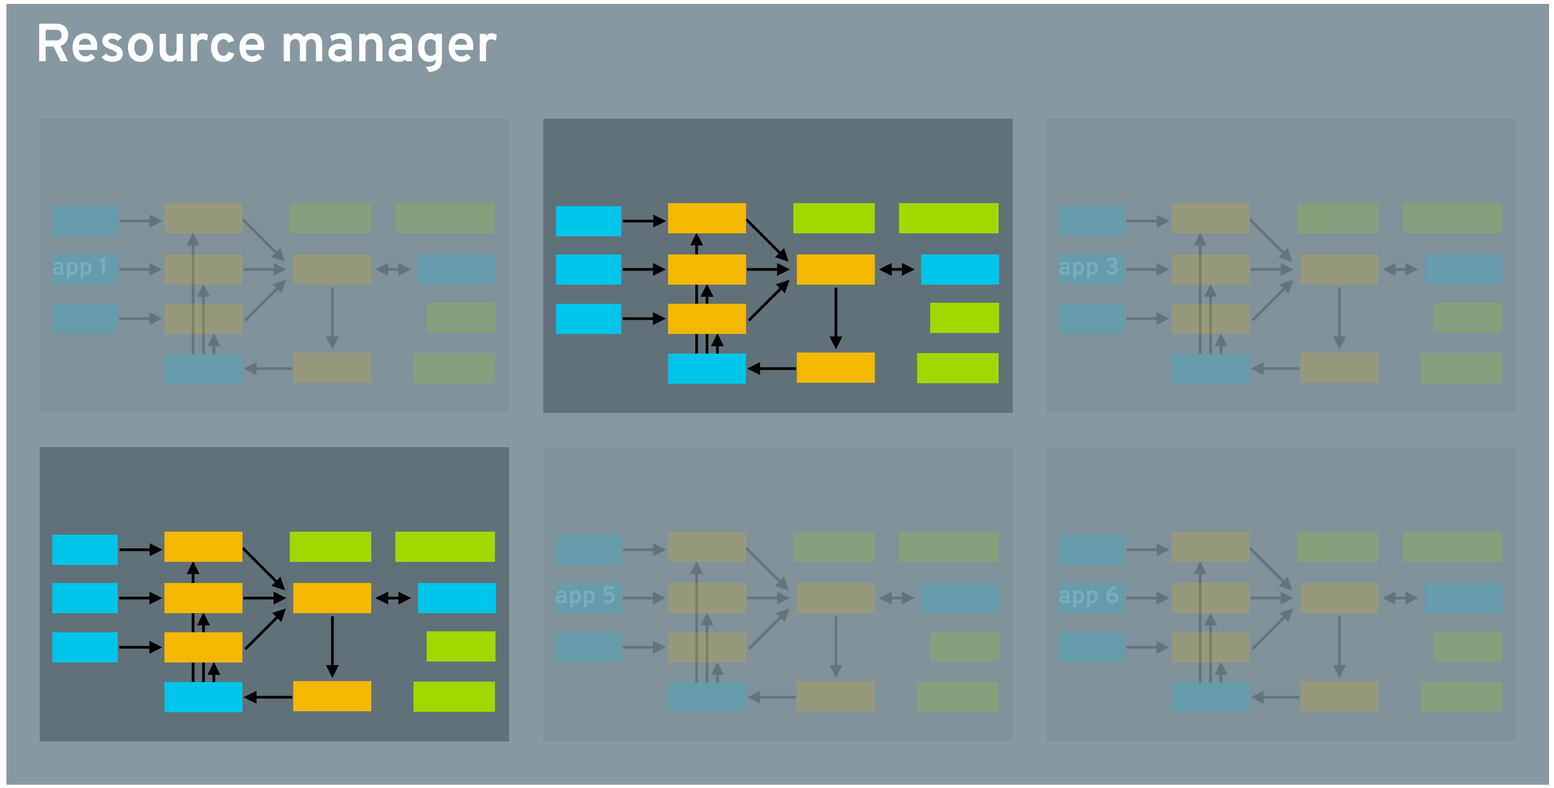

radanalytics.io

An **open-source community** enabling **intelligent applications** on **OpenShift**

Tooling to manage **Apache Spark**, **Jupyter notebooks**, and **TensorFlow training** and **model serving**

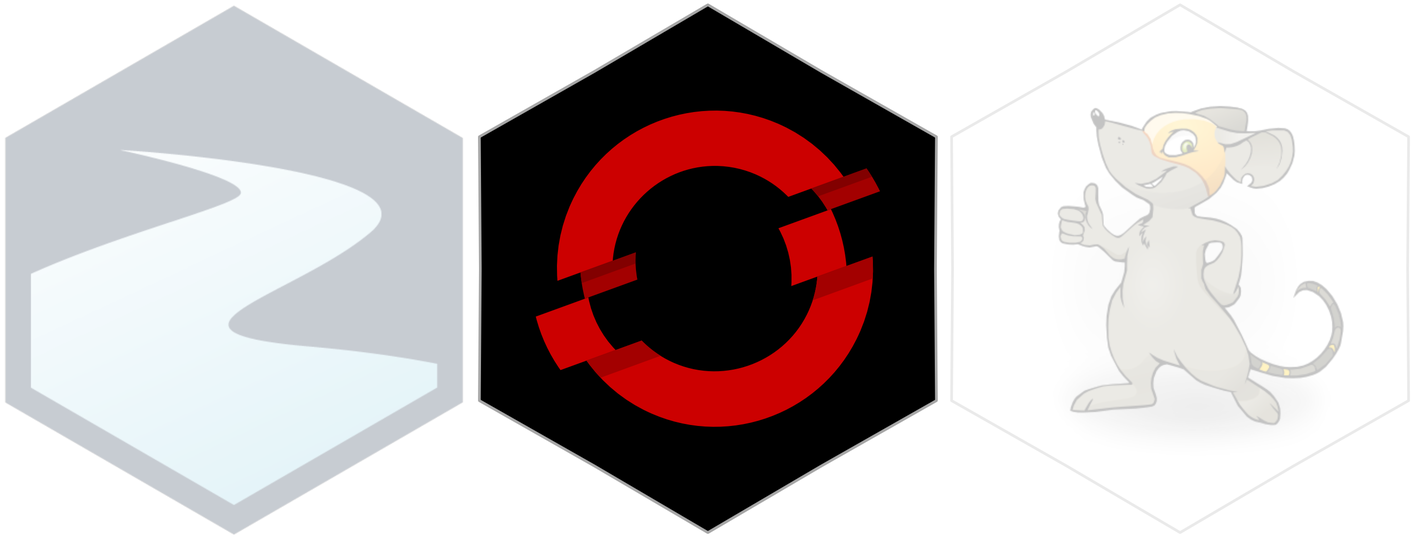

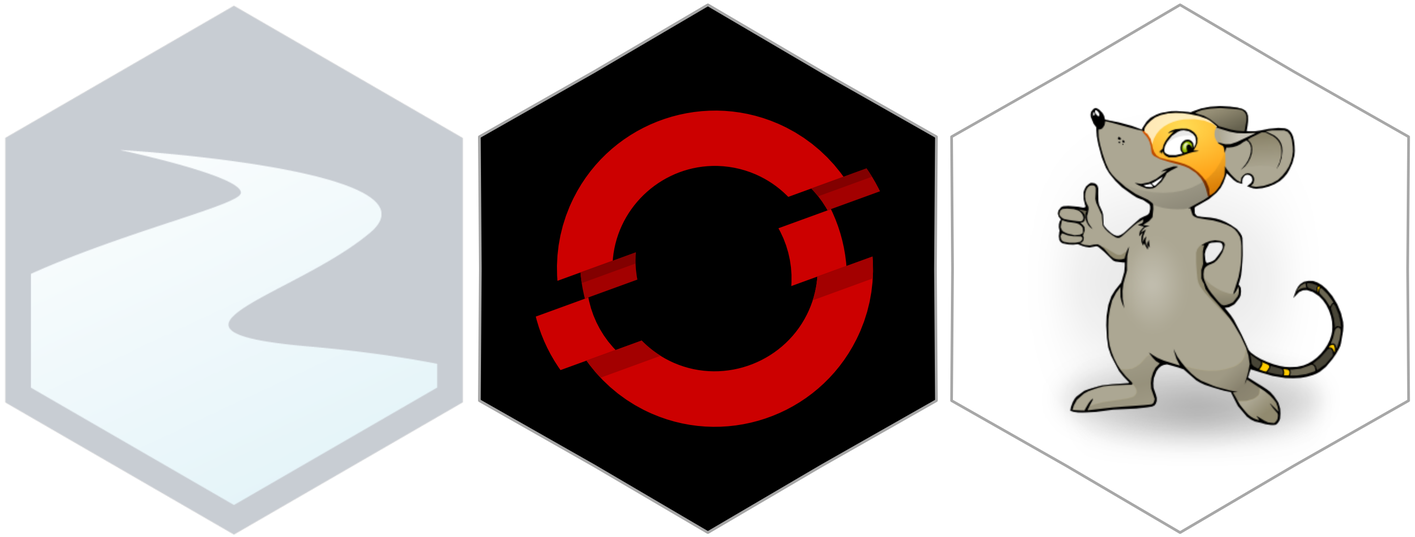

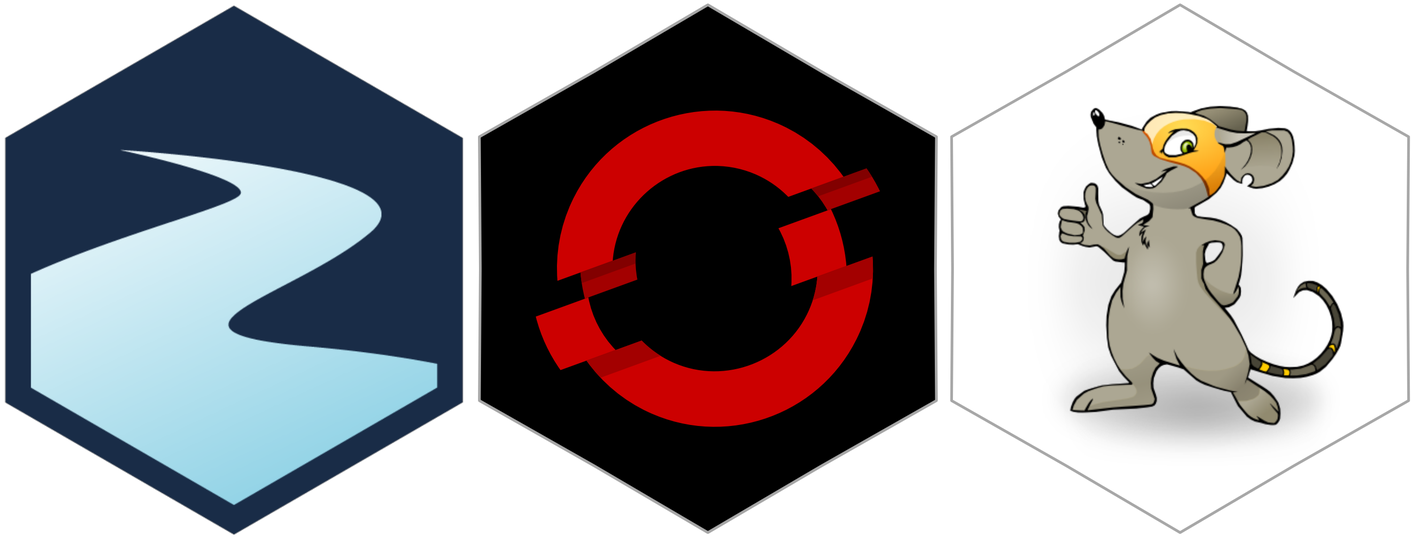

Properties of useful streaming algorithms

Interactivity time!

http://bit.ly/streaming-bds18

BIG DATA SPAIN

From notebooks and experiments to intelligent applications

Searching for solutions

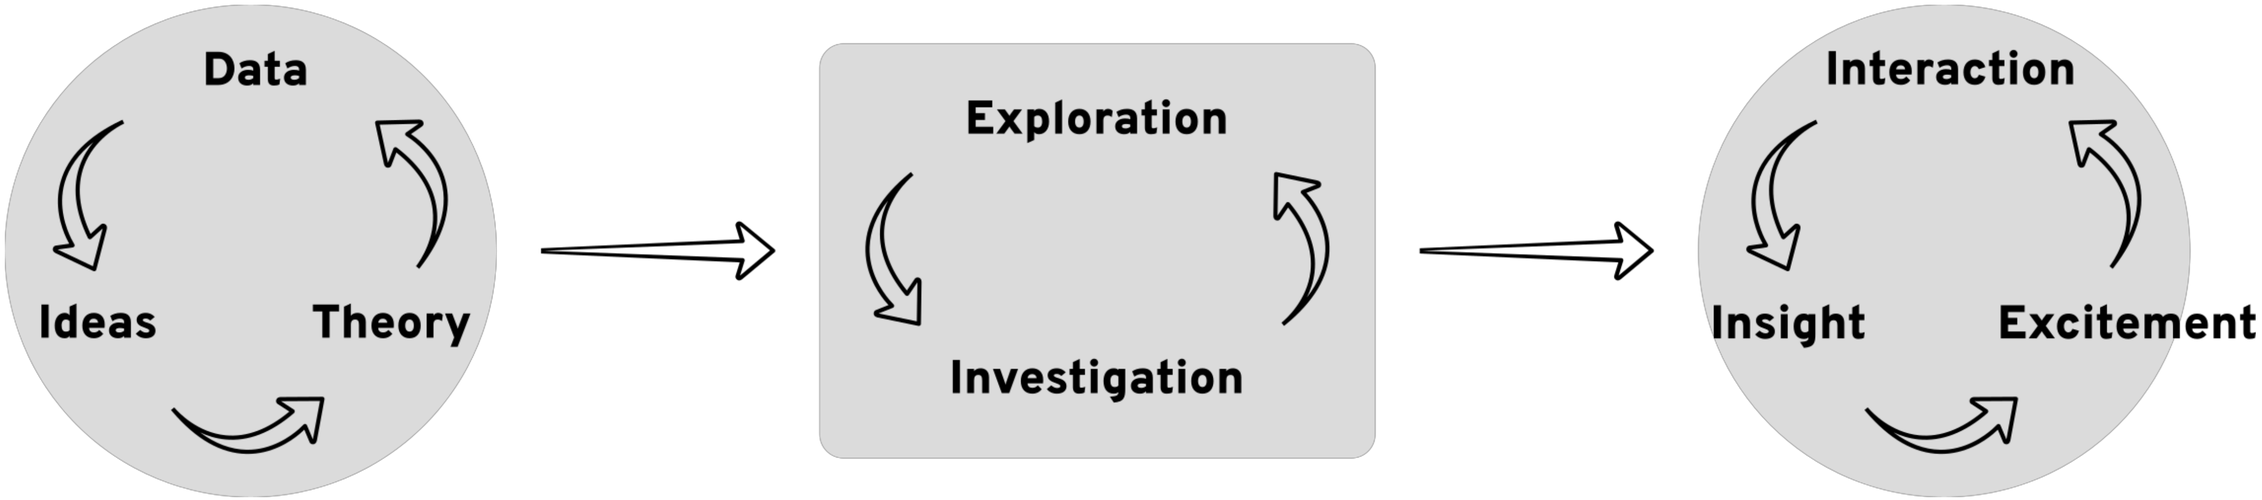

Consider possible interactions

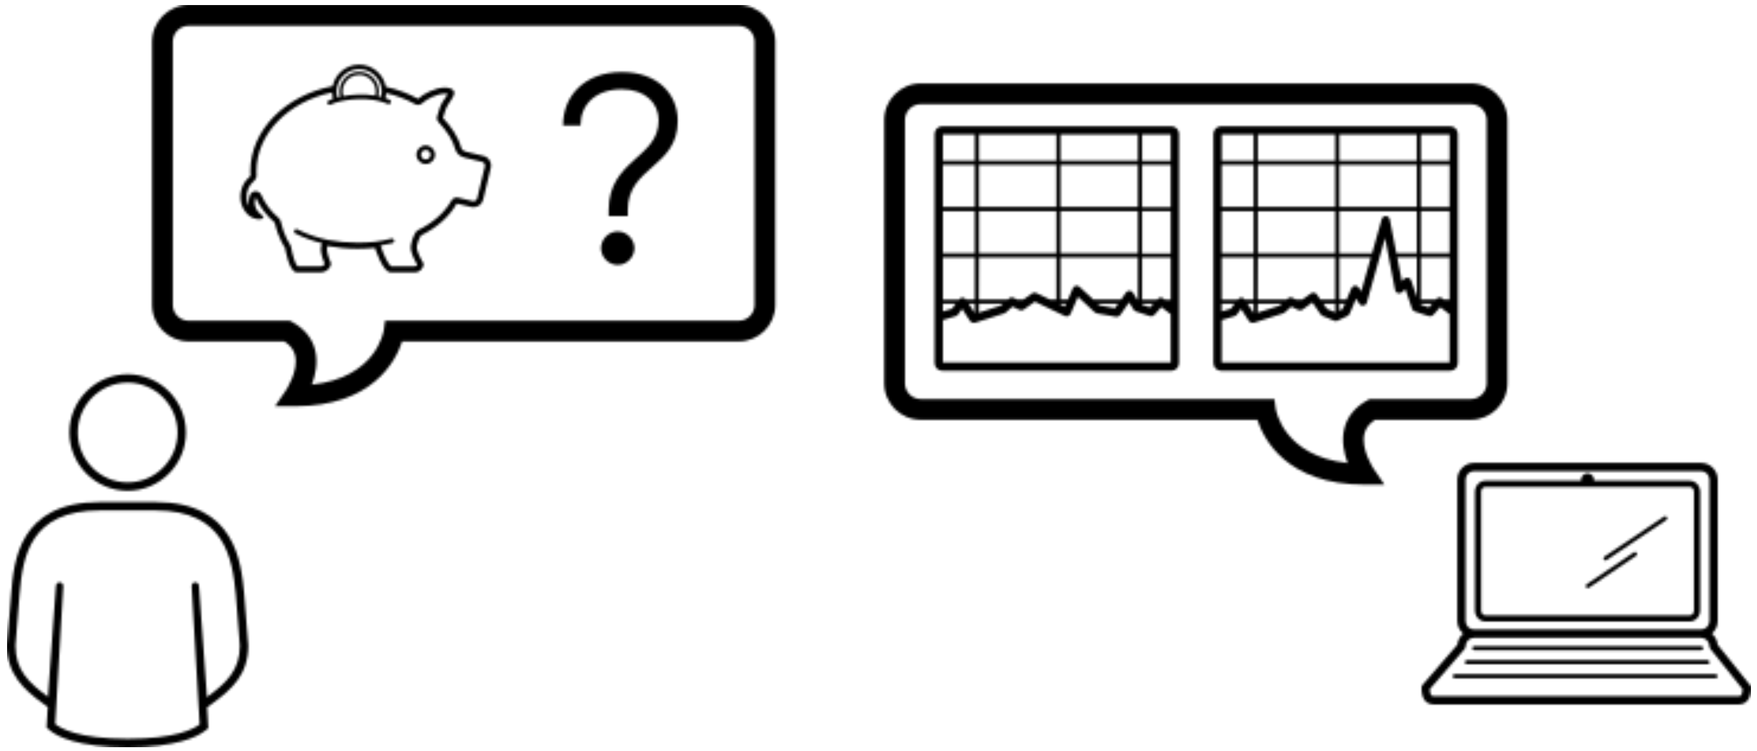

Plan for common patterns

Kappa in action

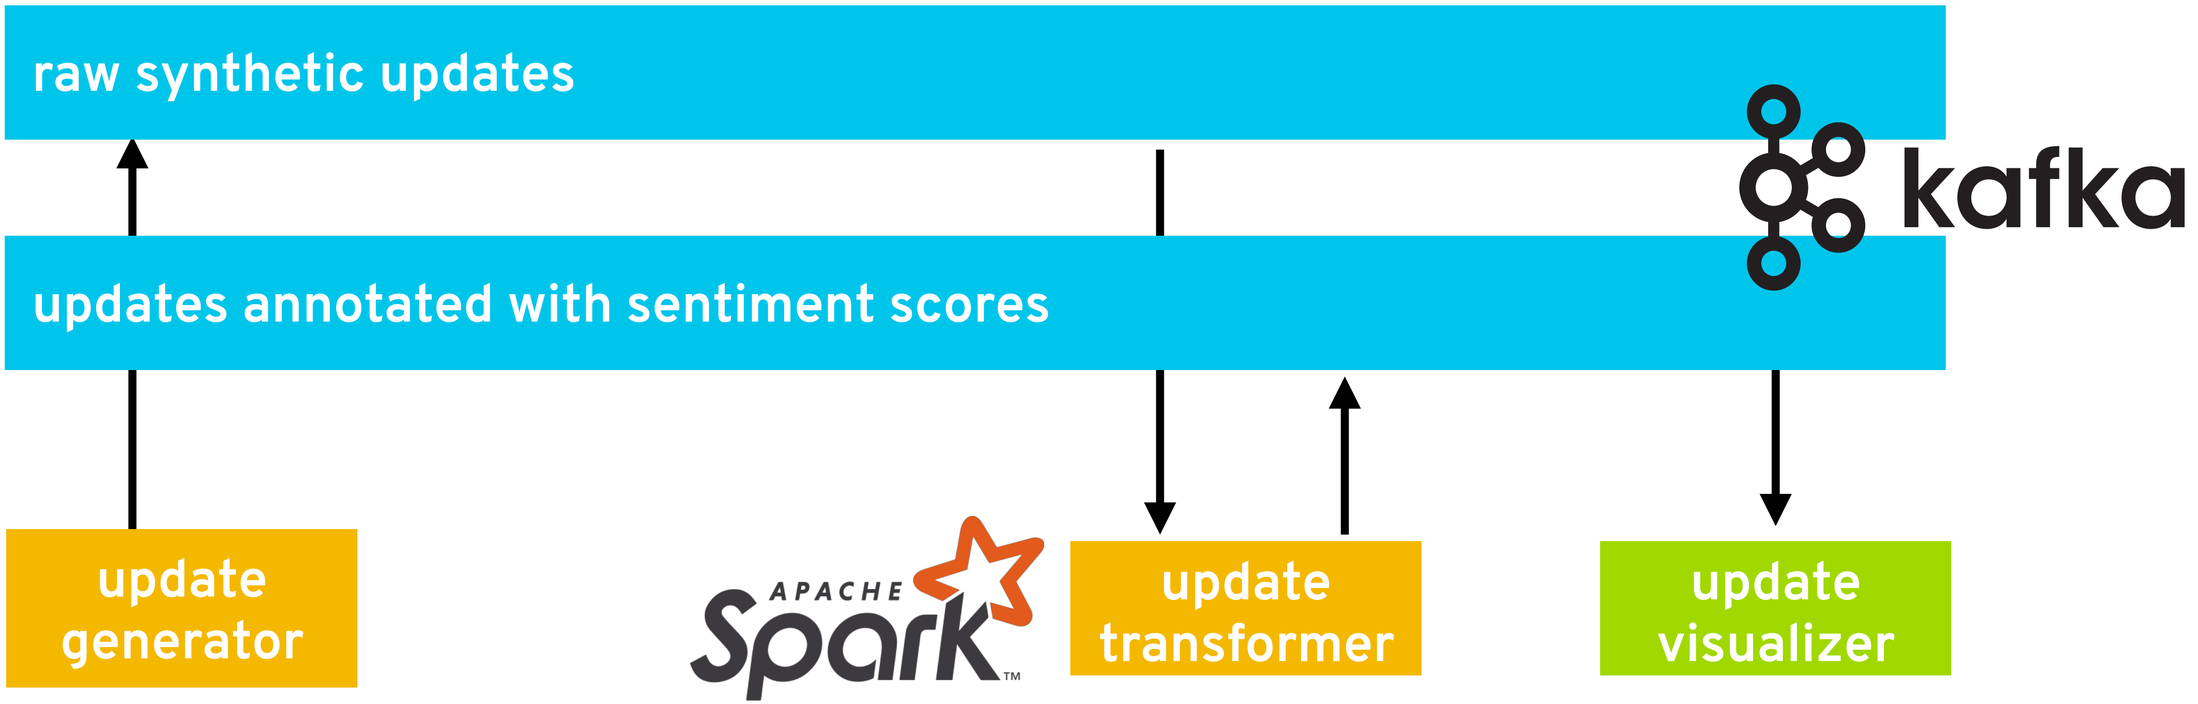

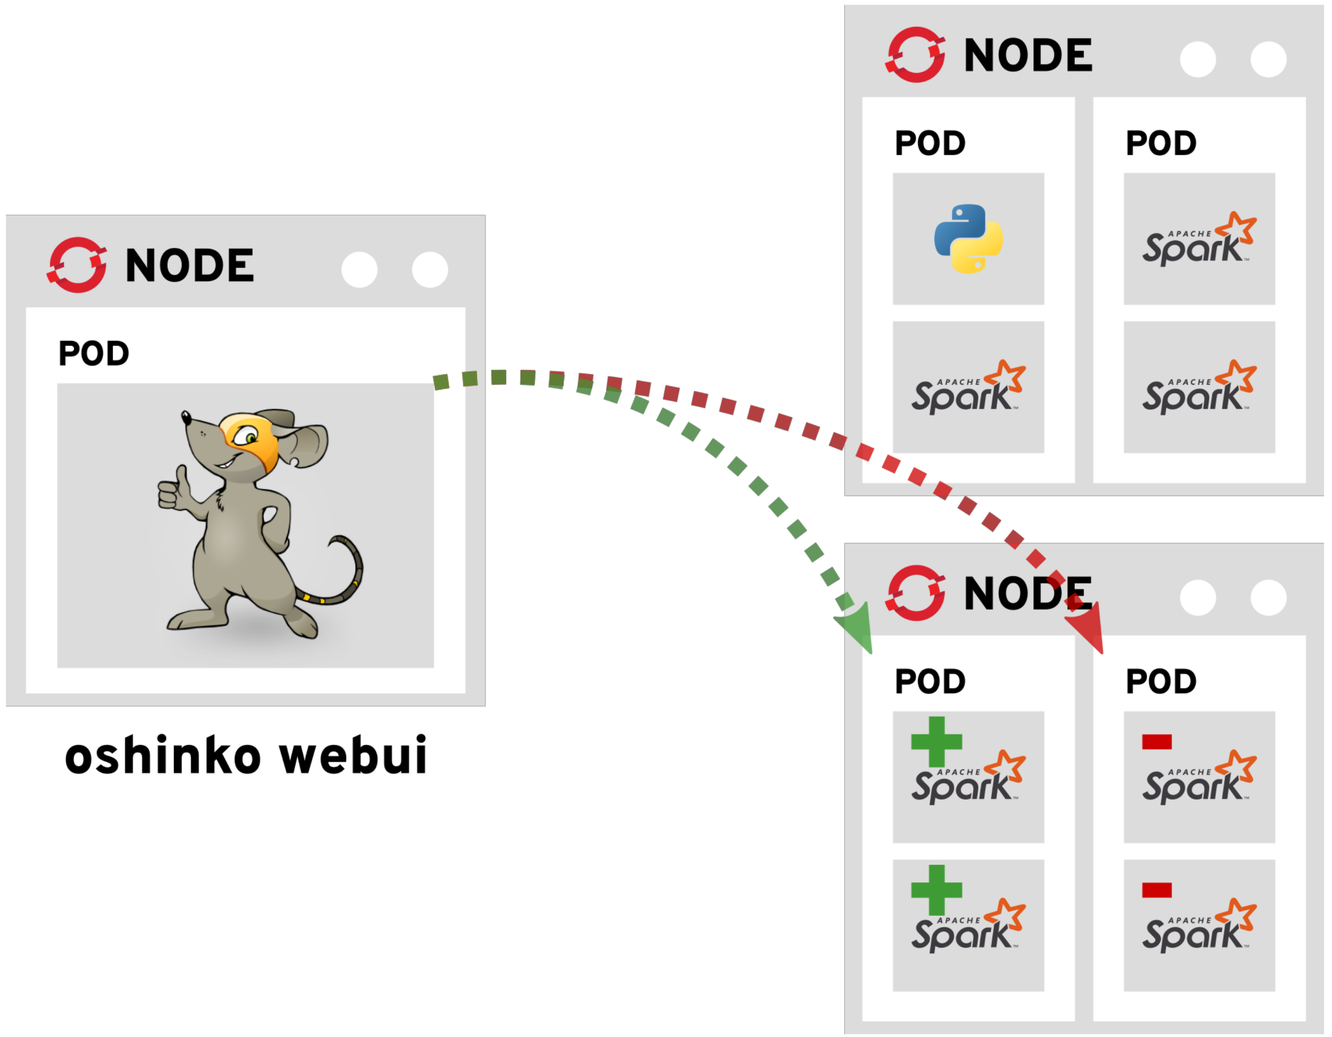
NODE
POD
oshinko webui
NODE
POD
POD
NODE
POD
POD

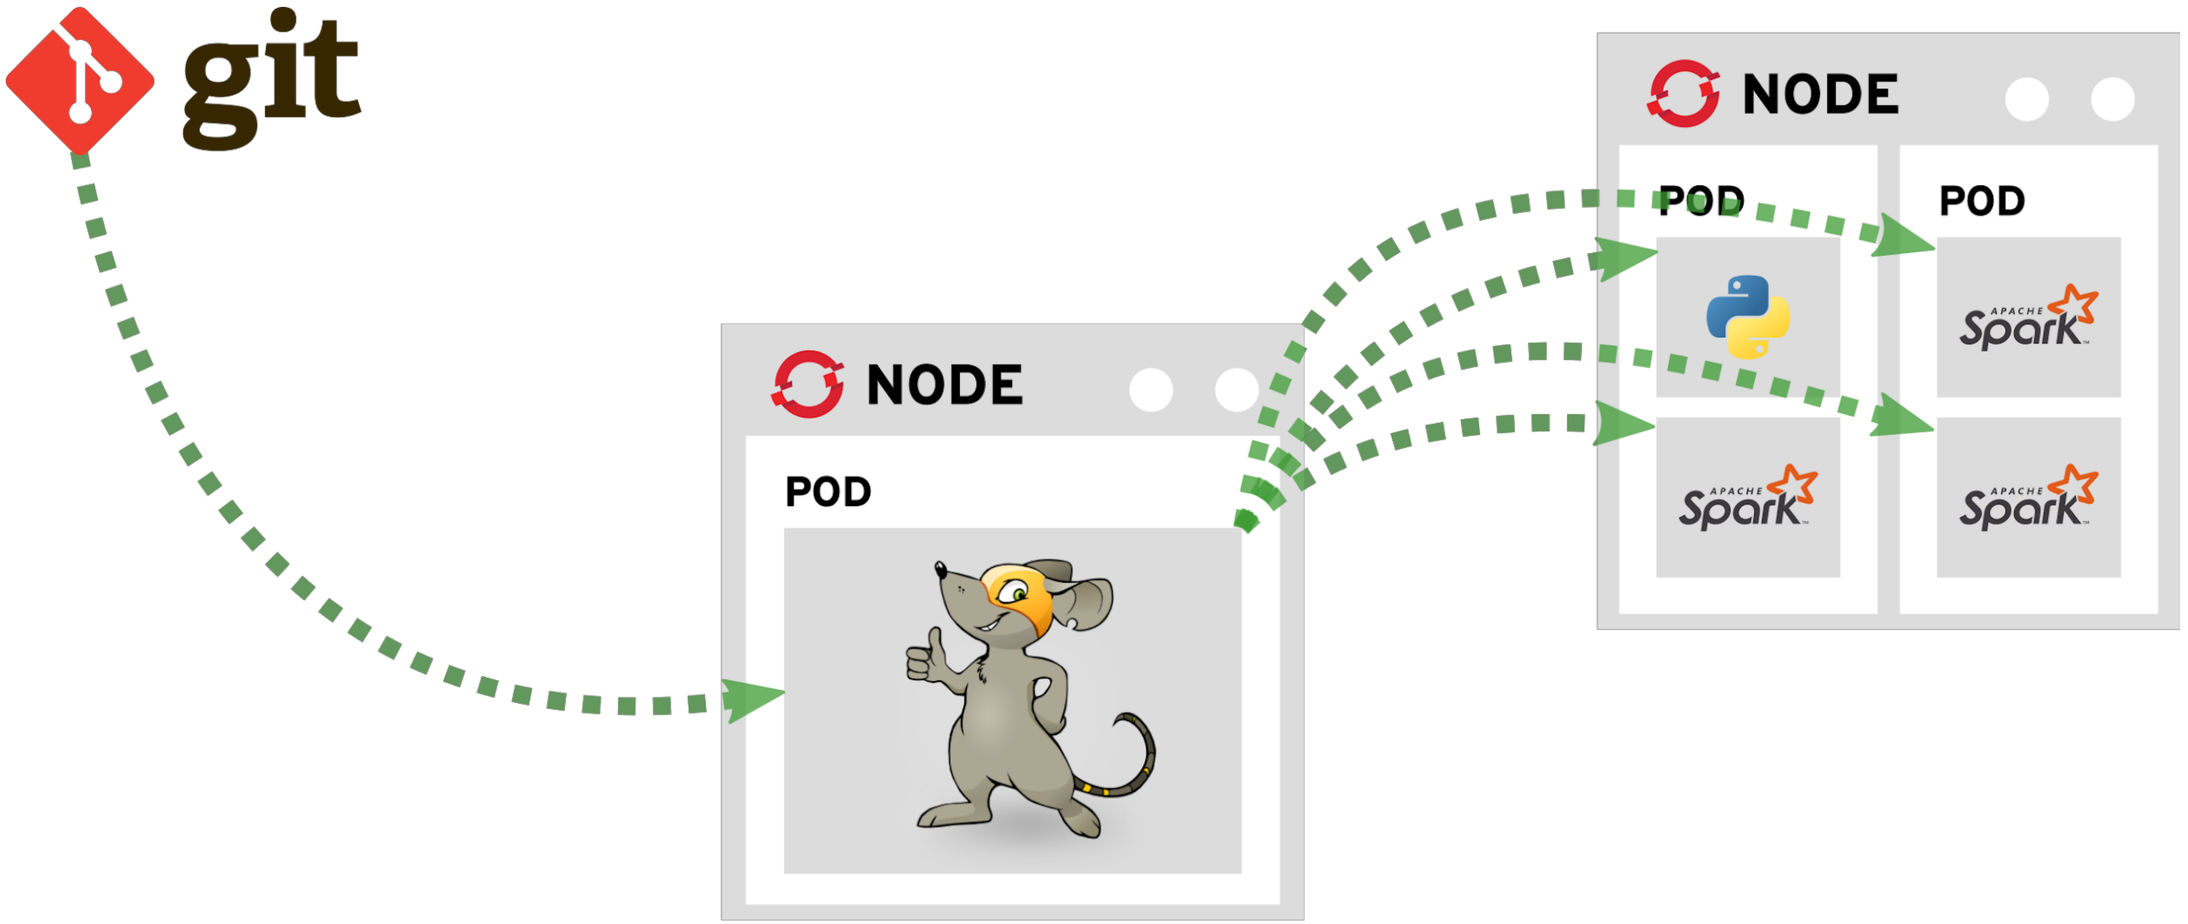

oshinko source-to-image